

Deep Learning

Equancy
November 7, 8, 9 2017

course given by
Hagop Boghazdekian



equancy

Nice to meet you!



Hagop Boghazdeklan
Data Scientist at Equancy
R&D in Deep Learning/Image Processing
Double M.Sc in Econometrics & Data Science

Disclaimer

- First steps in Deep Learning (DL)
- Goals:
 - Learning foundations of DL
 - Understanding/managing neural networks
 - Lead successfully DL projects
- You'll master not only the theory, but also see how to apply it in the industry by learning keras and tensorflow

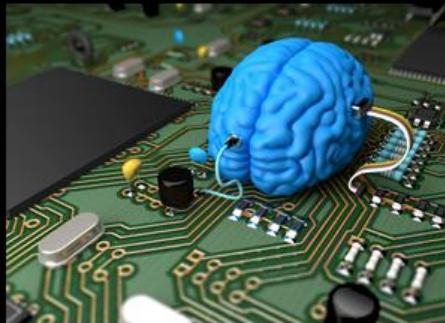
Disclaimer

- No need to be afraid, it's not that complicated !

Deep Learning



What society thinks I do



What my friends think I do



What other computer
scientists think I do



What mathematicians think I do



What I think I do

```
In [1]:  
import keras  
Using TensorFlow backend.
```

What I actually do

AGENDA

- ❖ Introduction
- ❖ Challenges Motivating Deep Learning: ML vs DL
- ❖ Theory of Deep Neural Networks
 - Neural Networks
 - Regularization for DL
 - Optimization for Training DL models
 - CNN
 - RNN
 - Transfer Learning, Reinforcement Learning
- ❖ Practicing : train your own networks using keras
 - Image Classification
 - Sentiment Analysis
 - Fisheries kaggle competition
- ❖ Conclusion



1

Introduction

What's Deep Learning ?

A subset of Machine Learning

ARTIFICIAL INTELLIGENCE

Machines that have some of human characteristics (very specific)

MACHINE LEARNING

Training a model using data so that it learns a specific task.

DEEP LEARNING

A particular way to train models that recently gave great results

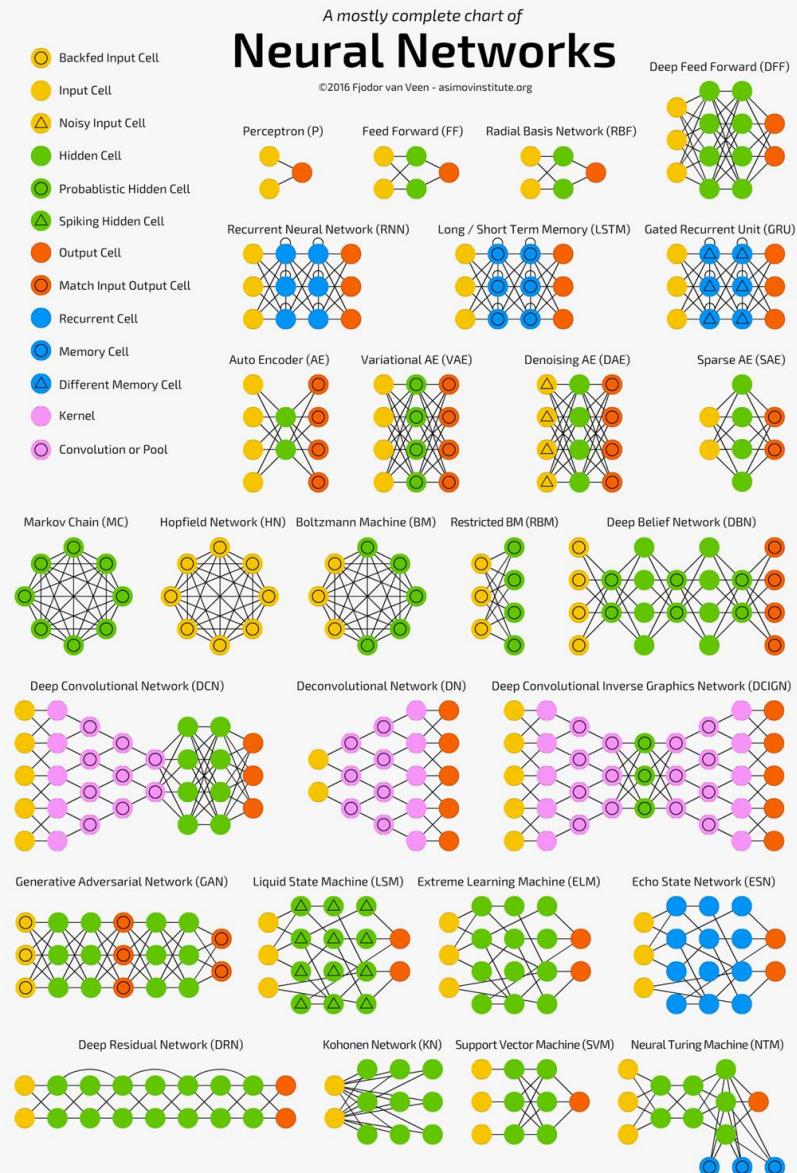


- Used when classical ML algorithms do not succeed
- To solve central problems in AI:
 - Image recognition
 - Natural Language Processing (Part of Speech, Named Entity Recognition, ...)

Neural Networks

The core of Deep Learning

- Using Deep Learning often implies training Deep Neural Networks
- Incredibly large number of possible networks ...
- Heavy theory & mathematical foundations
- Highly flexible & sensitive algorithms



Why Deep ?

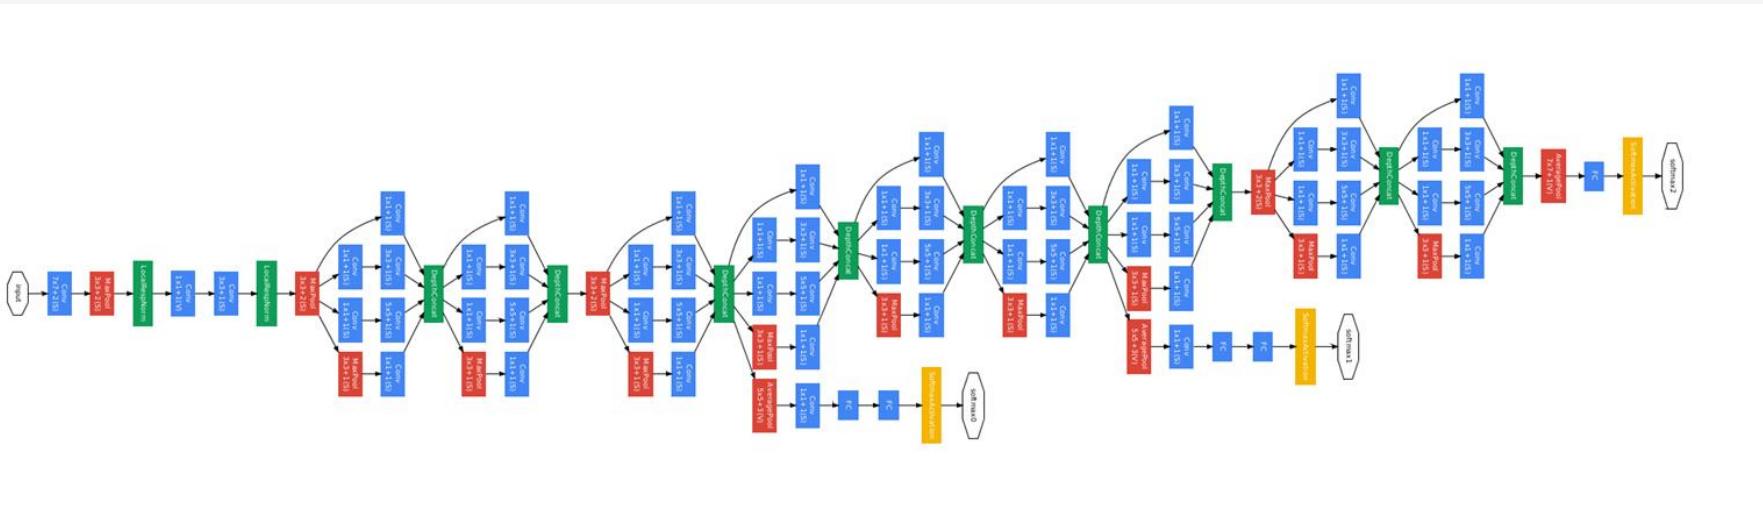
Complex and deep models

- The Deepness refers to the complexity of the Neural Networks
- To simplify! Each time we want to describe a modern deep network

It looks like this ...



GoogleNet - Inception



Who's using DL ?

All major companies



YAHOO!

Google



IBM



Baidu 百度

Who's using DL ?

Specialized start-ups ... it's booming!



60+ STARTUPS USING DEEP LEARNING

CORE AI: COMPUTER VISION

clarifai Tractable netra HYPERVERGE
captricity deepomatic sentient Sighthound
terraLOUPE imagenil piloPai

CORE AI: OTHER

TERADEEP affectiva* Wave Computing
DigitalGenius LEAPMIND Arya
heuro twentybn

BI, SALES & CRM

SKYMIIND TalkIQ intranetum ripjar
FRACTAL INDUSTRIES MARIANA DEEPGRAM
True AI KOVID ditto Gridspace

CORE AI: VOICE INTERFACE

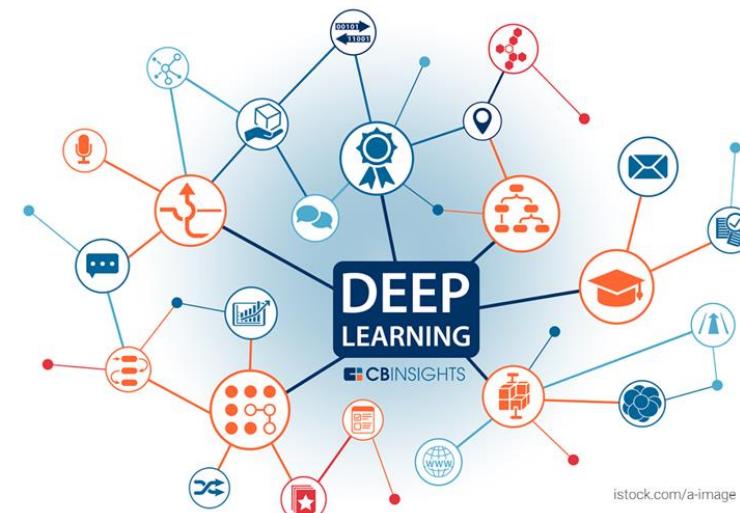
VIV Mobvoi Maluuba
OPEN

ROBOTICS & AUTO

netadyne AdasWorks
drive.ai Rokid 图灵机器人 TURING ROBOT
comma.ai

HEALTHCARE

imagiA Mindshare MEDICAL BAYLABS
BUTTERFLY Network, Inc. deep genomics Atomwise zebra MEDICALVISION
PATHWAY GENOMICS AVALON Behold.ai MediMatch
ADRENIO enlitic Lunit SIG TUPLE



SECURITY

deepinstinct Umbo CV
signalsense

OTHER

Alpaca Indico
Iris Automation

E-COMMERCE

Reflektion CORTEXICA
VISENZE STAQU

ACQUIRED (2014-2016YTD)

Etsy Blackbird Orbeus AlchemyAPI
intel Movidius nervana SYSTEMS
Perceptio vocaliq turi
GOO DEEPMIND DNNresearch Dark Blue Labs Vision Factory
MAGICPONY TECHNOLOGY MADBITS MetaMind

CB INSIGHTS

What for ?

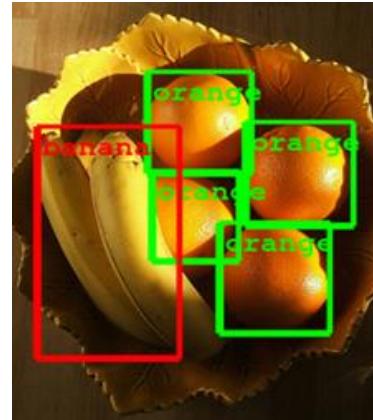
Image recognition

Image Classification



Persian cat : 0.45
Egyptian cat : 0.22
Hamster : 0.09
Lynx : 0.04

Object detection



Artist Imitations



Image Segmentation



Caption Generation



The man at bat readies to swing at the pitch while the umpire looks on.

Image Generation



What for ?

Natural Language Processing

Part of Speech Tagging

Text:

The grand jury commented on
a number of other topics

Legend:

article, adjective, noun,
preposition, verb

Named Entity Recognition

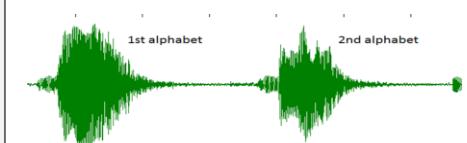
Text:

Apple CEO Tim Cook introduces
the latest iphone version at the
Cupertino Flint Center event in
California

Legend:

Organization, Location, Person

Speech Recognition

Audio:**Text:**

Winter is coming !

Translation

French

Bonjour à tous !

Chinese



Sentiment Analysis

Text 1:

Today was a great class about
Deep Learning ! #cepe #dl

**Text 2:**

We've had trouble understanding
what the guy was talking about
#equancy #cepe. It was bad.



Is it new ?

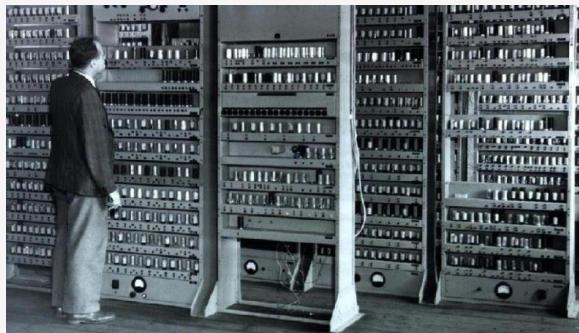
Not really ... but we can finally exploit it!

- Neural Networks - studied since the end of the 1950's
- Computationally heavy
- A reborn in the second part of 2000's:
 - Significant computational power
 - Explosion of Data
 - Highly implicated community

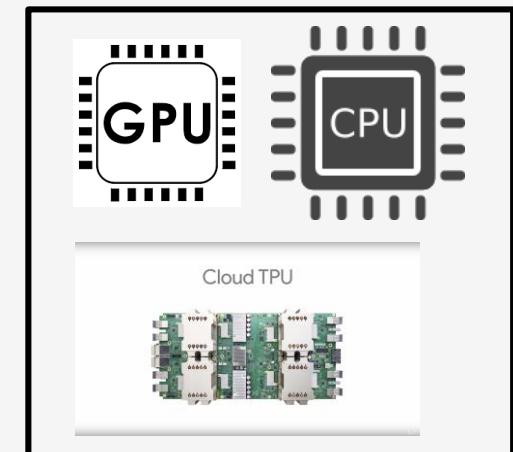
Is it new ?

Not really ... but we can finally exploit it!

Significant computational power



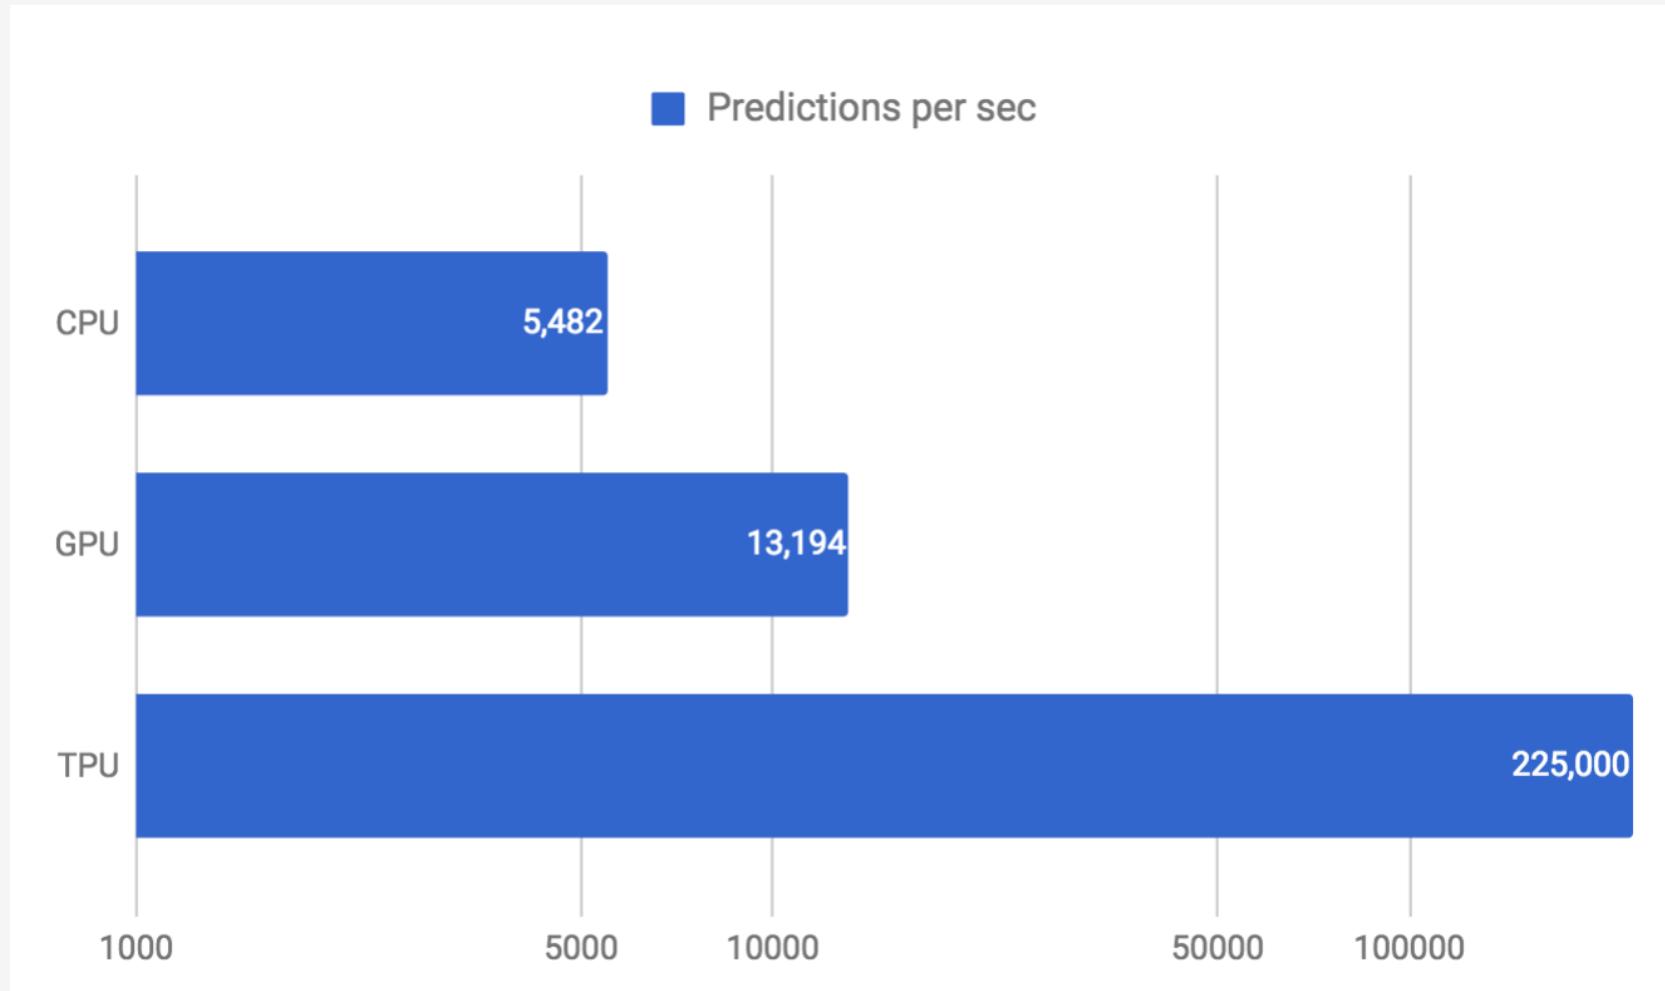
...



Is it new ?

Not really ... but we can finally exploit it!

Significant computational power
TPU ... always faster



Is it new ?

Not really ... but we can finally exploit it!

Explosion of DataSets



1627
Hosted
Datasets

276
Commercial
Competitions

1919
Student
Competitions

1M
Data
Scientist

4M
ML models
submitted

Is it new ?

Not really ... but we can finally exploit it!

Active and implicated community of deep learners
companies, researchers, scientists ...

large choice of available tools !





A black and white photograph of a Parisian skyline, featuring the iconic Eiffel Tower rising above the city's rooftops. The foreground shows the architectural details of several buildings, including a prominent church tower on the left.

2

Challenges motivating Deep Learning: ML vs DL

Reminders



Reminders

ML reminders/basics

- Remember DL is a subset of ML
- To understand DL, strong ML basics are necessary:
 - complexity, overfitting and underfitting
 - model selection & model generalization
 - regularization
 - bias/variance tradeoff
 - gradient-descent optimization
- Let's quickly review these key concepts

Reminders: Complexity, overfitting and underfitting

- Central challenge in ML: perform well on *new unseen* inputs (i.e. **generalization**)
- In optimization \Rightarrow Minimize the **training error**
- In ML \Rightarrow Minimize **test/generalization error** as well
- In other terms, the goals are to get:
 1. smallest training error
 2. smallest gap between training and test error
- Let's define $L(y, f(x))$ and $EPE \equiv E[L(y, f(x))|x = x_0]$

Reminders: Complexity, overfitting and underfitting

- Let's take the example of MSE : $L(y, f(x)) = \frac{1}{N} \sum_i^N (y_i - \hat{y}_i)^2$
- Training error** can be expressed as:

$$\frac{1}{N} \sum_i^N (y_i - \hat{y}_i)^2$$

- Test error** can be expressed as:

$$EPE = \sigma^2 + [Bias^2(\hat{f}(x_0)) + Var_{\tau}(\hat{f}(x_0))]$$

- How to assess these error terms ?

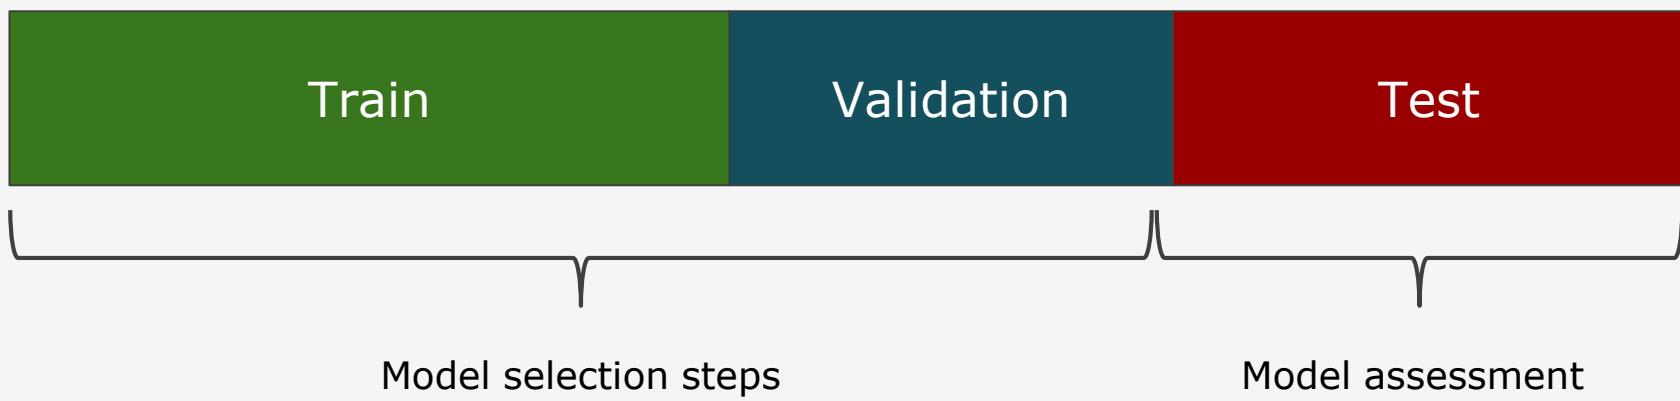
Reminders: Complexity, overfitting and underfitting

In rich data context

1 - Train all models
on the train set

2 - Estimate the prediction
error of all models

3 - Estimate the
generalization error
of the final model



$$\text{training error} = \frac{1}{N} \sum_i^N (y_i - \hat{y}_i)^2$$

$$EPE = \sigma^2 + [Bias^2(\hat{f}(x_0)) + Var_{\tau}(\hat{f}(x_0))]$$

Reminders: Complexity, overfitting and underfitting

K-folds cross-validation



Reminders: Complexity, overfitting and underfitting

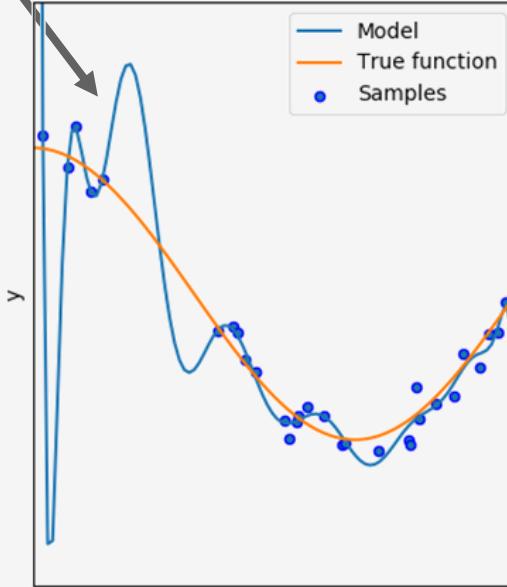
- **Underfitting:** "model unable to obtain a sufficiently low training error"
- **Overfitting:** "training error sufficiently low, but gap between training and test errors is too large"
- **Complexity:** model complexity (i.e. capacity)

Reminders: Complexity, overfitting and underfitting

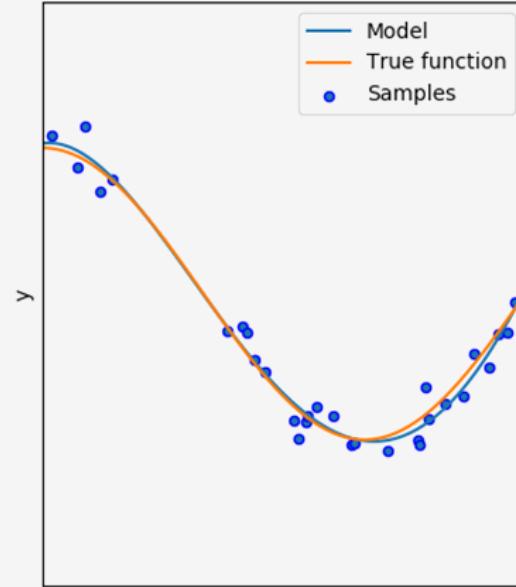
Illustration: fitting a polynomial regression

Overfitting

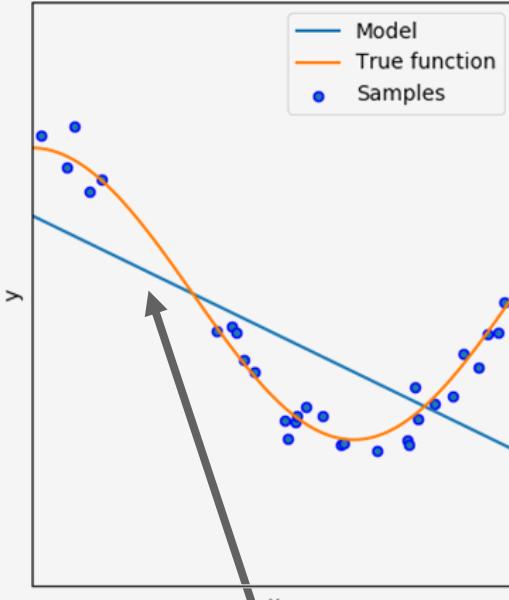
Degree 15
MSE = 1.82e+08(+/- 5.45e+08)



Degree 4
MSE = 4.32e-02(+/- 7.08e-02)



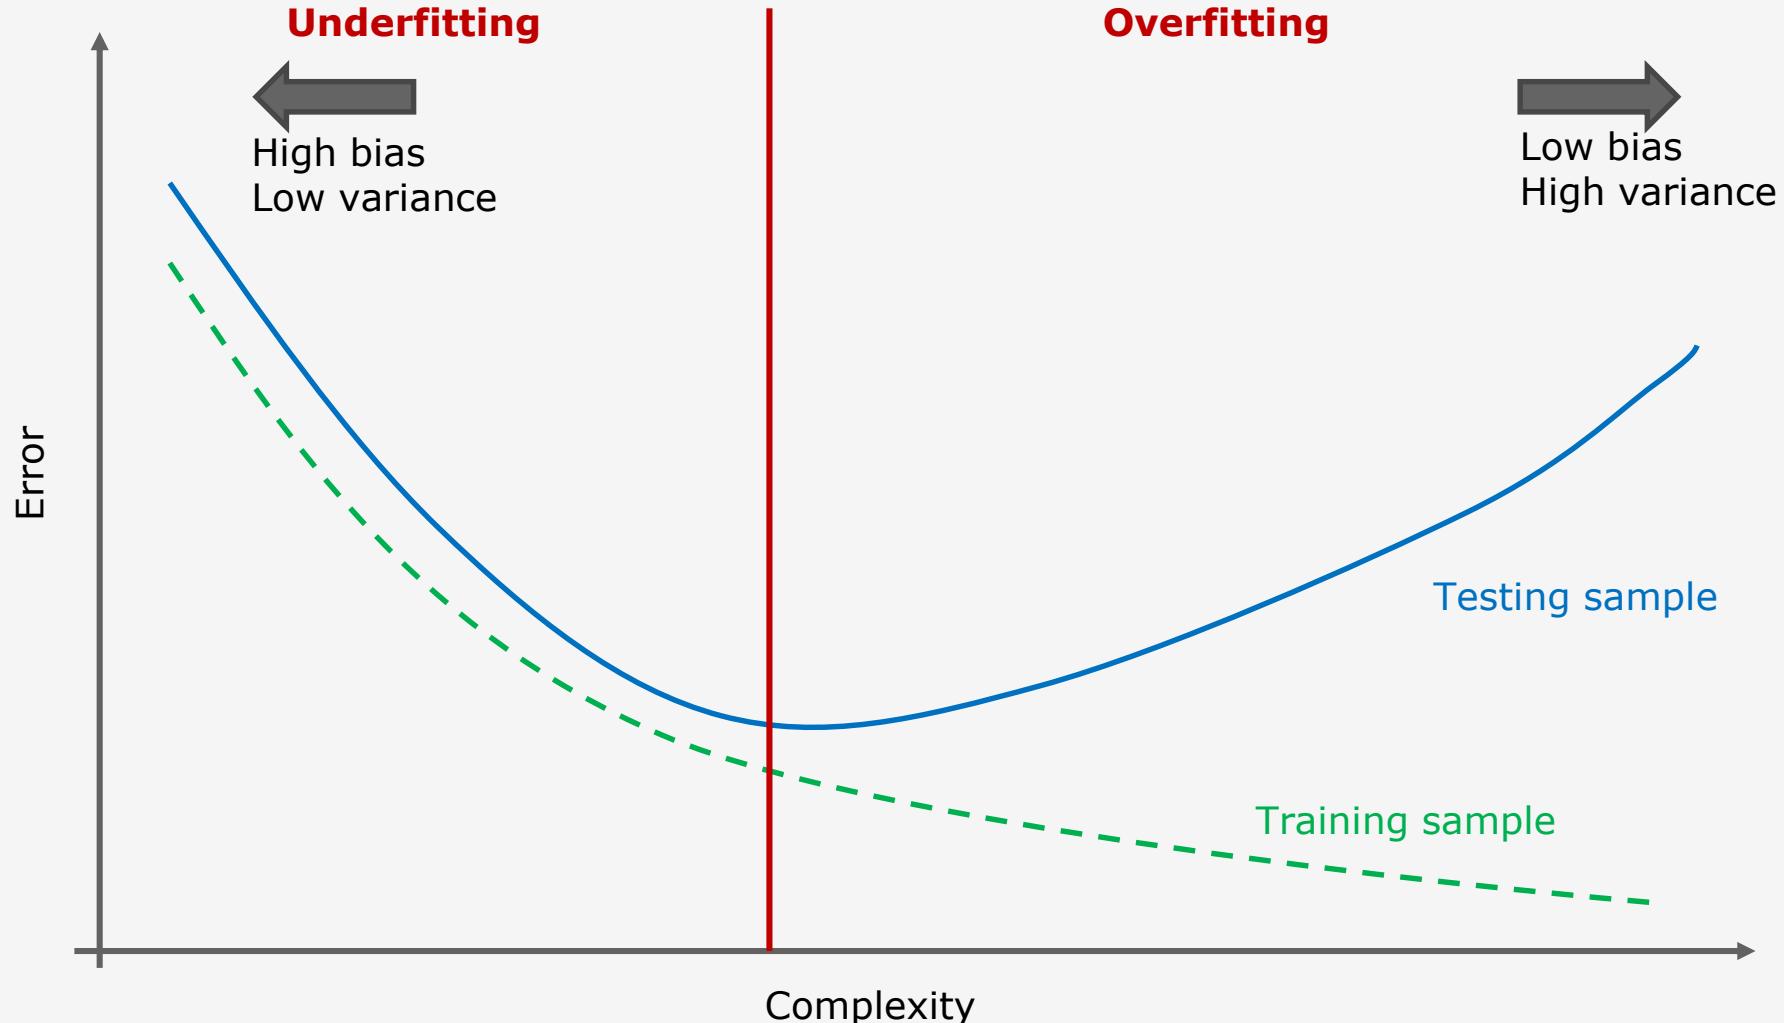
Degree 1
MSE = 4.08e-01(+/- 4.25e-01)



Underfitting

Reminders: Complexity, overfitting and underfitting

Overfitting and underfitting are linked with complexity



Reminders: Regularization

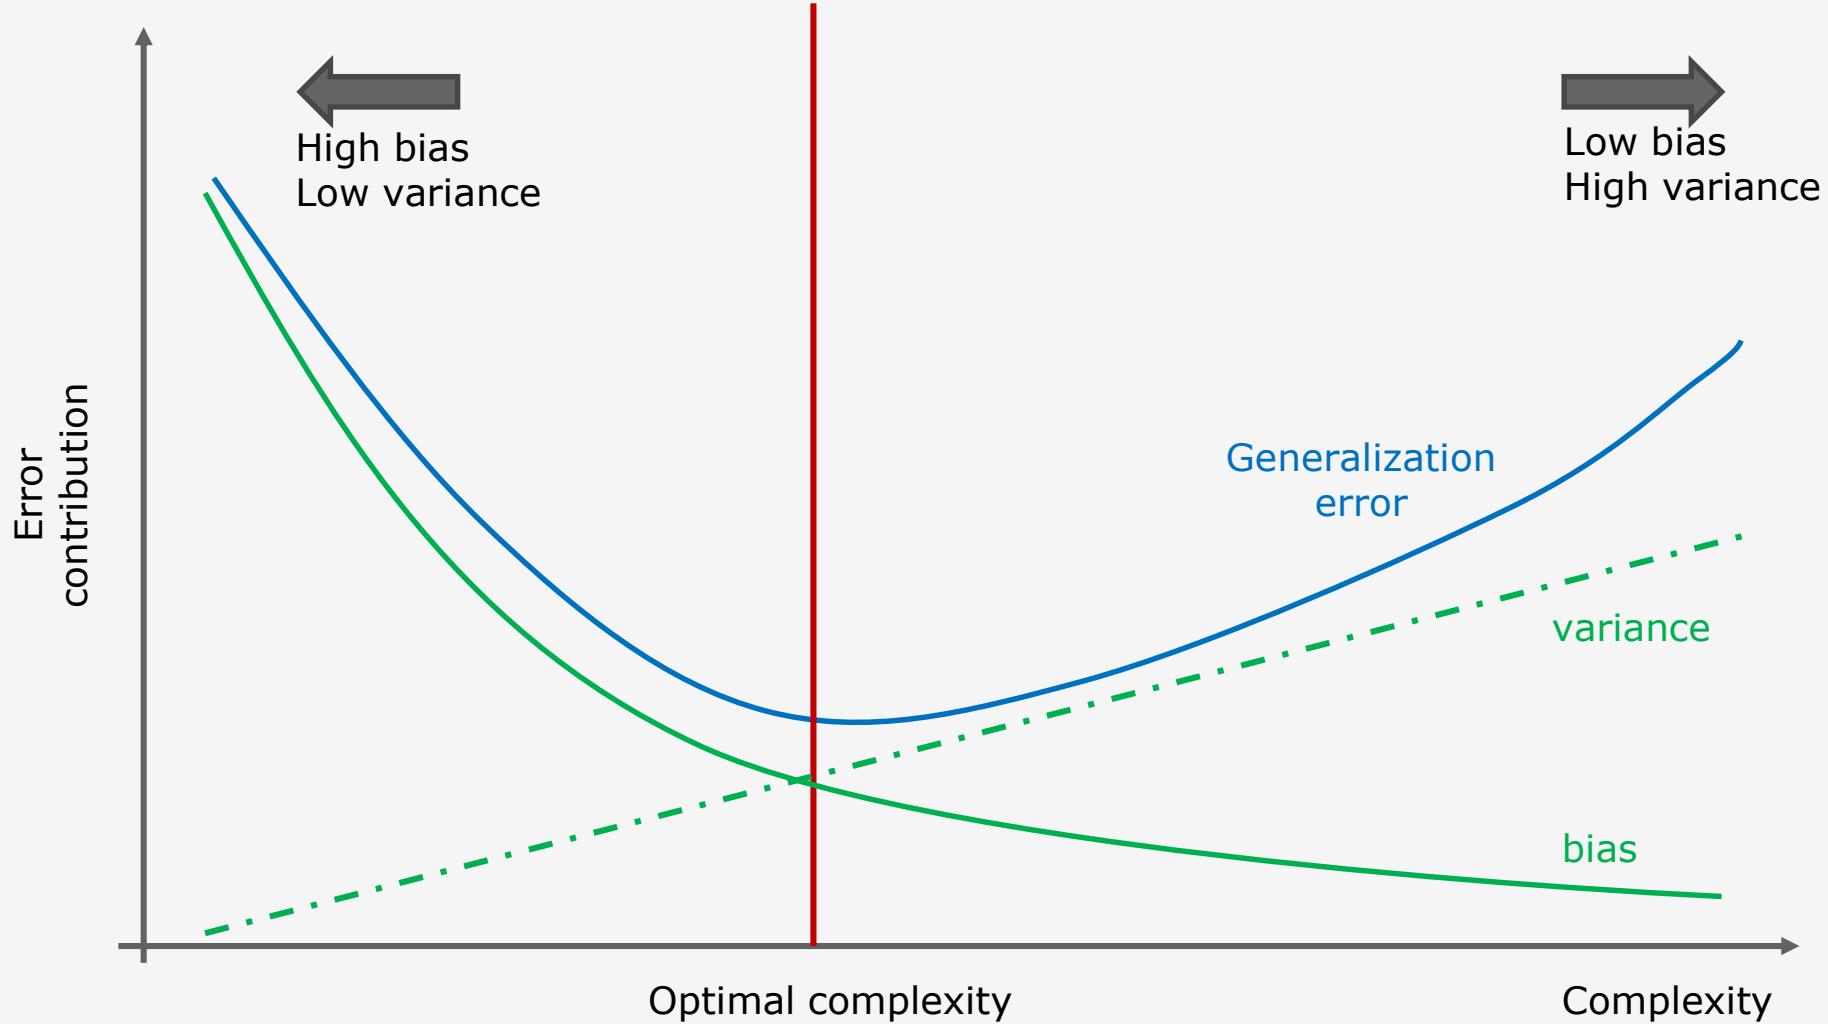
- To tackle overfitting and underfitting:
 - Increase or decrease the model capacity
 - Increase the training examples (more data)
 - Regularization
- **Regularization:** “any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error”
- Regularization to reduce overfitting

$$L(y, f(x)) = MSE_{train} + \lambda w^T w$$

- L1 or L2 regularization

Reminders: bias/variance tradeoff

A question of 'arbitrage'



Reminders: Gradient Descent optimization

Basic Gradient Descent

- Gradient-based optimization is very important in Deep Learning
- Gradient Descent to find *Global Minimum* of a given function with parameters
- Let's consider a cost/loss function

$L(\theta)$ with input parameters $\theta \in \Theta$

- Such as

$$\min L(\theta)$$

- Gradient Descent:

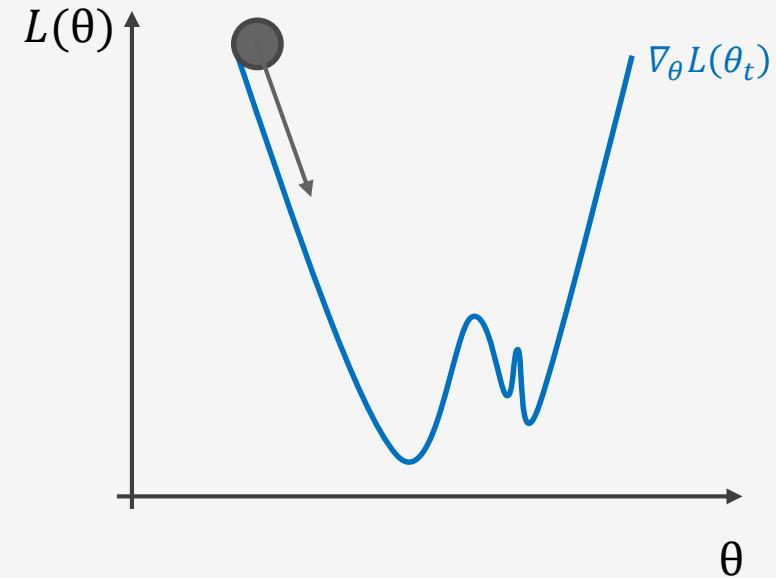
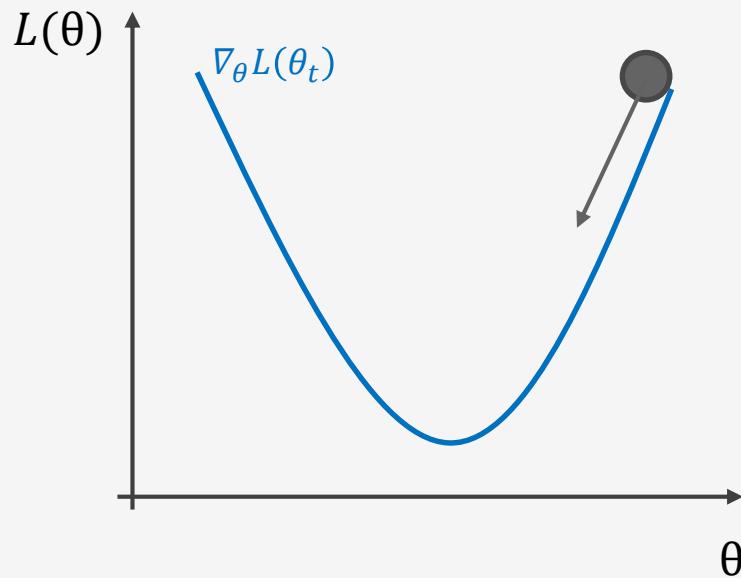
$$\theta_{t+1} = \theta_t - \lambda \nabla_{\theta} L(\theta_t)$$

with $\nabla_{\theta} L(\theta_t)$ as the gradient of $L(\theta)$ w.r.t. to θ
and λ as the learning rate

Reminders: Gradient Descent optimization

Basic Gradient Descent

See it as a ball rolling down ...

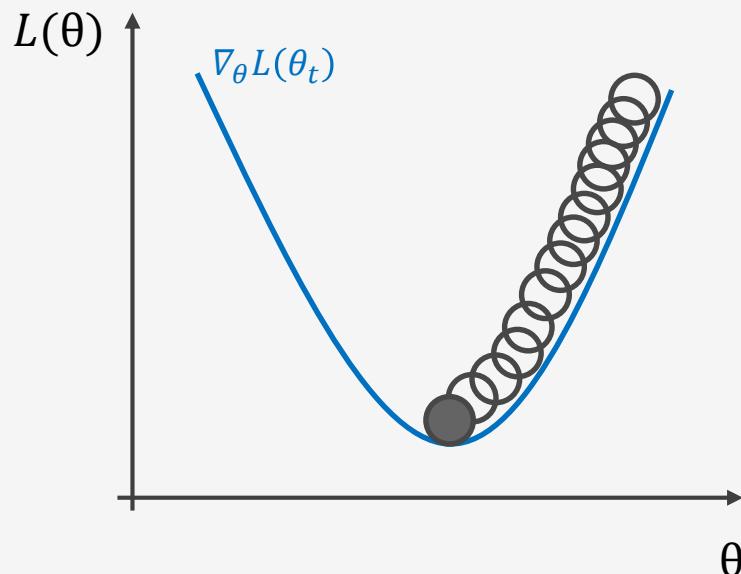


Reminders: Gradient Descent optimization

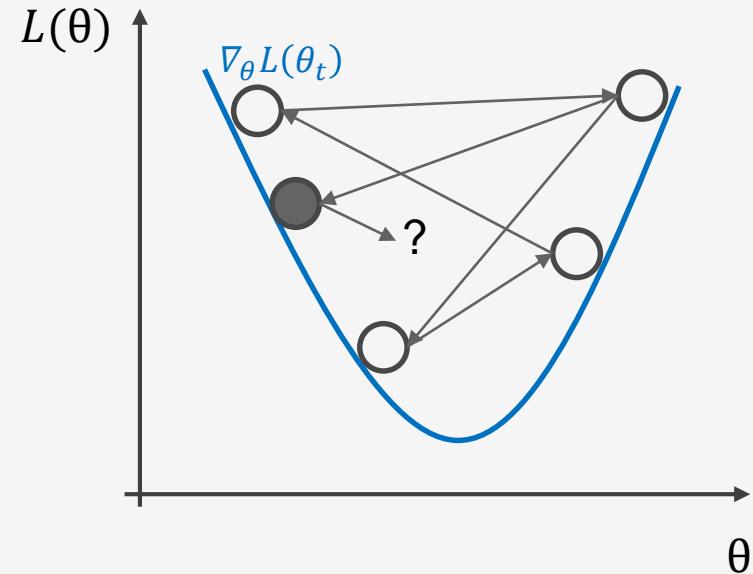
Basic Gradient Descent

Choice of the learning rate λ

λ too **small** Descent too **slow**



λ too **large** Target **missed**

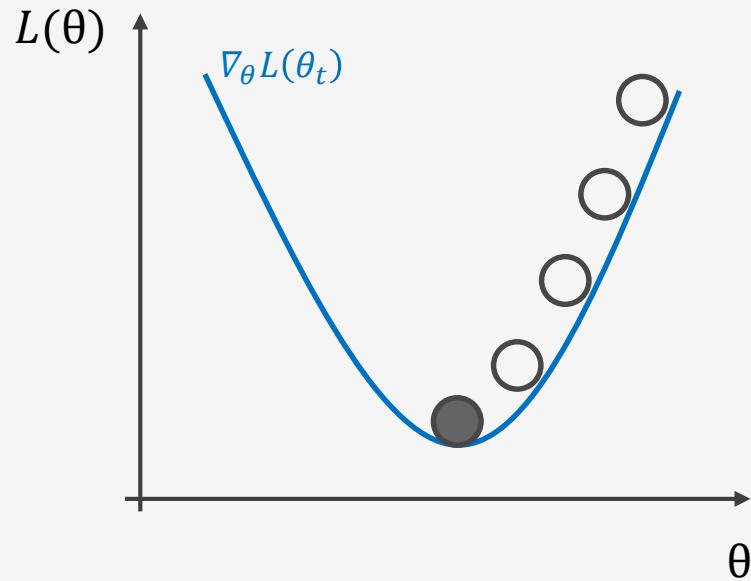


Reminders: Gradient Descent optimization

Basic Gradient Descent

Choice of the learning rate λ

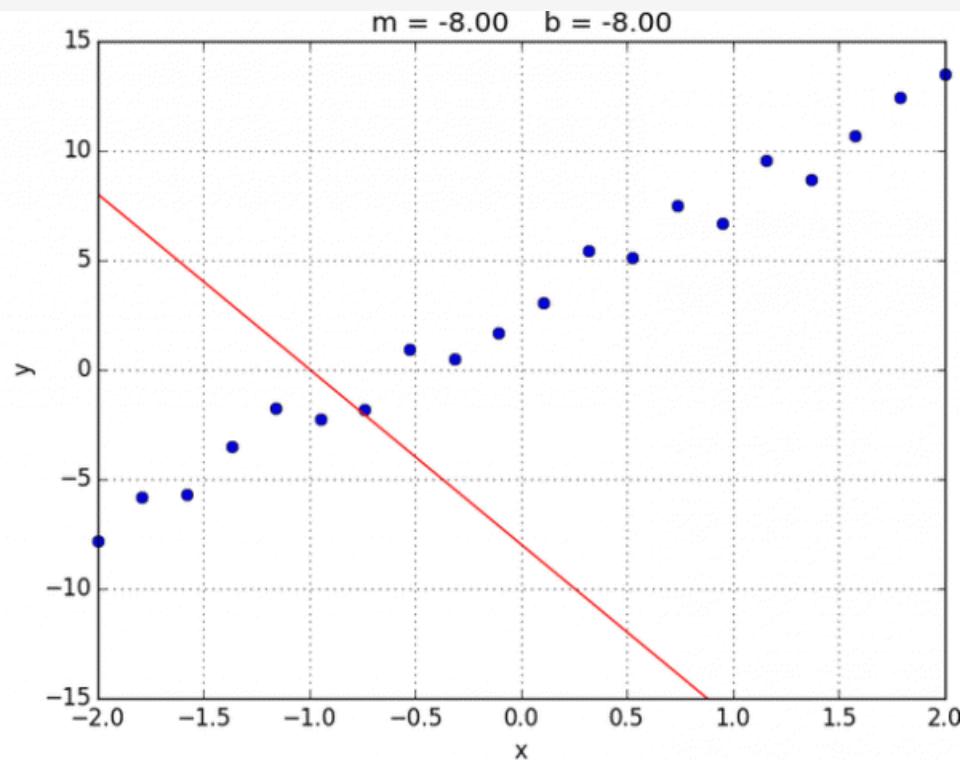
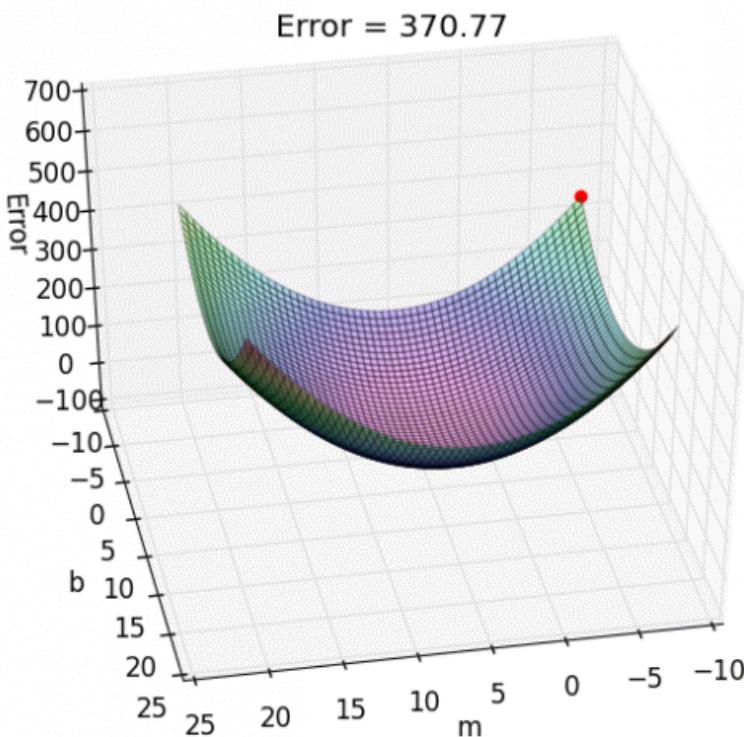
Tradeoff between
speed and precision



Reminders: Gradient Descent optimization

Illustration of Gradient Descent

- Let's consider a linear regression with $f(x) = m * x + b$
- With $m, b \equiv \theta$ and $L(\theta) = MSE$



The Challenges motivating DL



Challenges motivating Deep Learning

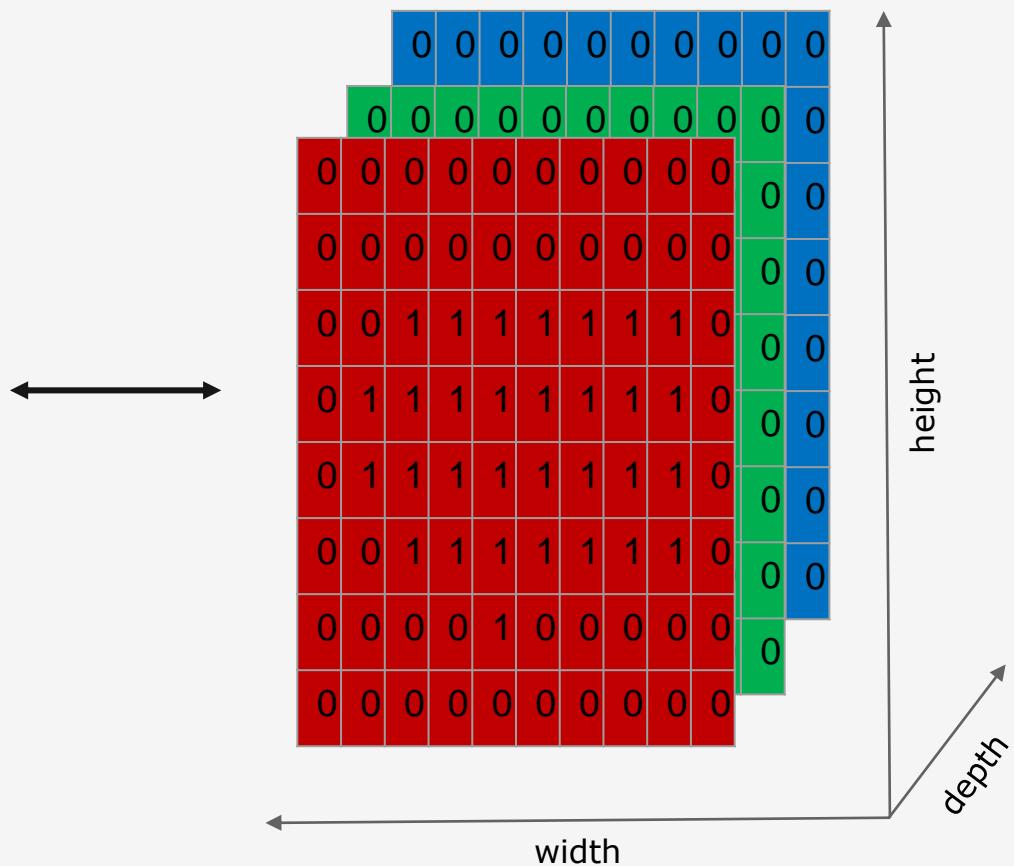
- ML works well in wide variety of problems
- Unfortunately, they haven't succeed in solving core AI problems (e.g. image recognition, speech recognition, etc ...)

But Why ?

Challenges motivating Deep Learning

Failure of ML: example of image recognition

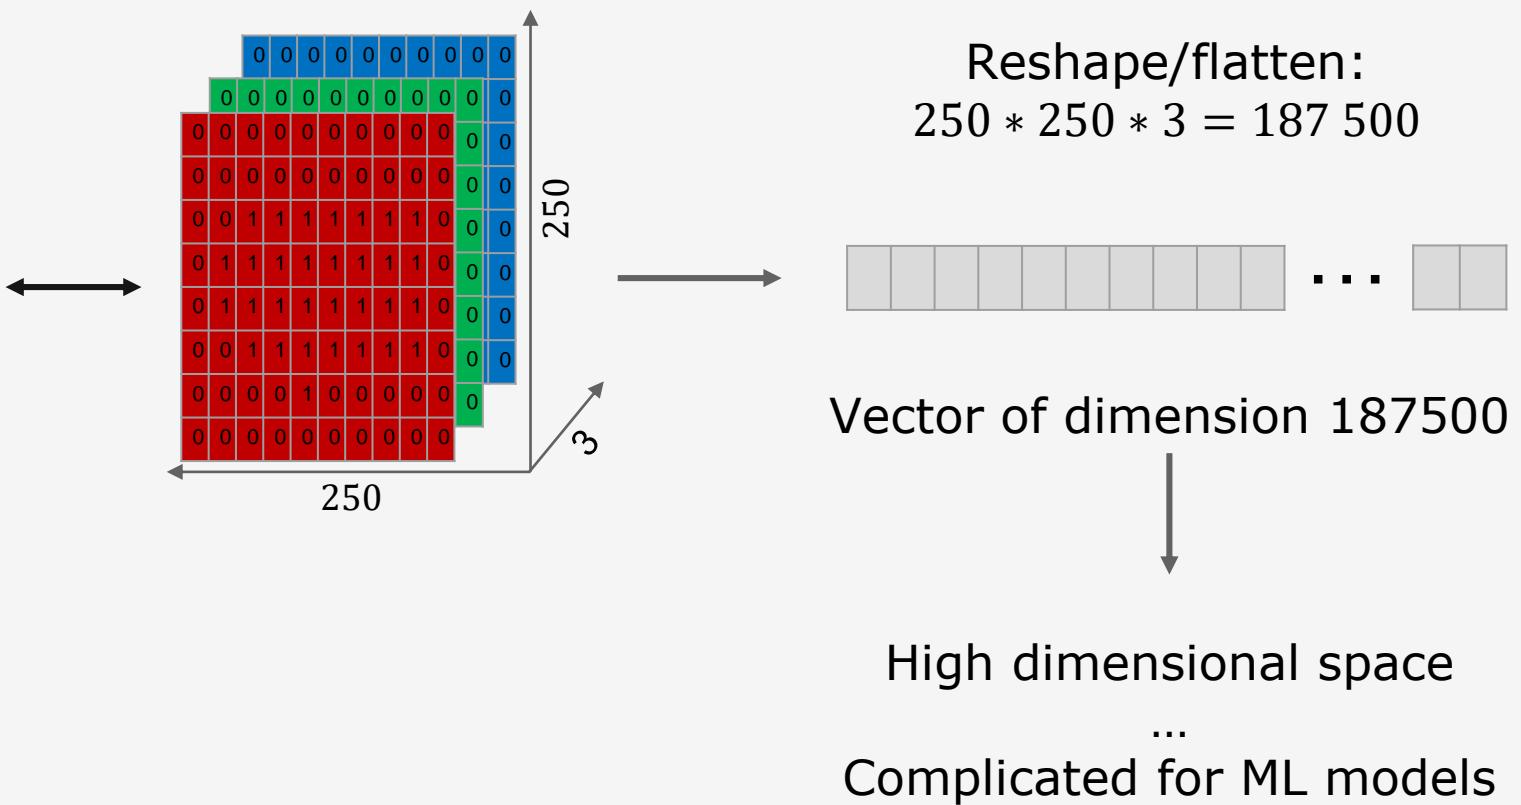
Image = RGB matrix



Challenges motivating Deep Learning

Failure of ML: example of image recognition

With classical ML model



Challenges motivating Deep Learning

Challenge 1: The curse of dimensionality

- Many ML problems' complexity increase with the number of dimensions in the data
- The number of possible set of configurations increases exponentially with the number of variables
- When can we qualify the issue of high-dimensional ?

number of variables(x) \gg number of training examples (n)

Challenges motivating Deep Learning

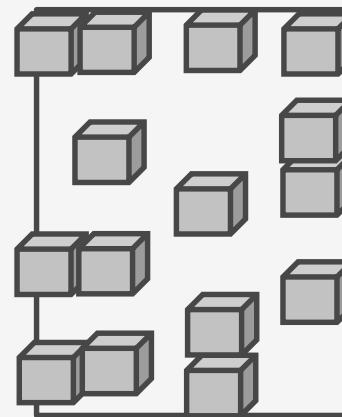
Challenge 1: The curse of dimensionality

1 dim



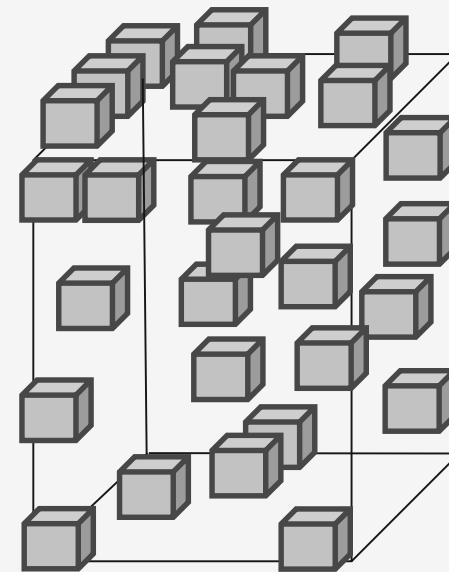
4
regions of
interest

2 dims



$4 * 4 = 16$
regions of
interest

3 dims



$4 * 4 * 4 = 64$
regions of
interest

Challenges motivating Deep Learning

Challenge 1: The curse of dimensionality

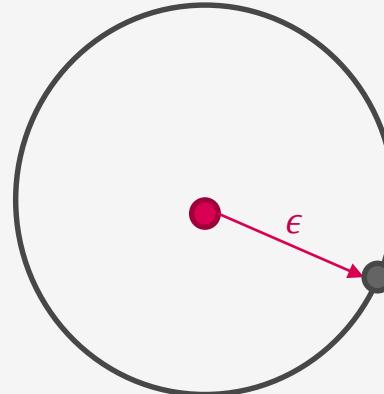
- Traditional ML approach insufficient for these spaces
- **DL challenge:** generalize well in high-dimensional spaces while keeping the computational cost low

Challenges motivating Deep Learning

Challenge 2: Local constancy and smoothness regularization

- ML algorithms must be guided by *prior beliefs* about the kind of function to learn
- **Local constancy prior** : “the function we learn should not change very much within a small region”:

$$f^*(x) \approx f^*(x + \epsilon)$$



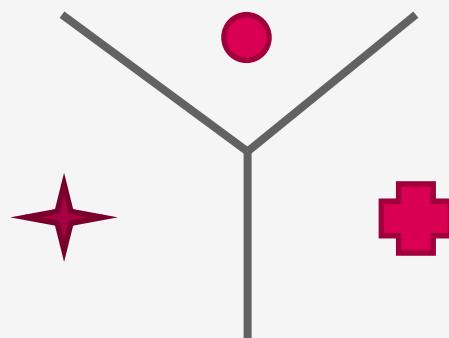
- Intuition: if we know a good answer for an input x , the answer is also good in the small neighborhood of x

Challenges motivating Deep Learning

Challenge 2: Local constancy and smoothness regularization

- **Limitation:** “to distinguish $O(k)$ regions in the input space, we need $O(k)$ examples”
- *Nearest neighbors, local kernel and tree based* methods suffer from the limitations of this prior

Example of knn



Challenges motivating Deep Learning

Challenge 2: Local constancy and smoothness regularization

- Challenge
 - Can we represent a complicated function efficiently?
 - Can it generalize well to new inputs ?

- DL allows to represent

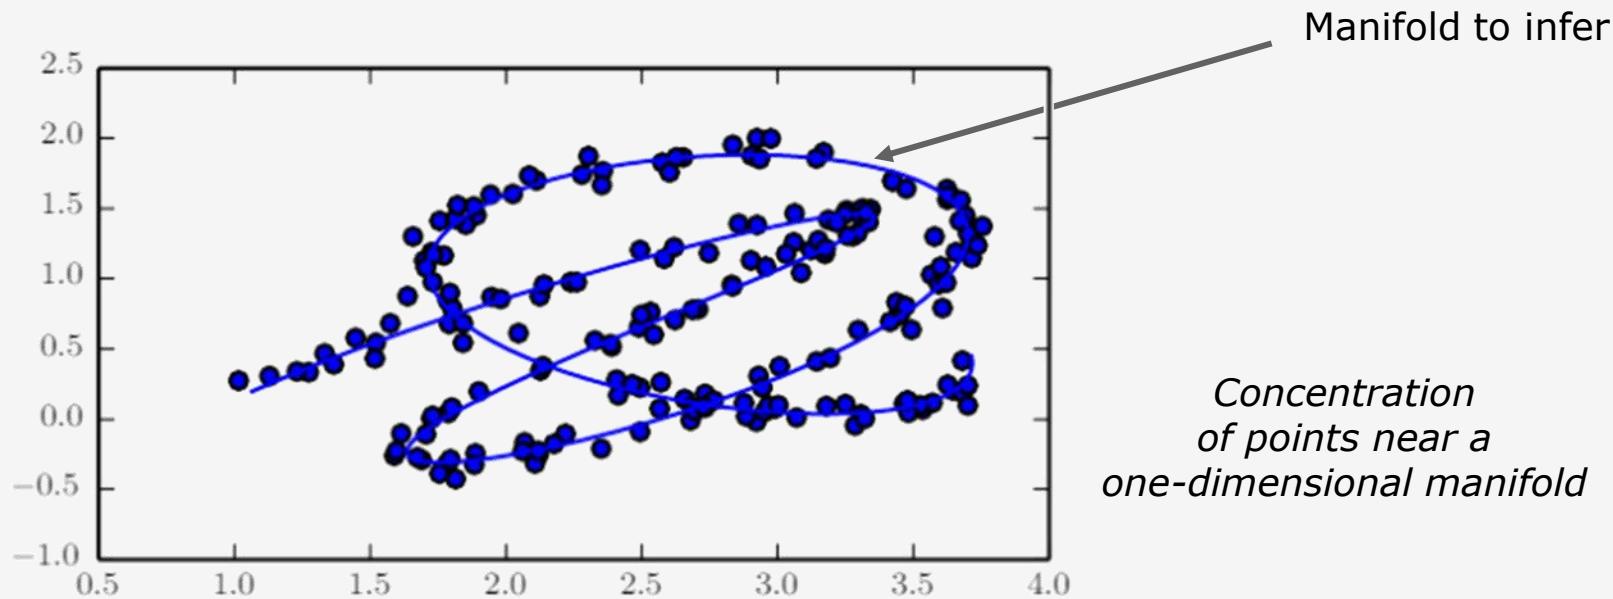
$$O(2^k) \text{ regions with } O(k) \text{ examples}$$

as it can introduce additional assumptions about the underlying data-generating distribution

Challenges motivating Deep Learning

Challenge 3: Manifold learning

- **Manifold** is a connected region: a set of points associated with a neighborhood around each point. It is an Euclidean space.
- It is possible to move from neighborhood to another one on the manifold



Challenges motivating Deep Learning

Challenge 3: Manifold learning

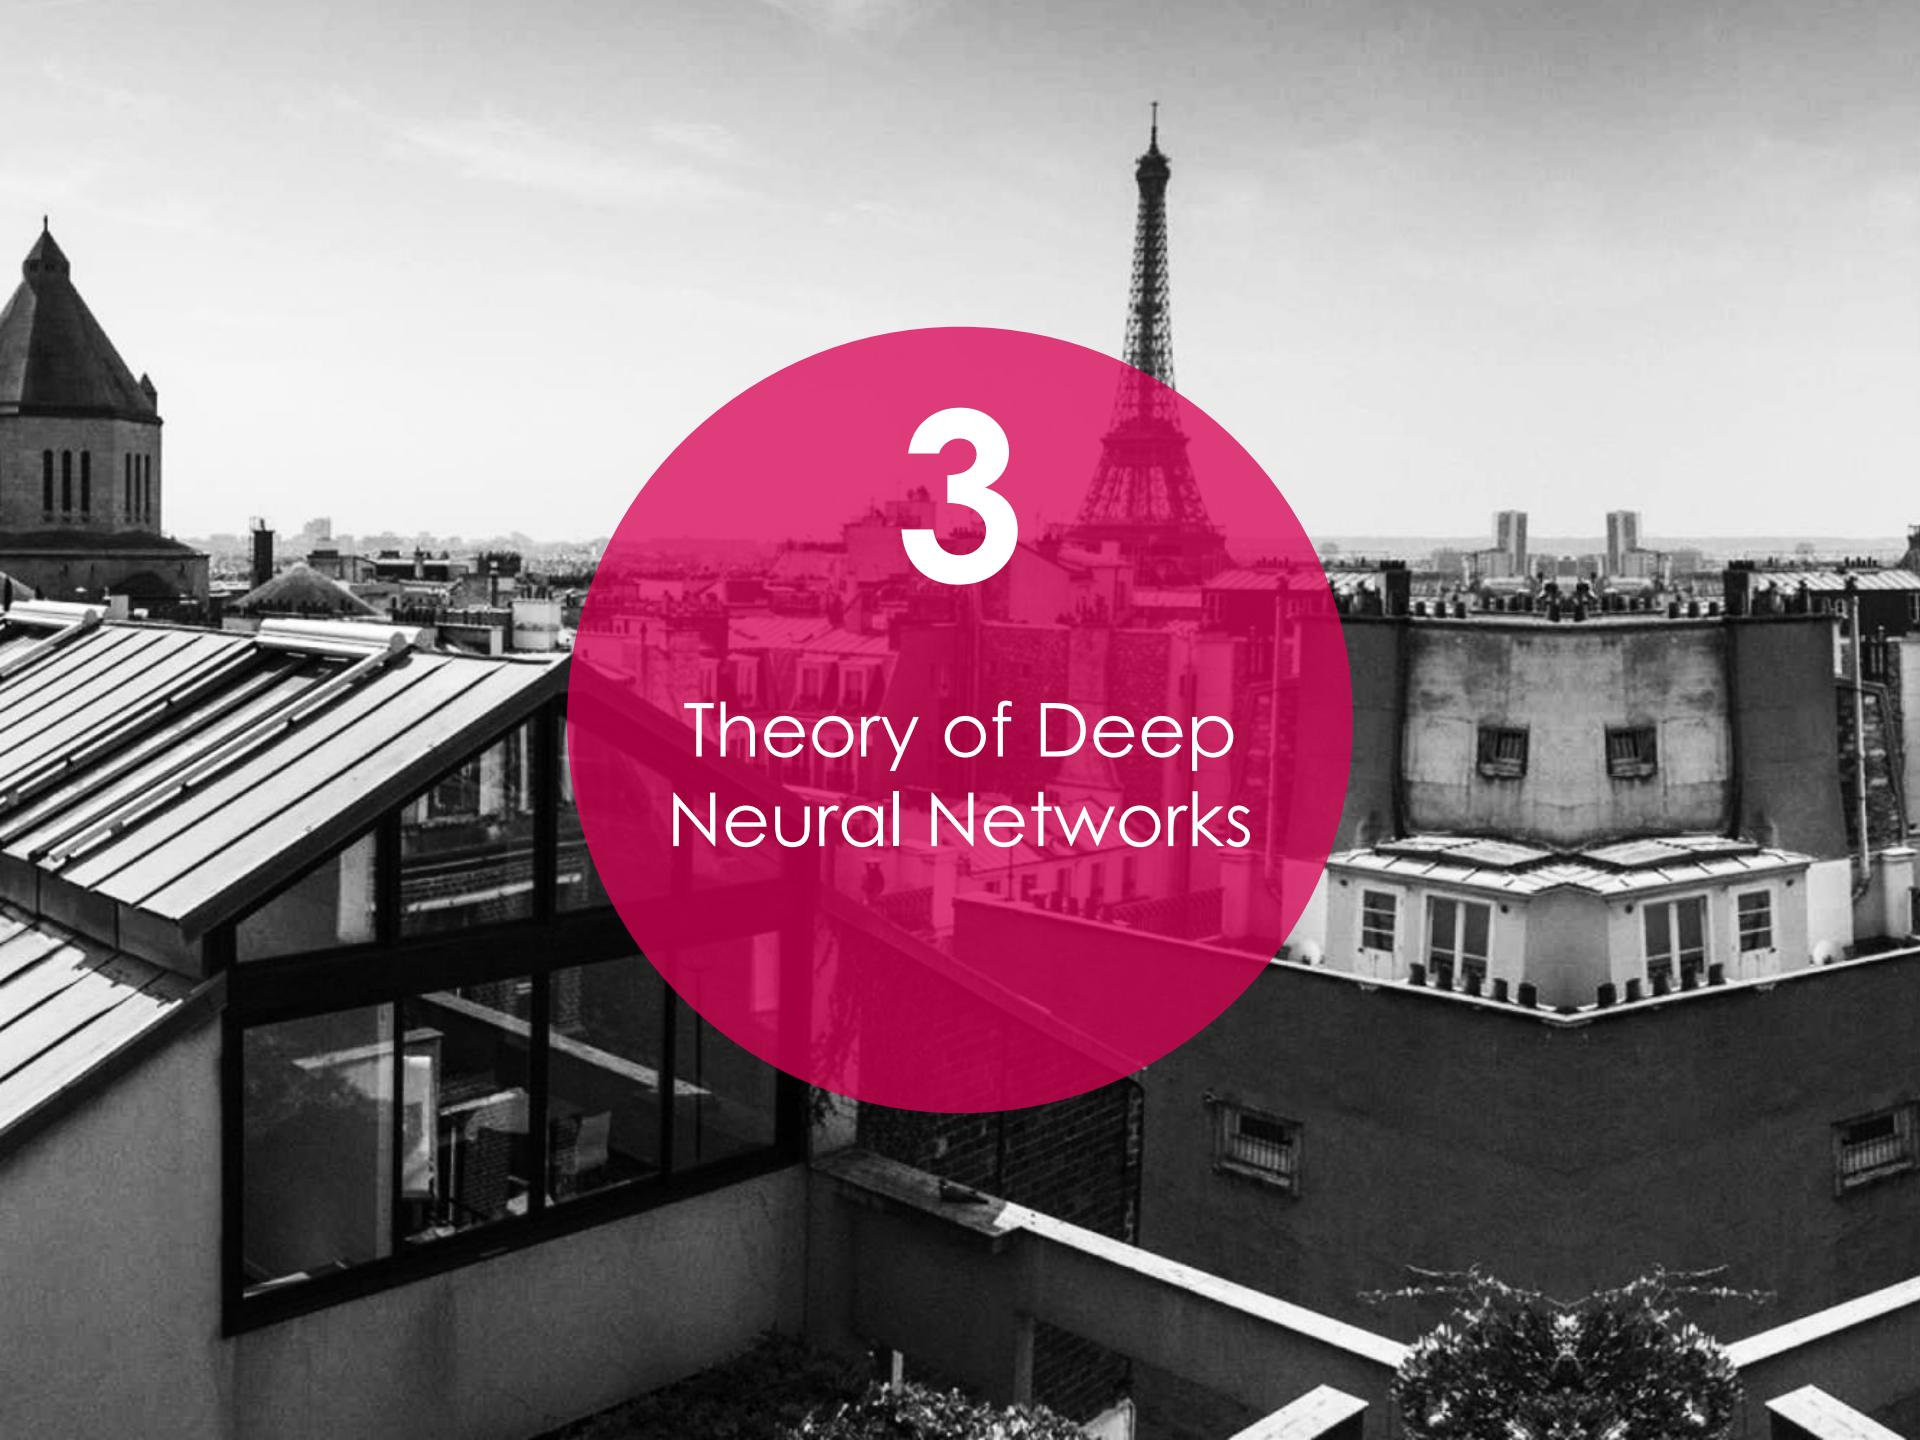
- ML seem hopeless if want to learn interesting variations across all R^n
- Manifold learning algorithms consider that
 - R^n inputs are invalid
 - Interesting variations occur only along a collection of manifolds
 - Along the directions of the manifold
 - When passing from one manifold to another
- Necessary assumption for processing images, audio or texts

Challenges motivating Deep Learning

To summarize so far ...

- Deep learning is motivated by its ability to
 - Learn in high-dimensional spaces
 - Learn very complicated functions efficiently and still be able to generalize well
 - Allow manifold learning
- Solve core AI issues

- We've seen:
 - Basic reminders of ML
 - Reasons to interest in DL



3

Theory of Deep Neural Networks

Deep Forward Neural Networks



Deep Forward Neural Networks

Perceptron

- Simplest neuron possible
- Introduced by *Rosenblatt* in the late 50's
- « *A device that makes decisions by weighting up evidence* »
- Takes input X , transforms it and fires a binary output (0 or 1)

Deep Forward Neural Networks

Perceptron

What does it mean in maths ?

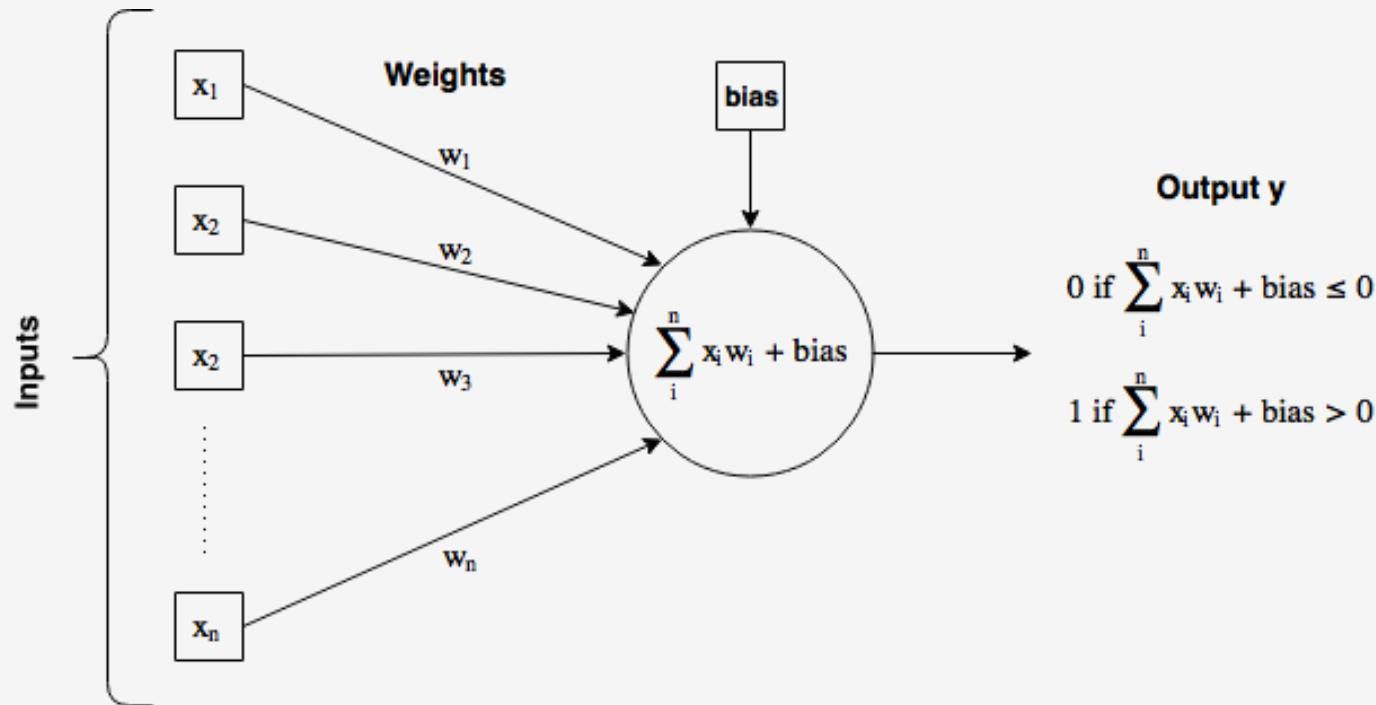
- Let's consider:
 - Inputs $X (x_1, x_2, x_3, \dots, x_n)$
 - Binary output $f(x)$
 - Set of weights $W(w_1, w_2, w_3, \dots, w_n)$
 - Bias b
- The perceptron can be expressed:

$$f(x) = \begin{cases} 0 & \text{if } \sum_{i=1}^n x_i w_i + b \leq 0 \\ 1 & \text{if } \sum_{i=1}^n x_i w_i + b > 0 \end{cases}$$

Deep Forward Neural Networks

Perceptron

Flow of information in a perceptron with weights W and bias b



Deep Forward Neural Networks

Perceptron

- A perceptron fires only binary outputs
- How can we get continuous outputs ?

Deep Forward Neural Networks

Neuron with activation

- **Activation** term σ
- such as

$$f(x) = \sigma \left(\sum_{i=1}^n x_i w_i + b \right)$$

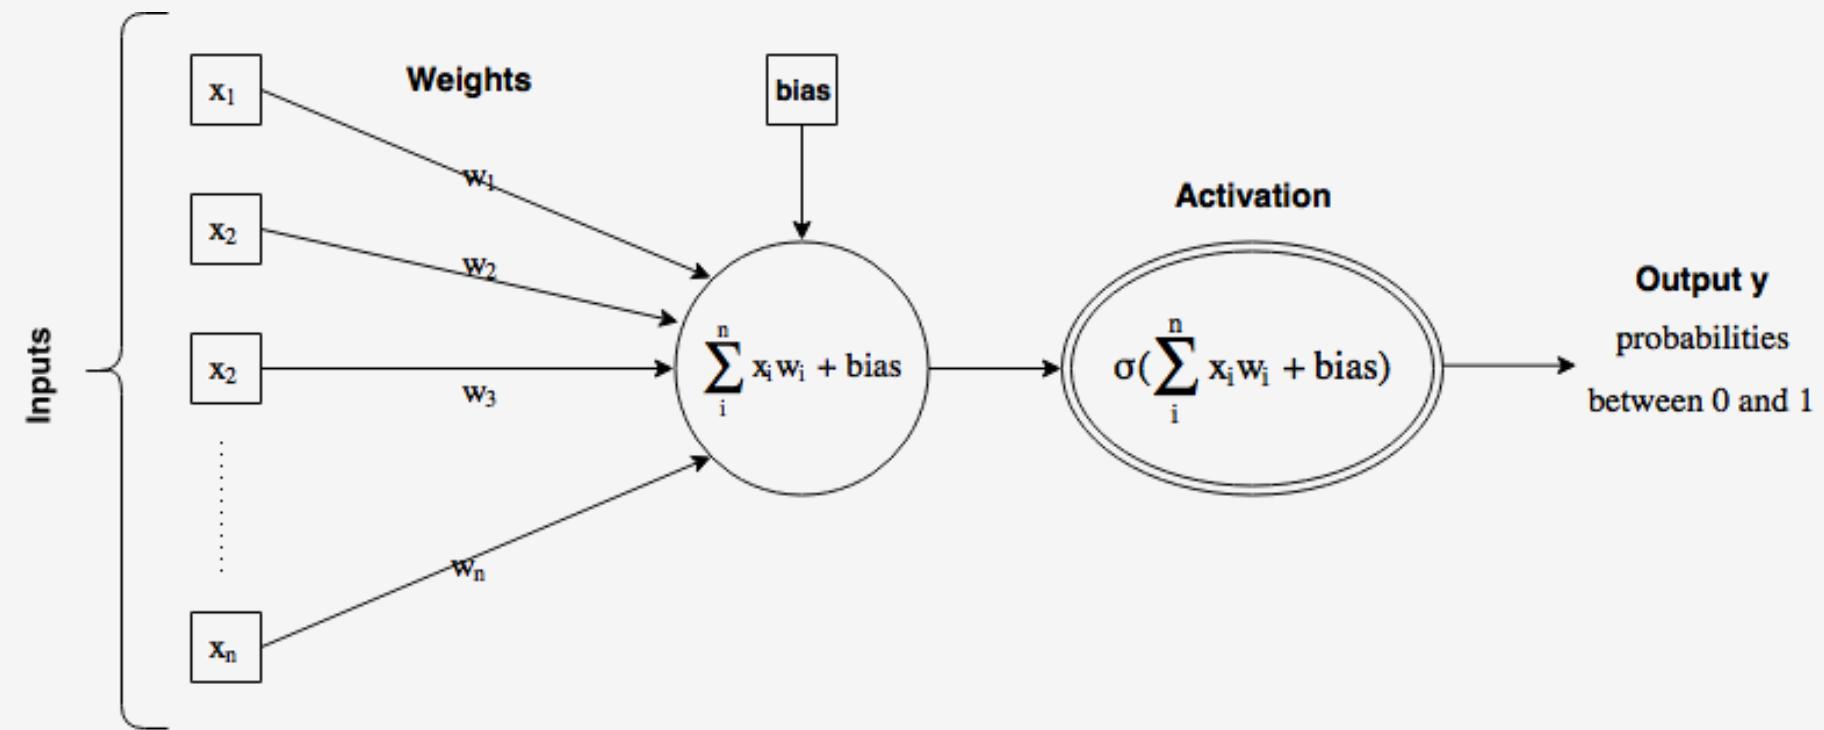
- Sigmoid activation function (i.e. logit) \rightarrow famous

$$\sigma \left(\sum_{i=1}^n x_i w_i + b \right) = \frac{1}{1 + e^{-\sum_{i=1}^n x_i w_i + b}}$$

Deep Forward Neural Networks

Neuron with activation

Flow of information in a neuron with activation



Deep Forward Neural Networks

Neuron with activation

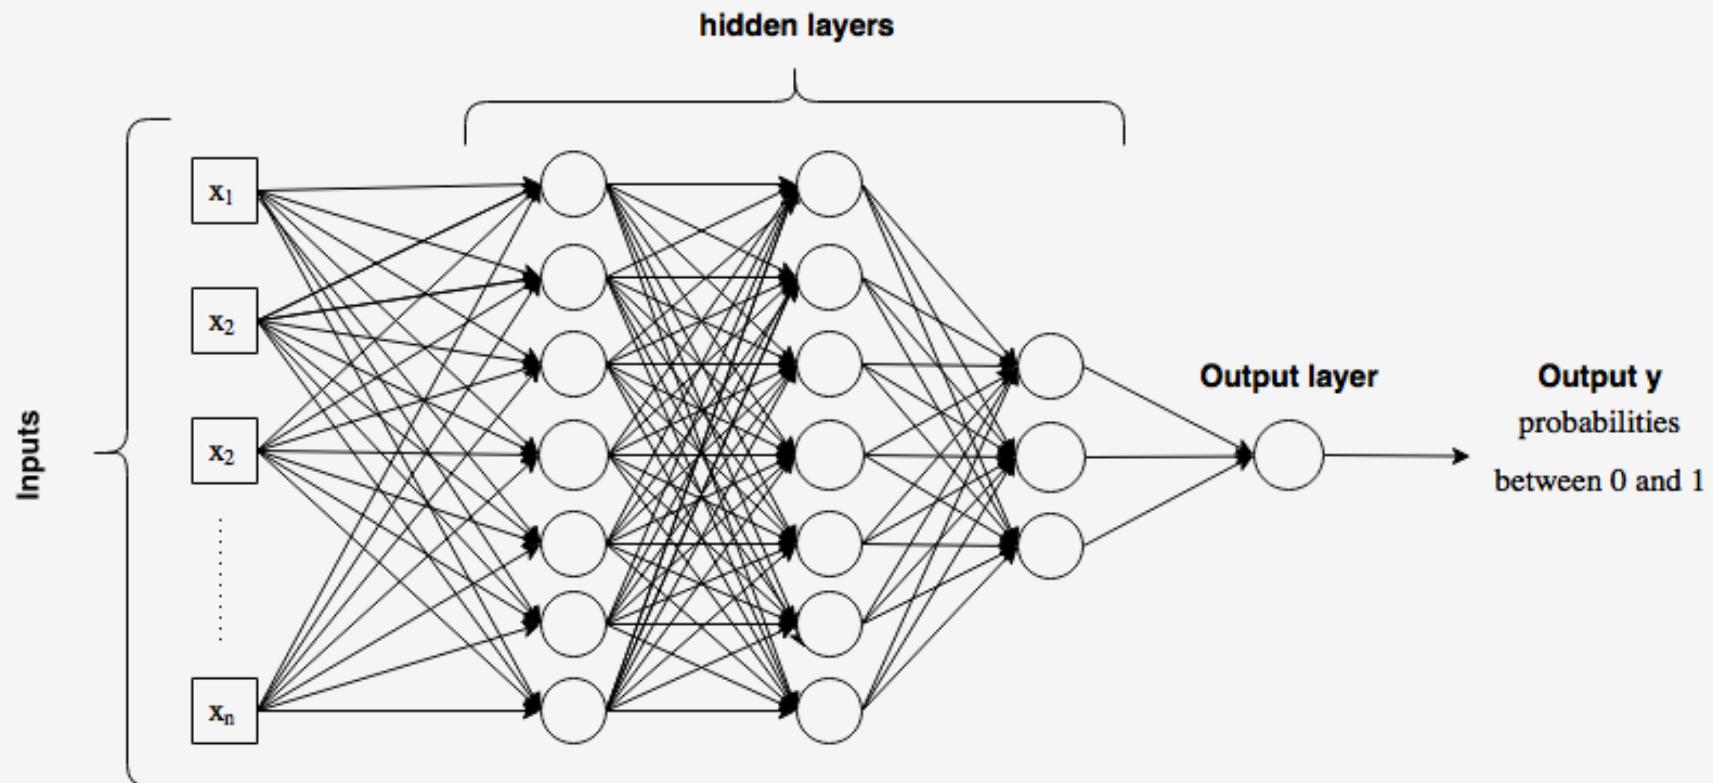
A single neuron might not be the strongest model of decision making

...

However, a set of neurons constituting together a sufficiently complex network could make sophisticated decisions

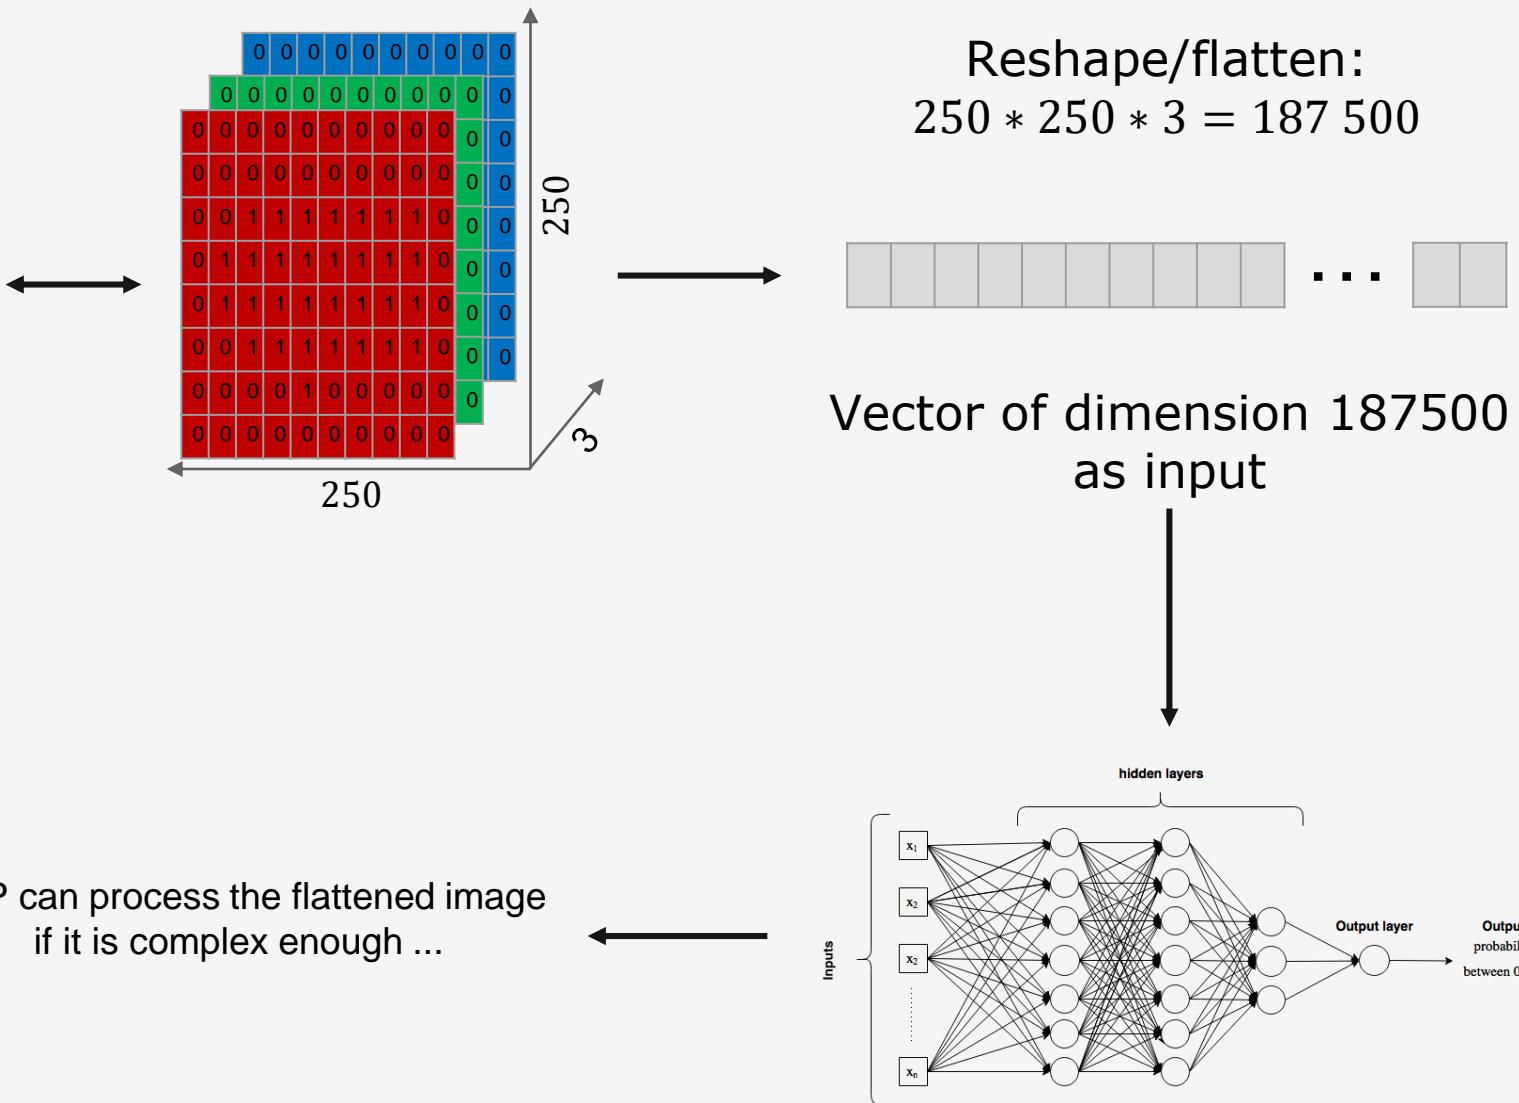
Deep Forward Neural Networks

Multilayer Perceptron



Deep Forward Neural Networks

Multilayer Perceptron



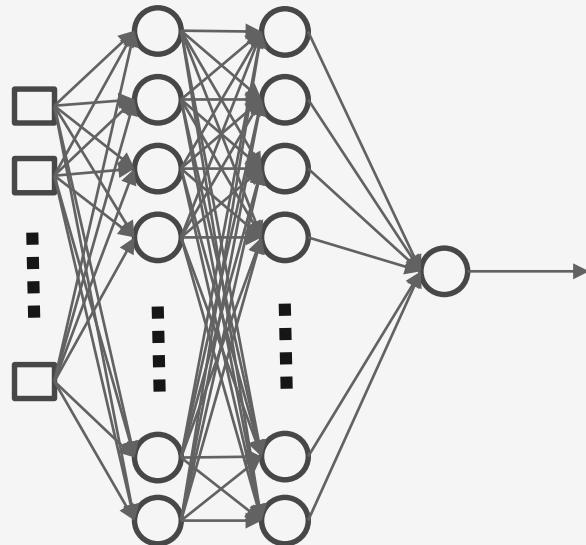
Deep Forward Neural Networks

Multilayer Perceptron - design

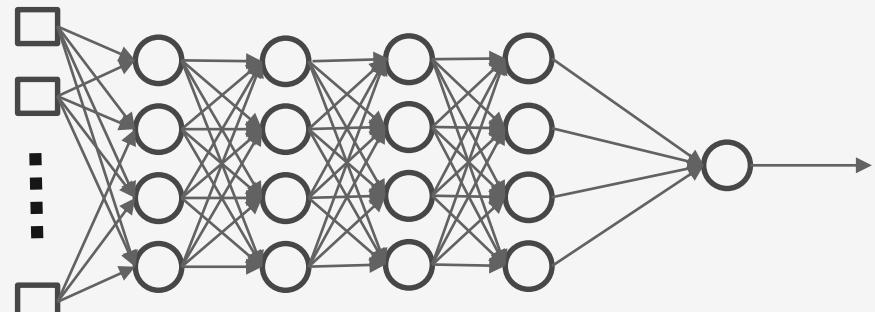
If it is complex enough ?!

...
How to make it 'more' complex ?

More neurons
per layer



More layers



Deep Forward Neural Networks

Multilayer Perceptron – complexity, overfitting and underfitting

Be careful !!

If **too complex**

...

if too many layers,
too many neurons



Risk of **overfitting**

If **not sufficiently complex**

...

too few layers,
too few neurons



Risk of **underfitting**

Deep Forward Neural Networks

Universal approximation theorem

Feedforward networks (such as MLP) allow the learning of highly non-linear function.

"According to the **universal approximation theorem** (Hornik et al. (1989)) any feedforward network with a linear output layer and at least one hidden layer with any "squashing" activation function (such as the logistic sigmoid activation function) **can approximate any Borel measurable function** from one finite-dimensional space to another with any desired non-zero amount of error, **provided that the network is given enough hidden units**"

Deep Forward Neural Networks

Universal approximation theorem

- The learning process is critical as it can fail because:
 - Underfitting → unable to find the optimal parameters
 - Overfitting → the algorithm chooses the wrong function
- Can lead to bad accuracy and bad generalization

Deep Forward Neural Networks

Limitations of Multilayer Perceptron

The overparametrized model has 3 main inconvenient:

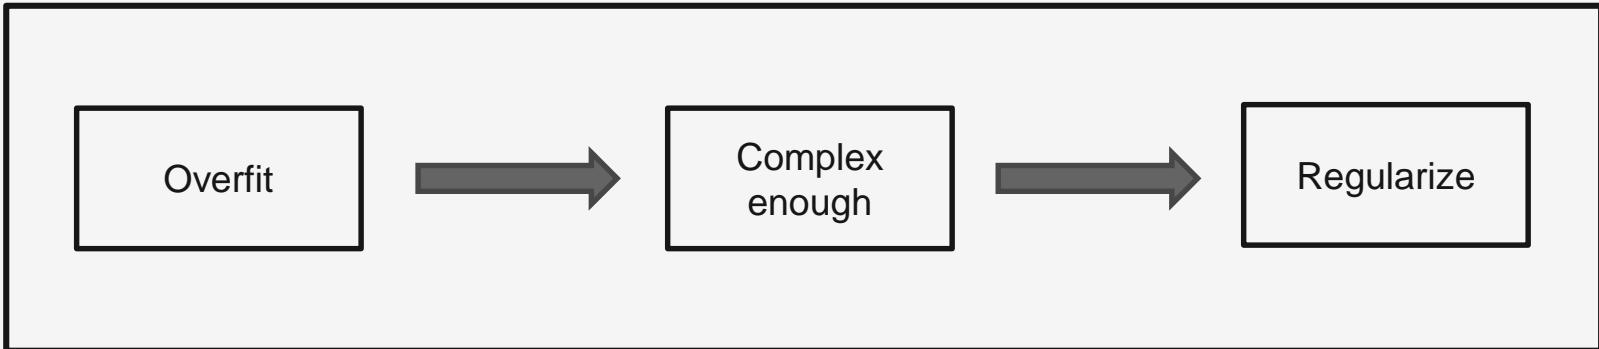
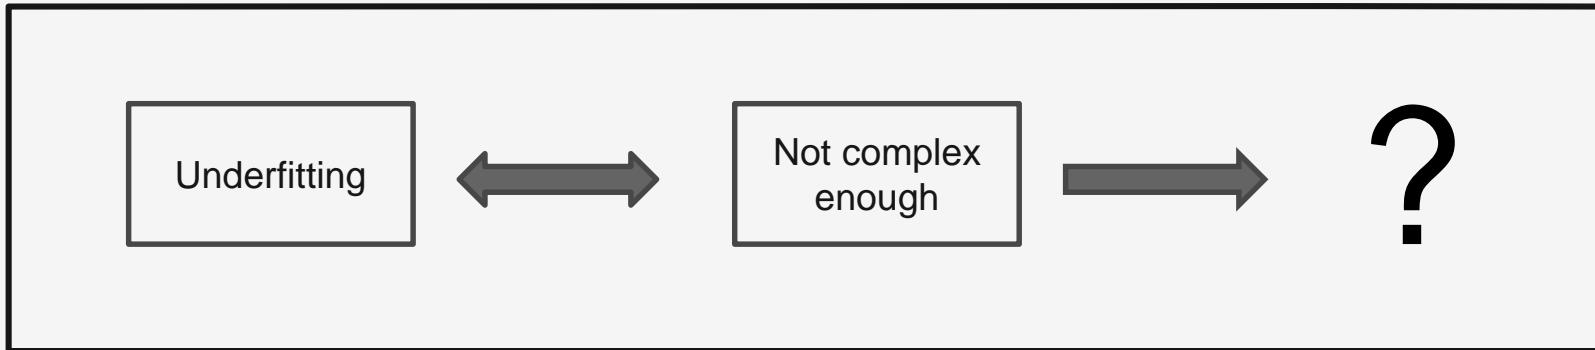
- **Non-convex and unstable loss function**
more than one local minima possible – random weight initialization can lead to different accuracy
- **Overfitting**
- **Sensitivity to feature scaling**
between 0 and 1

Regularization for DL



Regularization for DL

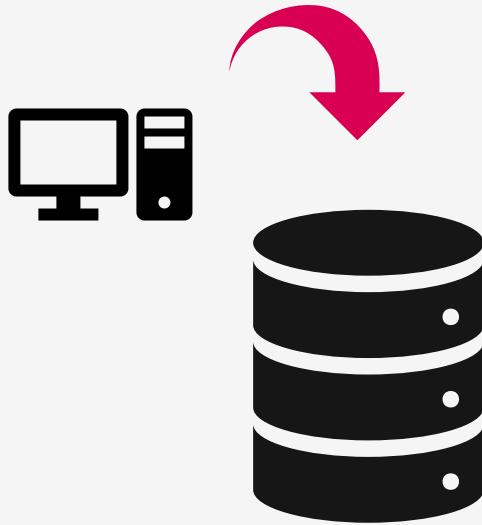
Should we first look to overfit ?



Regularization for DL

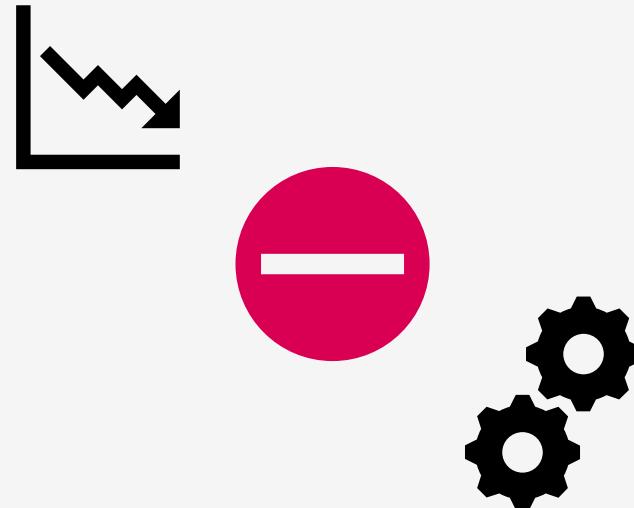
First easy options

Add more data



Not always possible ...

Stop the process when validation accuracy drops or stagnates



Regularization for DL

Data Augmentation

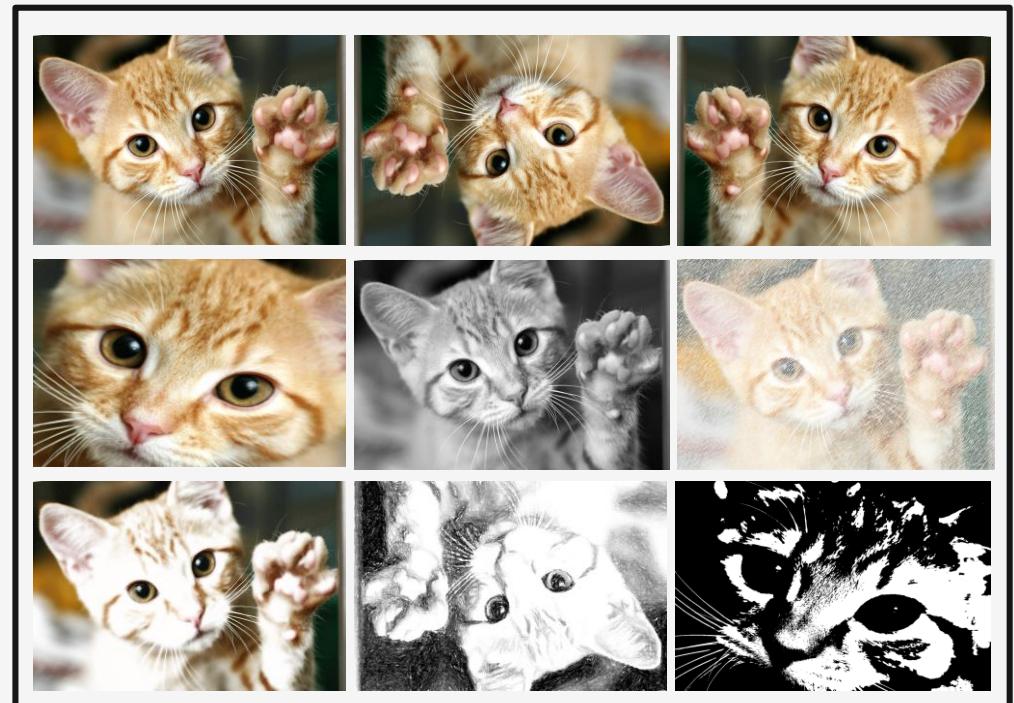
Blow artificially your training set !



Random
transformations



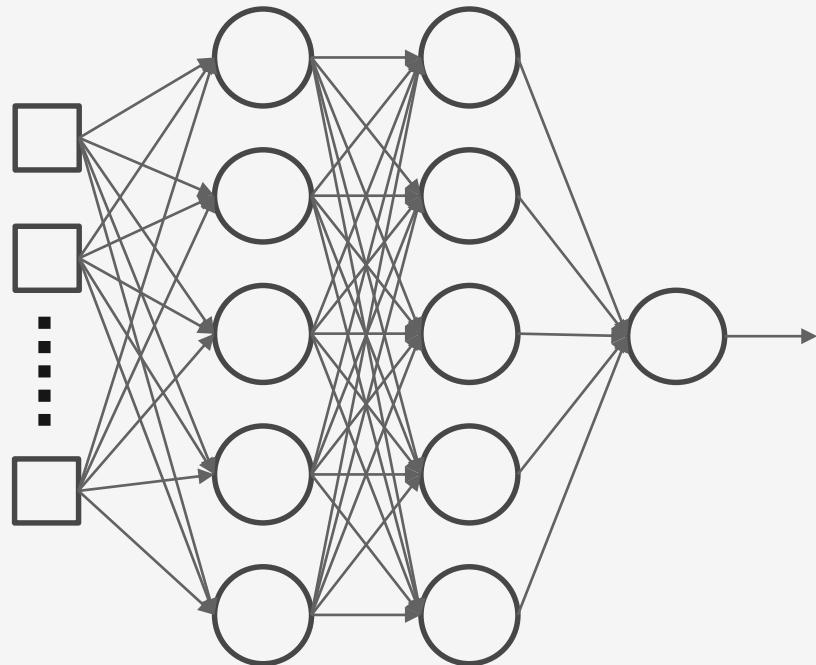
New set of images



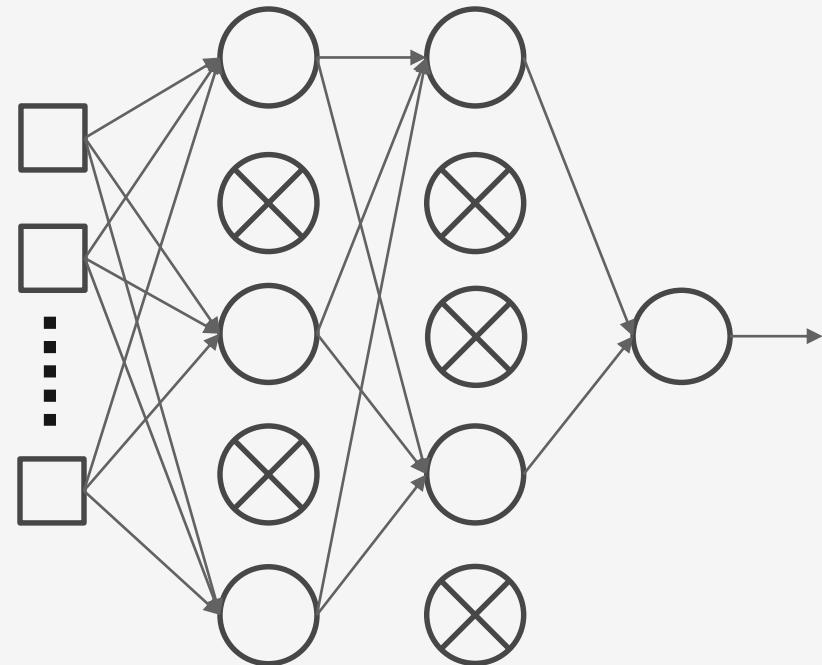
Regularization for DL

Dropout

Overfitting model



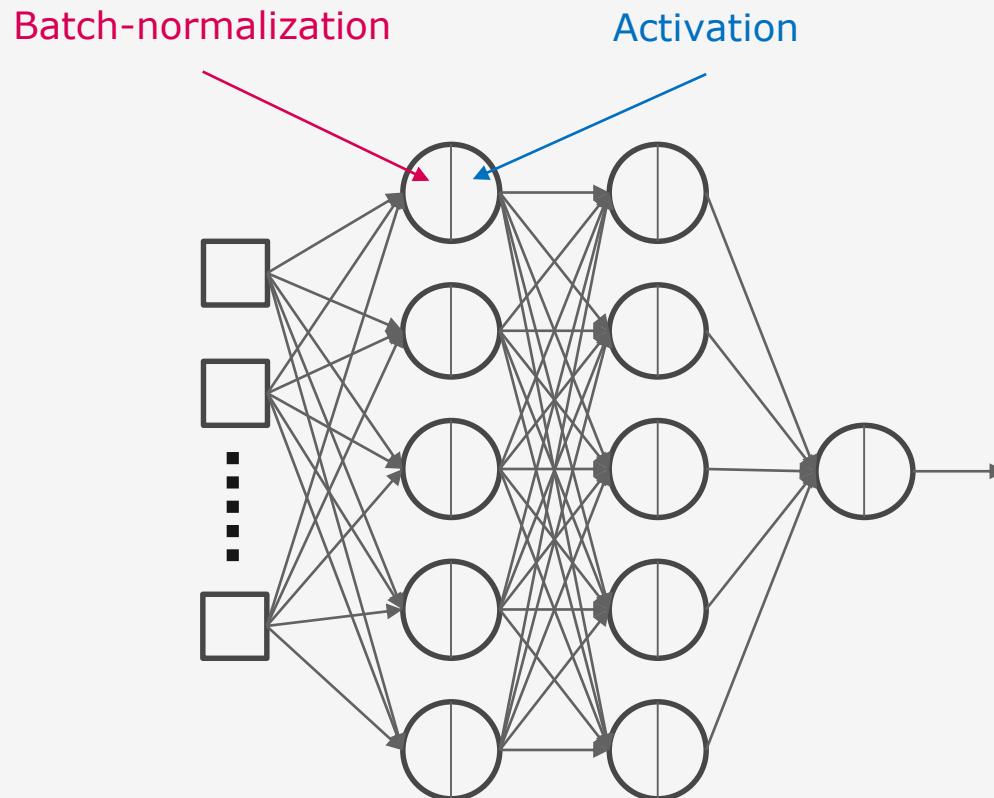
Dropout regularized model



It avoids units
too co-adapt too
strongly

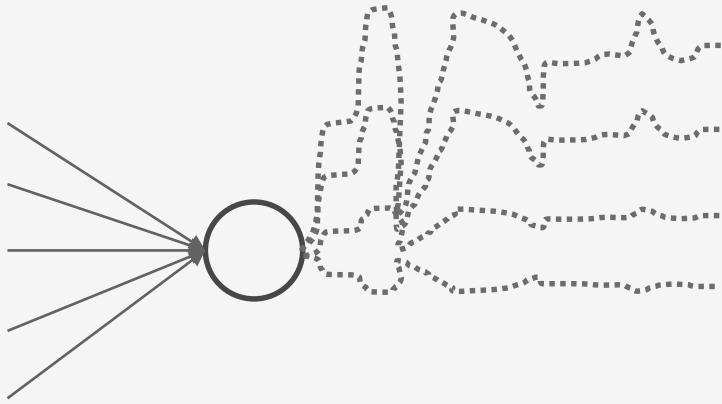
Regularization for DL

Batch-Normalization



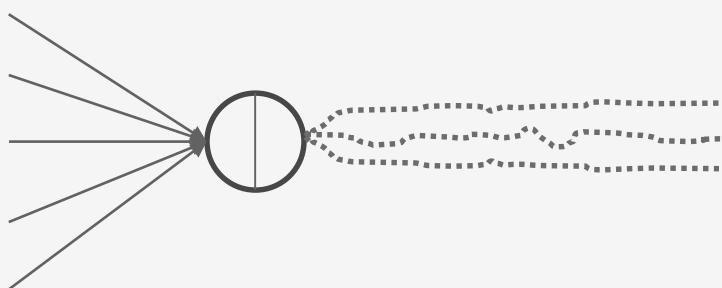
Regularization for DL

Batch-Normalization



Normal case without BN

- Internal covariance
 - Layer's need to adapt to inputs' distribution
- Learning requires lower $\lambda \rightarrow$ slower process



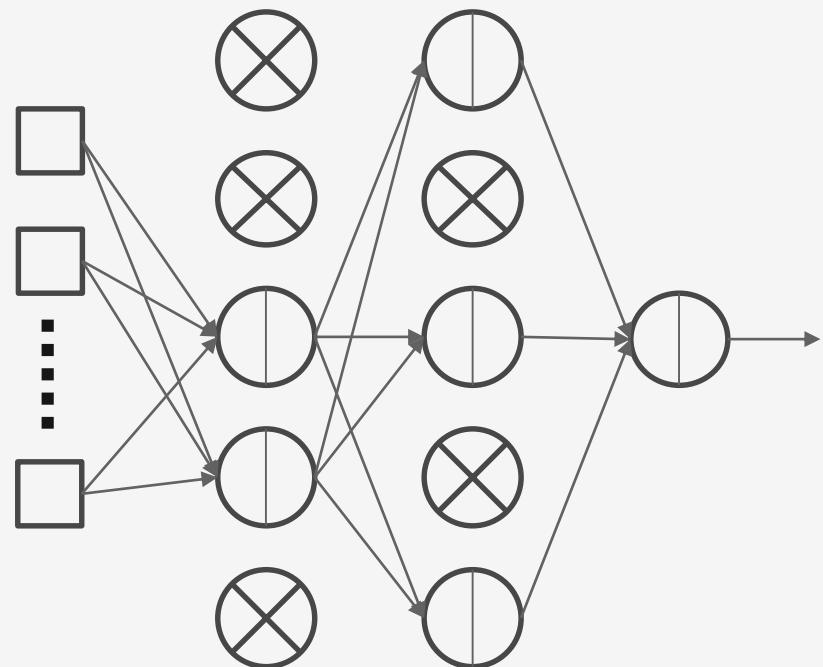
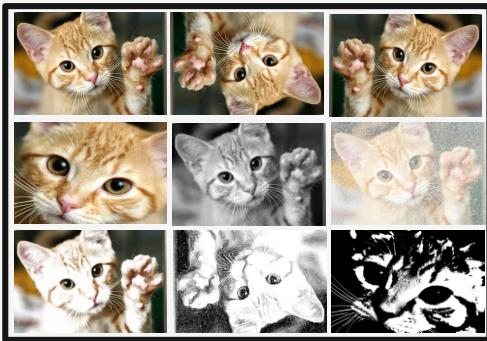
With BN

- Initially treats the internal covariance issue
- Allows higher $\lambda \rightarrow$ faster learning process
- It has also proven to significantly reduce overfitting

Regularization for DL

Combine all methods

Data Augmentation + BN + Dropout
It should work !

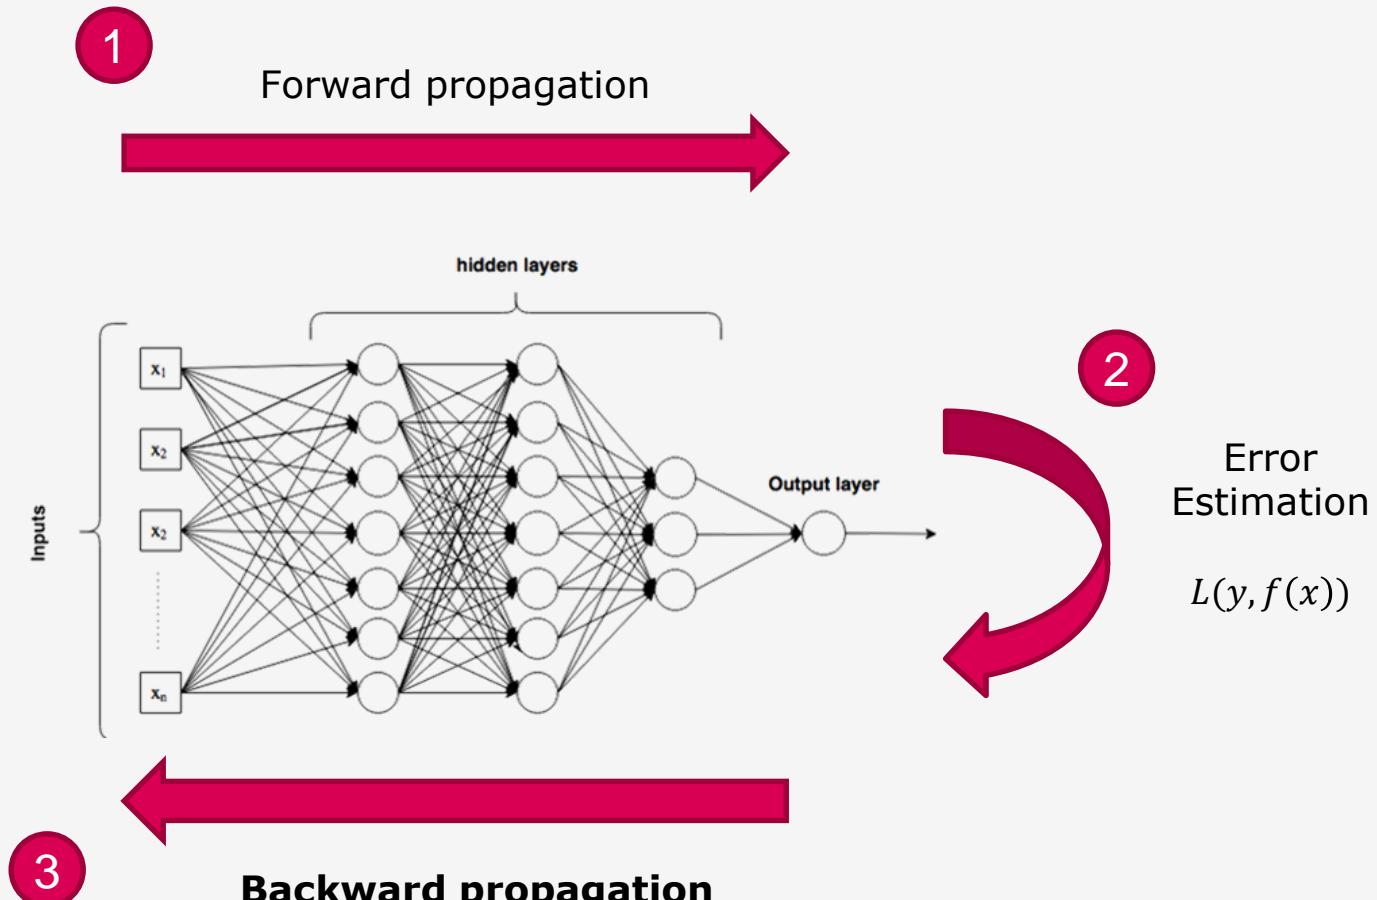


Training Deep Neural Networks



Training Deep Neural Networks

Back-propagation



- **Numerical computation of the gradient**
- **Weight actualization**

Training Deep Neural Networks

Stochastic Gradient Descent

- Recurrent issue in ML: large training set is good for generalization
- Gradient Descent with large datasets is computationally heavy



Stochastic Gradient Descent (SGD) with *mini-batches*

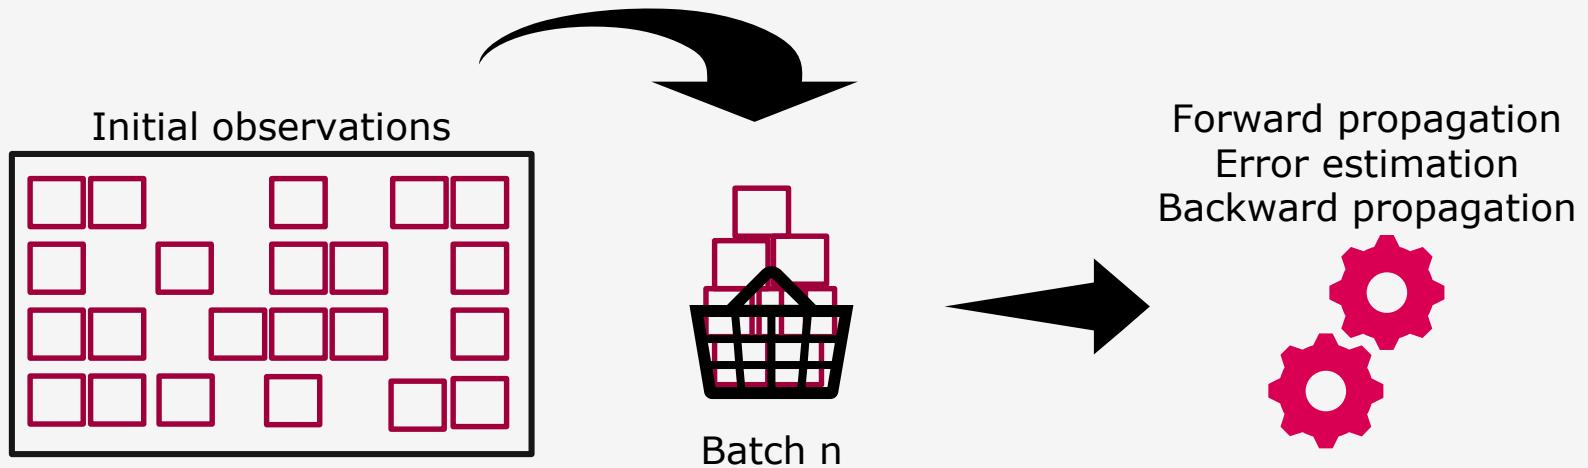
- Update of parameters for each mini-batch i

$$\theta_{t+1} = \theta_t - \lambda \nabla_{\theta} L(\theta_t; x_i; y_i)$$

Training Deep Neural Networks

Stochastic Gradient Descent

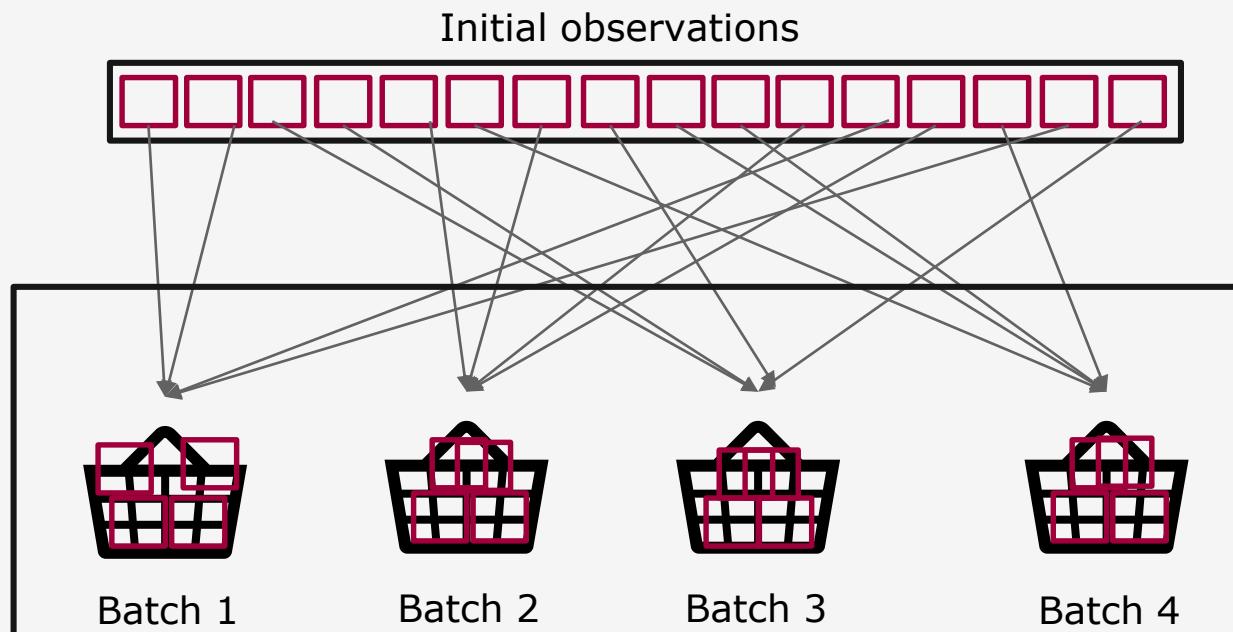
Batch/mini-batch learning



Training Deep Neural Networks

Stochastic Gradient Descent

Constituting an epoch



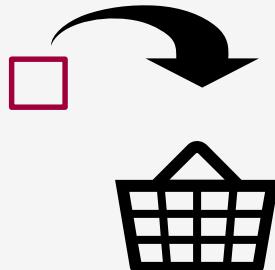
Training Deep Neural Networks

Stochastic Gradient Descent

Which batch size (bs) ?

If batch size too small

- Learning too small
- Regularizing effect expected
(generalization error best with bs 1)



If batch size too large

- Large bs → Better estimate of the gradient
- Faster learning
- Too large → degradation of the gen. accuracy

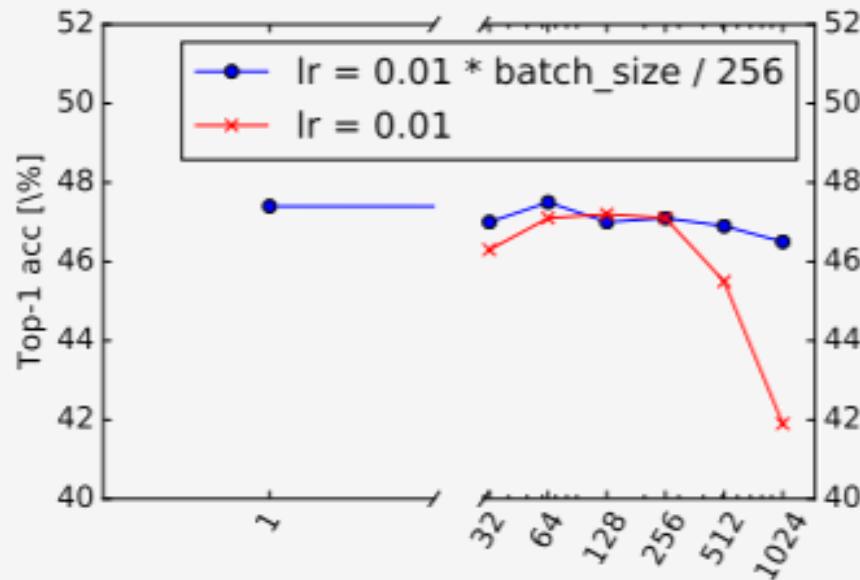


Training Deep Neural Networks

Stochastic Gradient Descent

Which batch size (bs) ?

Batch size and initial learning rate
impact to the accuracy

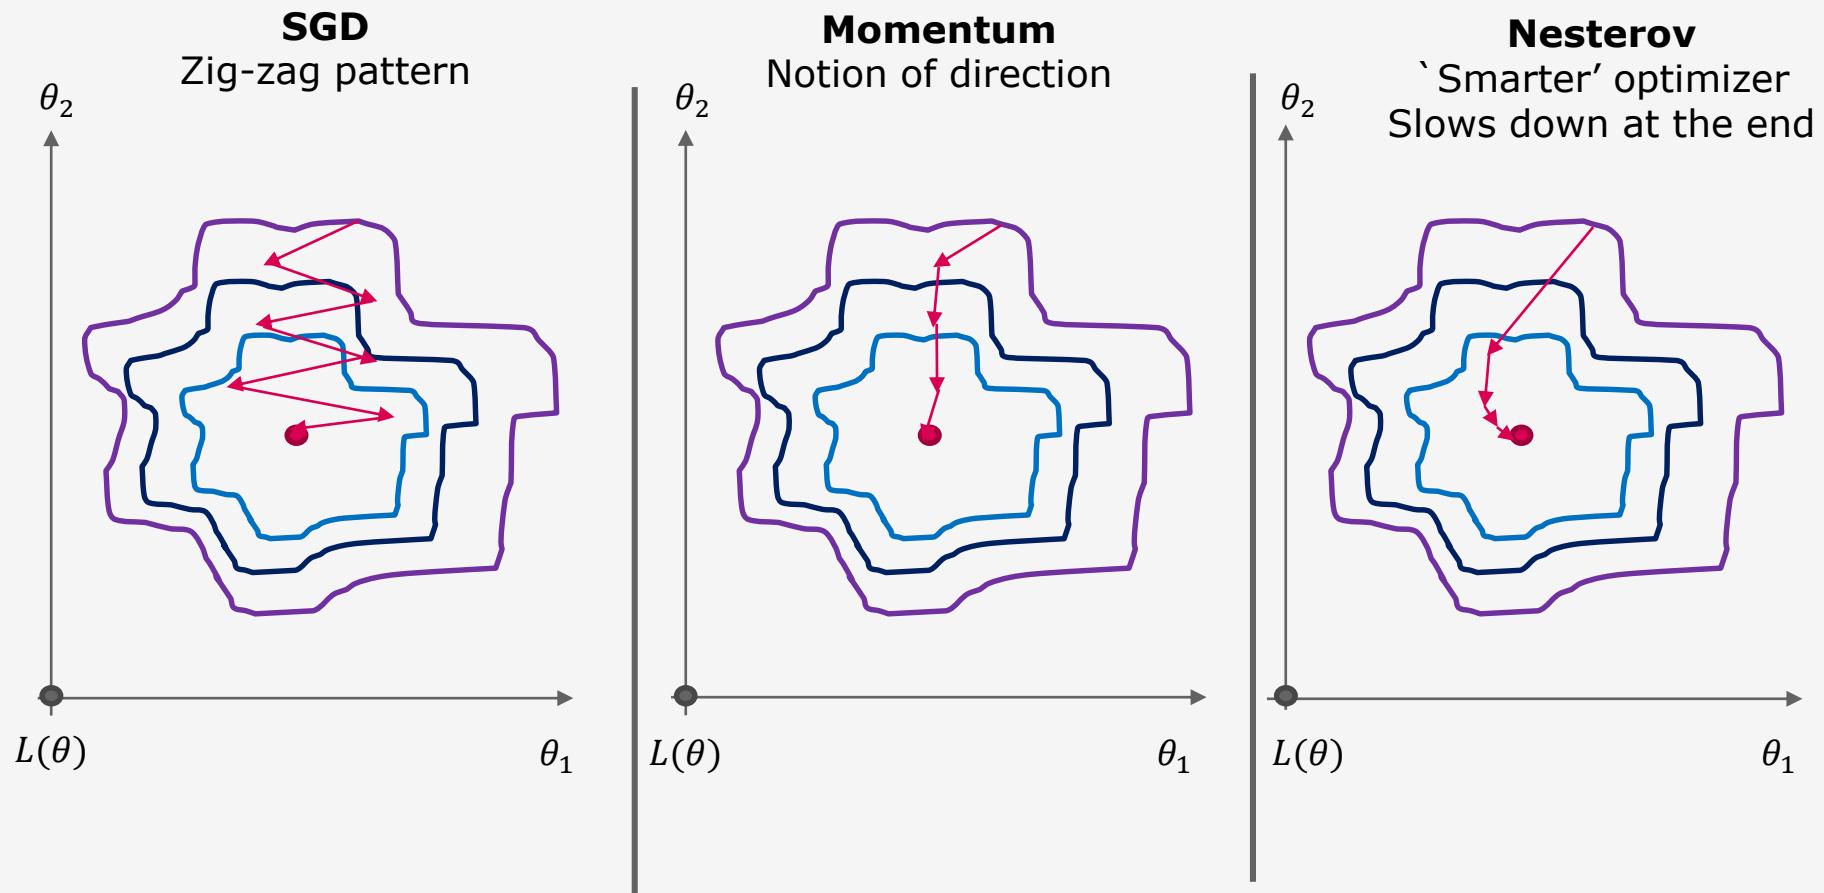


<https://arxiv.org/pdf/1606.02228.pdf>

Training Deep Neural Networks

Stochastic Gradient Descent

- Stochastic Gradient Descent is a “first order method”
- Such as **Nesterov** and **Momentum**



Training Deep Neural Networks

Adaptive optimizers

- *Second order moment* methods
- They minimize $J(\theta)$ not only by estimating the Gradient but also by estimating the **Hessian matrix** and using the **Newton's method**
- They **adapt the learning rate** λ at each step:

Adagrad, Adadelta, Adamax, RMSprop, Adam

.... For detailed description

Training Deep Neural Networks

Choosing the right optimization algorithm

Adaptive (second order moment)

Adapted for
Sparse input data

No tuning
Required for λ



Faster
convergence

First order moment

Slow
convergence

Can get stuck
In saddle paths

Too reliant on
Robust intialization



Adaptive = best option
In particular Adam

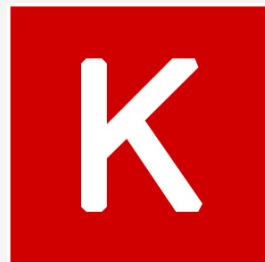
Introducing Keras

A little bit of practice !



Introducing Keras

Deep Learning specialized library



Wrapper designed for
fast implementation

On CPU or GPU



 TensorFlow™

The TensorFlow logo features a stylized orange 'T' composed of several small squares, followed by the word "TensorFlow" in a bold, sans-serif font, with a trademark symbol at the end.

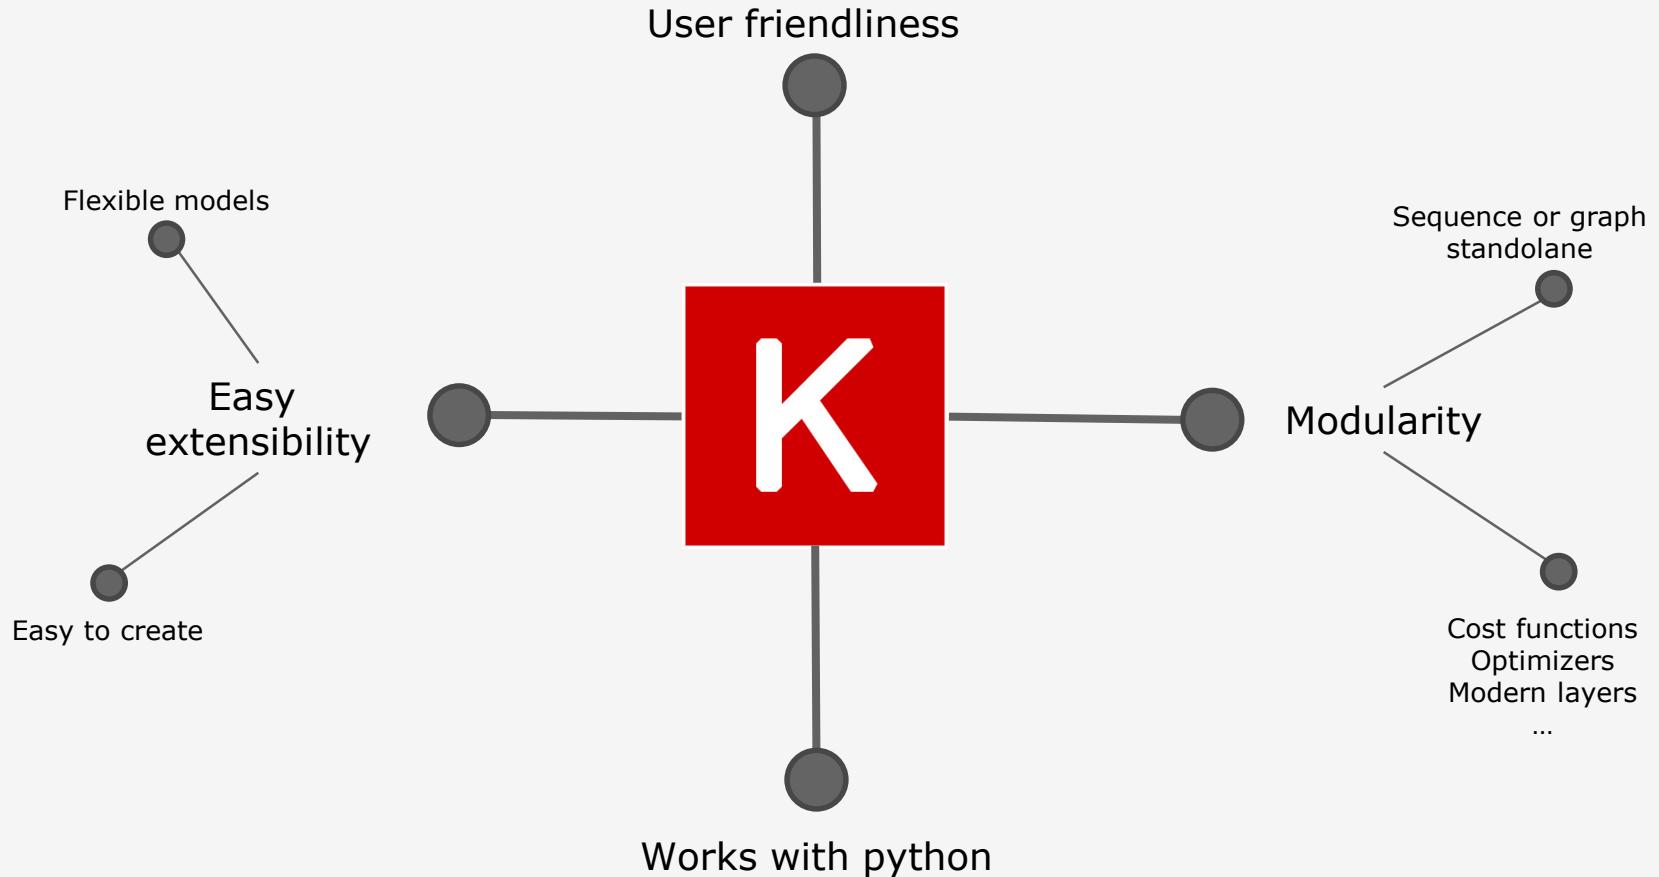
 Microsoft
CNTK

The Microsoft CNTK logo includes the Microsoft color scheme (blue, green, red, yellow squares) followed by the text "Microsoft" and "CNTK".

 theano

The theano logo is written in a large, blue, lowercase sans-serif font.

Introducing Keras



Introducing keras

Design, train and predict

1 – Declare a **Sequential** model

```
import keras  
from keras.models import Sequential  
  
model = sequential()
```

2 – Stack all layers

```
from keras.layers import Dense, Activation  
  
model.add(Dense(units=64,input_dim=100))  
model.add(Activation('Relu'))  
model.add(Dense(units=32))  
model.add(Activation('Relu'))  
model.add(Dense(units=10))  
model.add(Activation('softmax'))
```

3 – Configure its learning process

```
model.compile(loss='categorical_crossentropy',  
              optimizer='adam',  
              metrics=['accuracy'])
```

4.A – Fit your model – automatic batch

```
model.fit(x_train, y_train,  
          epochs=10,  
          batch_size=32)
```

4.B – Or fit on your own batch

```
model.train_on_batch(x_batch,y_batch)
```

5 – Evaluate on your validation set

```
loss_metrics = model.evaluate(X_val, y_val,  
                             batch_size=32)
```

6 – Make your own predictions

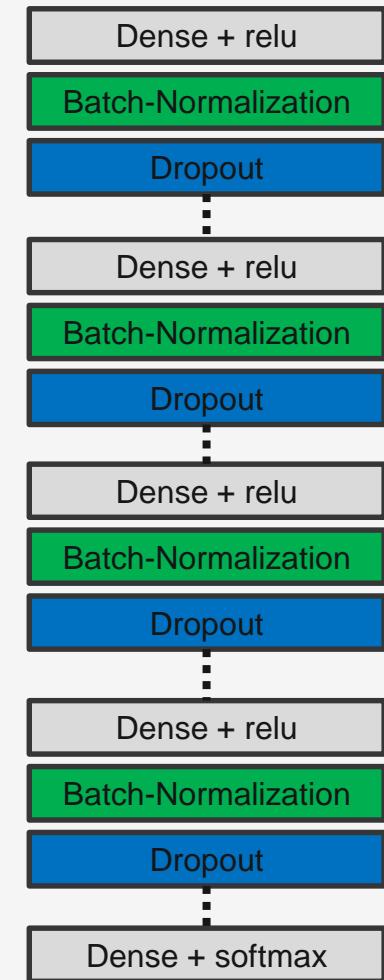
```
preds = model.predict(x_test,  
                      batch_size=32)
```

Introducing keras

Data augmentation, Dropout, Batch-Normalization

A - Dropout and Normalization

```
from keras.layers import Dense, Dropout  
from keras.layers.normalization import BatchNormalization  
  
model.add(Dense(units=128,input_dim=100, activation='Relu'))  
model.add(BatchNormalization())  
model.add(Dropout(0.25))  
  
model.add(Dense(units=128, activation='Relu'))  
model.add(BatchNormalization())  
model.add(Dropout(0.25))  
  
model.add(Dense(units=64, activation='Relu'))  
model.add(BatchNormalization())  
model.add(Dropout(0.5))  
  
model.add(Dense(units=64, activation='Relu'))  
model.add(BatchNormalization())  
model.add(Dropout(0.5))  
  
model.add(Dense(units=10, activation='softmax'))
```



Introducing keras

Data augmentation, Dropout, Batch-Normalization

B – Data Augmentation

1 – Declare all possible transformations

```
from keras.preprocessing.image import ImageDataGenerator  
  
train_datagen = ImageDataGenerator(rescale=1./255,  
                                   shear_range=0.2,  
                                   zoom_range=0.2,  
                                   horizontal_flip=True)  
  
test_datagen = ImageDataGenerator(rescale=1./255)
```

2 – Generate batches (flow from directory)

```
train_generator = train_datagen.flow_from_directory(  
    'data/train',  
    target_size=(150, 150),  
    batch_size=32,  
    class_mode='categorical')  
  
validation_generator = test_datagen.flow_from_directory(  
    'data/validation',  
    target_size=(150, 150),  
    batch_size=32,  
    class_mode='categorical')
```

3 – Fit on these generated batches

```
model.fit_generator( train_generator,  
                     steps_per_epoch=2000,  
                     epochs=50,  
                     validation_data=validation_generator,  
                     validation_steps=800)
```

4 – predict on the test set

```
test_generator = test_datagen.flow_from_directory(  
    'data/test',  
    target_size=(150, 150),  
    batch_size=32,  
    class_mode='categorical')  
  
model.predict_generator( test_generator,  
                       steps=total_size/batch_size)
```

Introducing keras

Design, train and predict

Quiet easy right ?!?



Just need to wait the
model to run !

....

It could take a while



Study case n°1

Design, train and predict

MNIST
Handwritten Digit
Recognizer

Using MLP



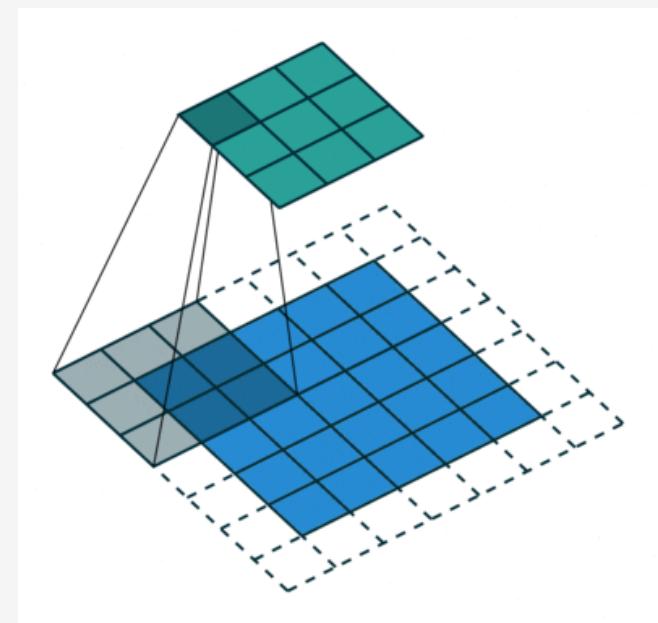
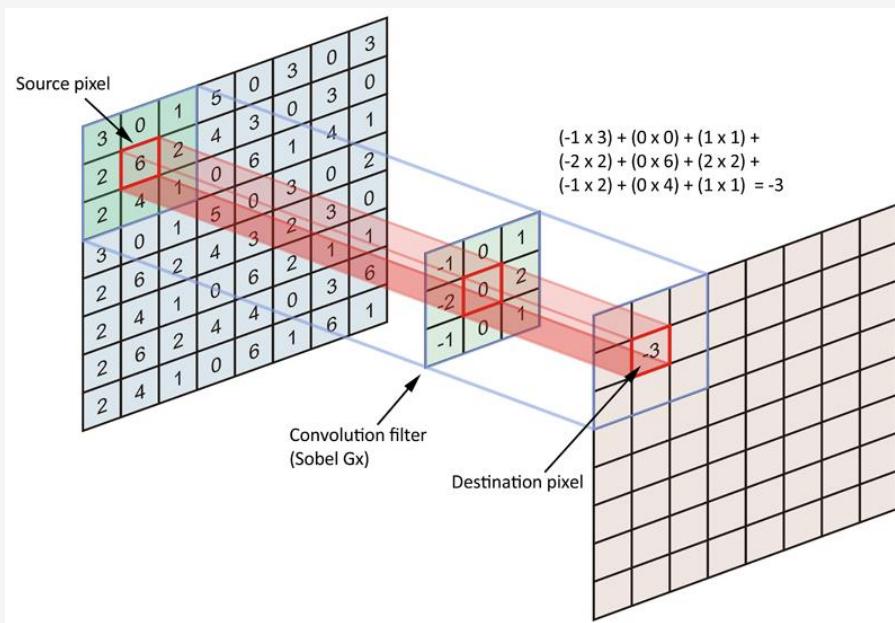
Convolutional Neural Networks



Convolutional Neural Networks

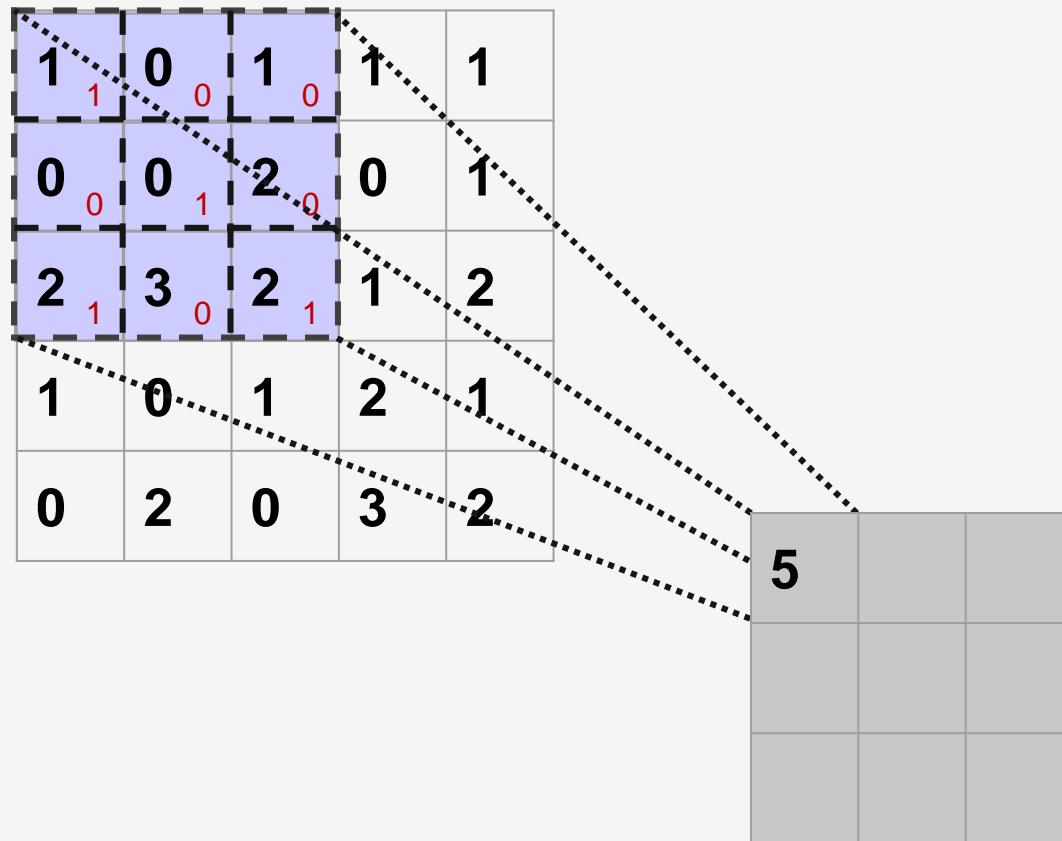
Convolution as core mechanism of CNNs

Convolution \equiv Cross-Correlation



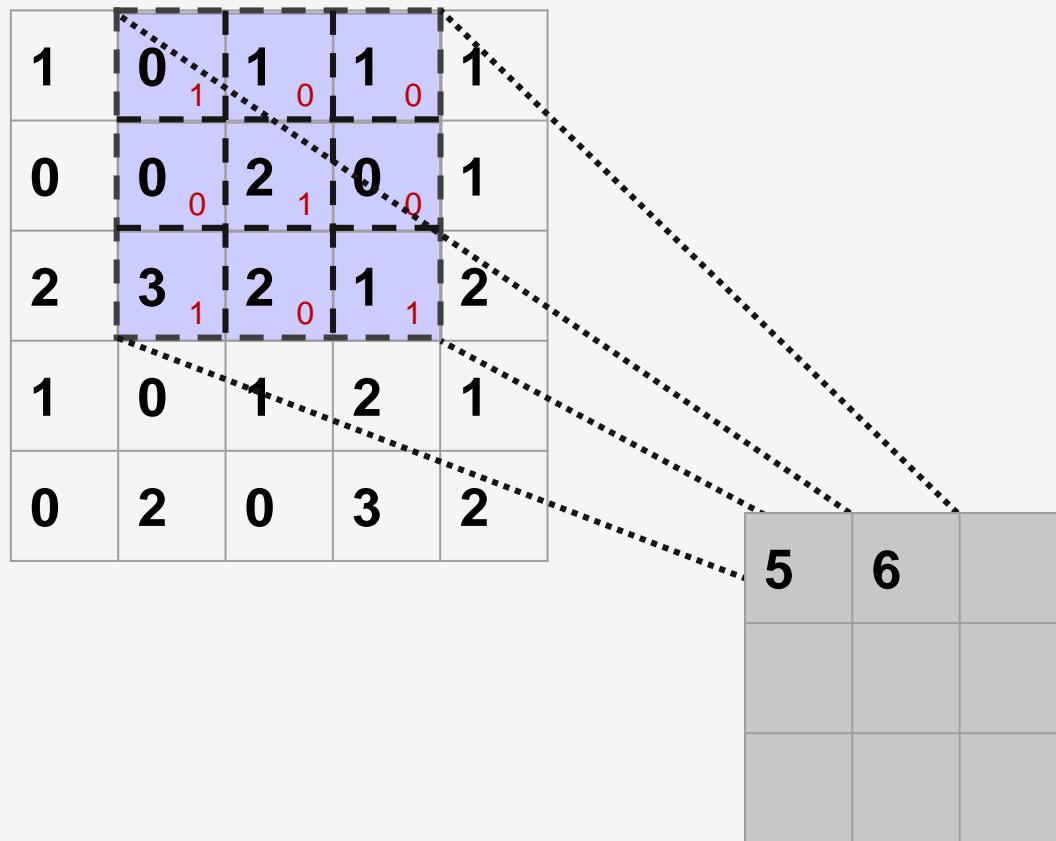
Convolutional Neural Networks

Convolution as core mechanism of CNNs



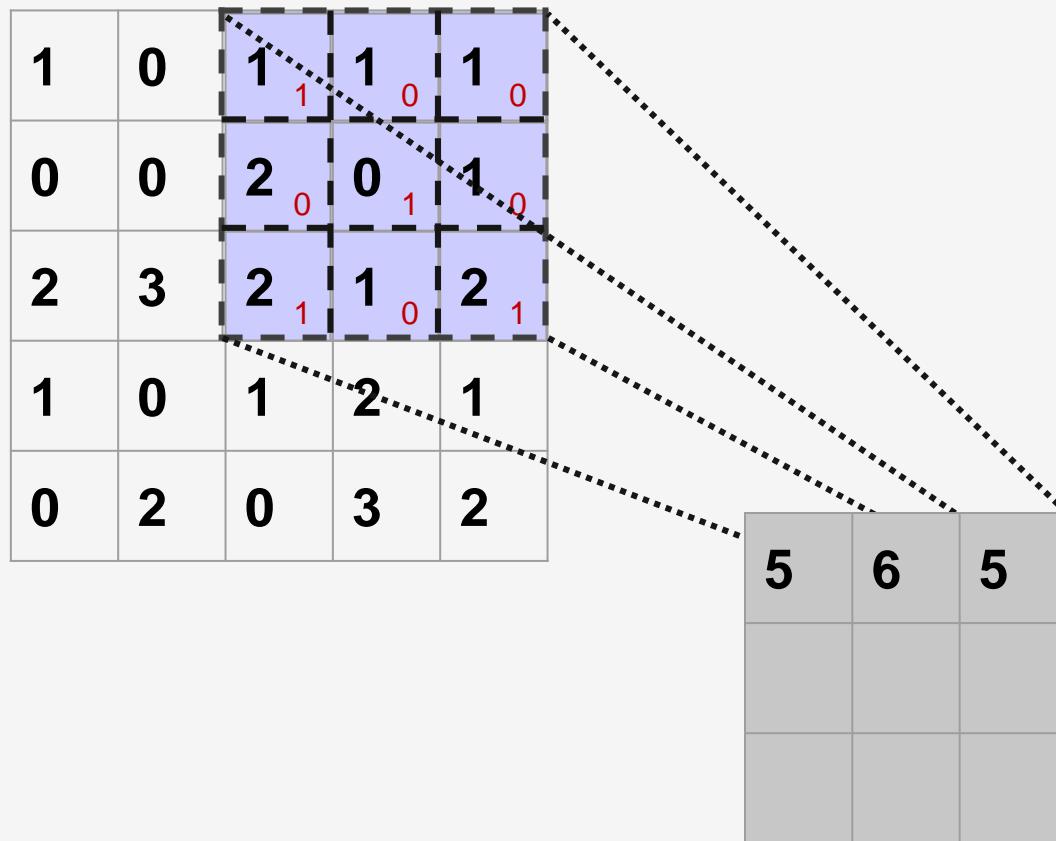
Convolutional Neural Networks

Convolution as core mechanism of CNNs



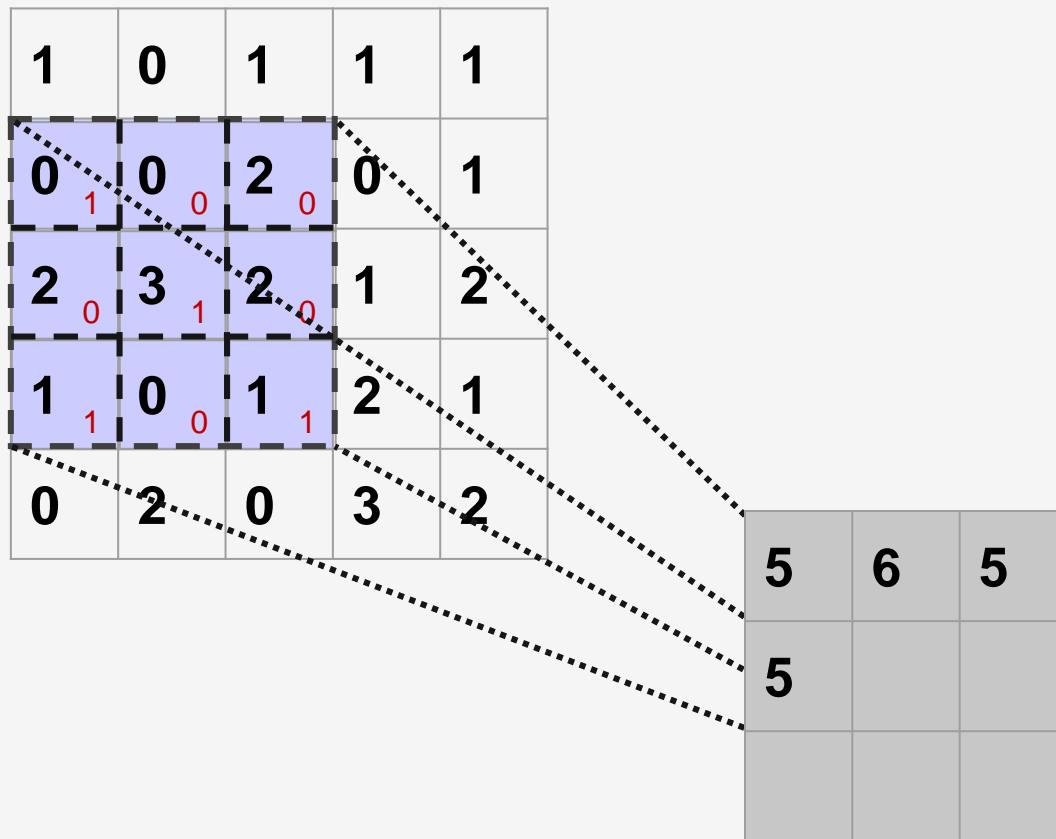
Convolutional Neural Networks

Convolution as core mechanism of CNNs



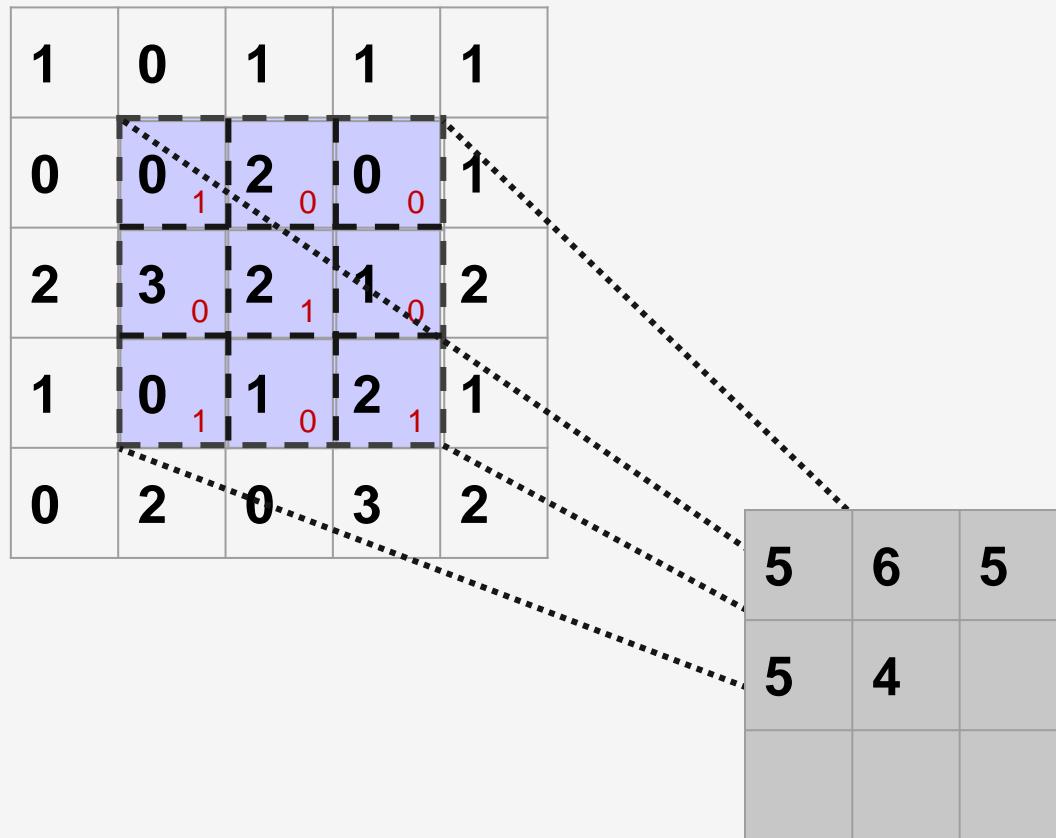
Convolutional Neural Networks

Convolution as core mechanism of CNNs



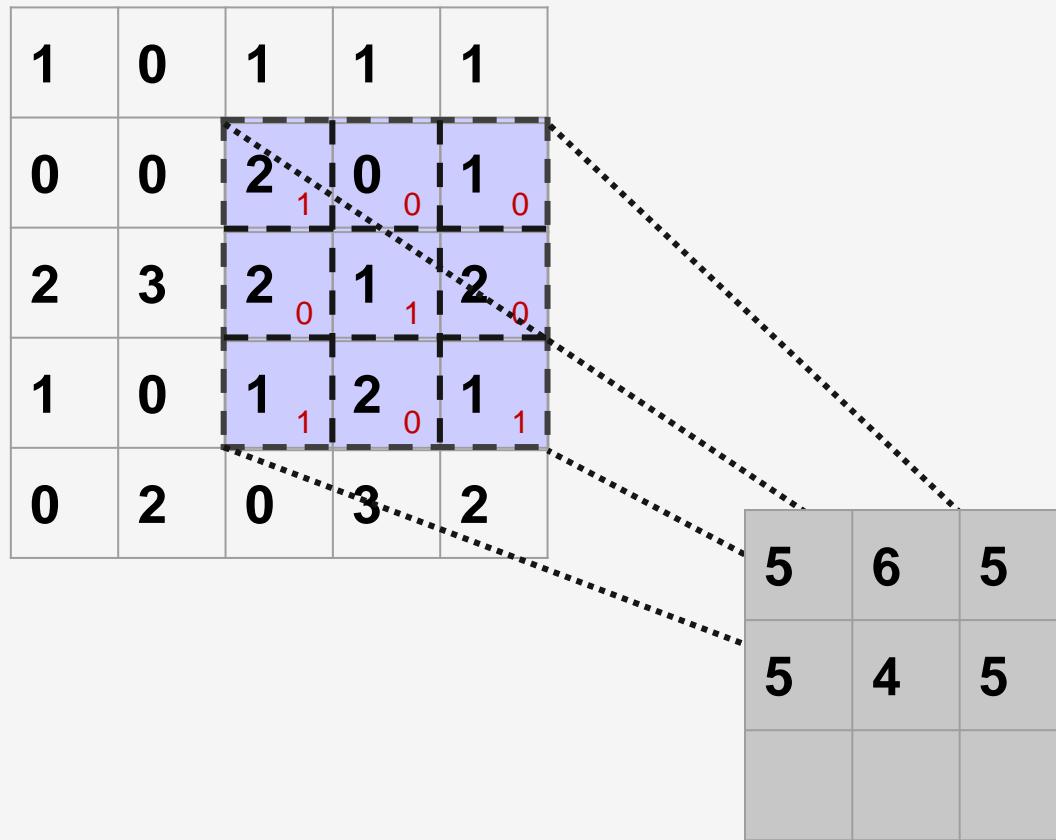
Convolutional Neural Networks

Convolution as core mechanism of CNNs



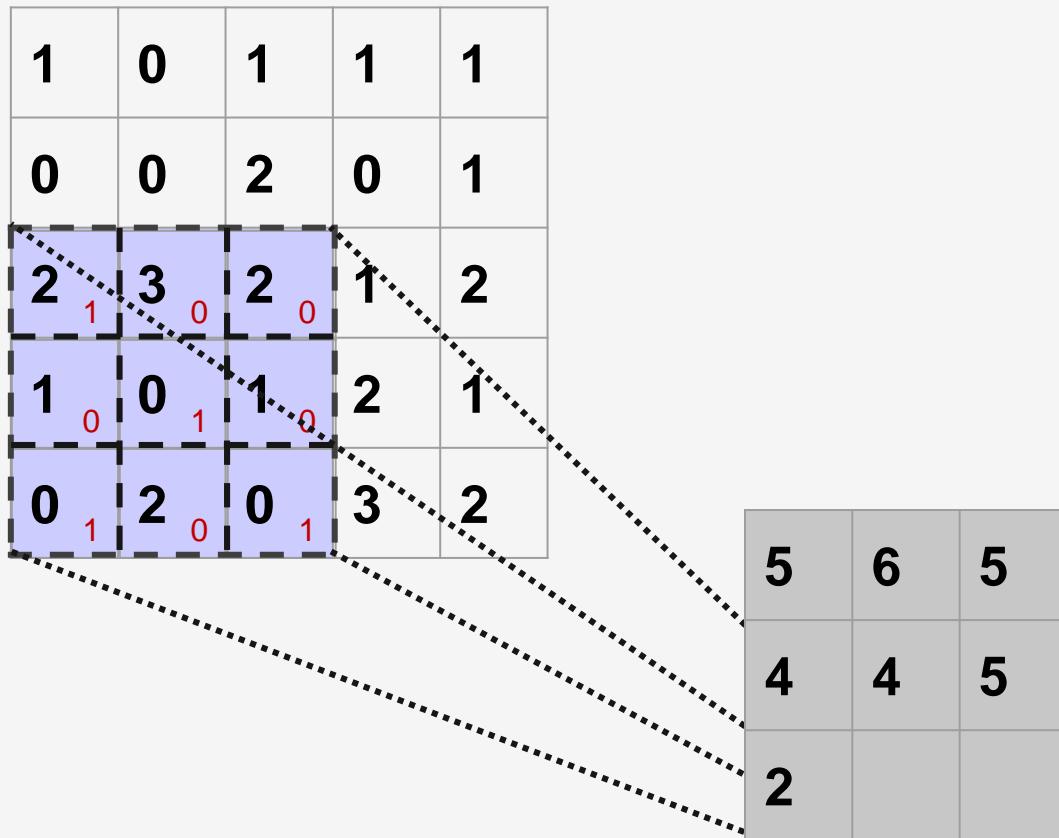
Convolutional Neural Networks

Convolution as core mechanism of CNNs



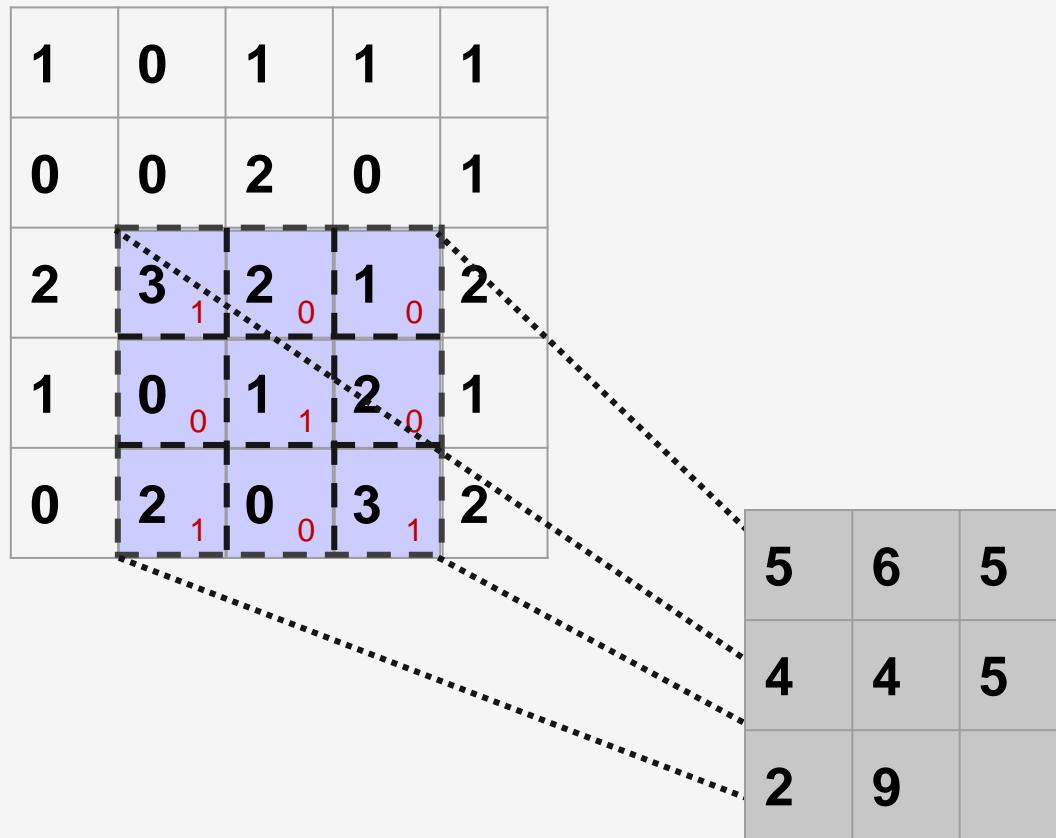
Convolutional Neural Networks

Convolution as core mechanism of CNNs



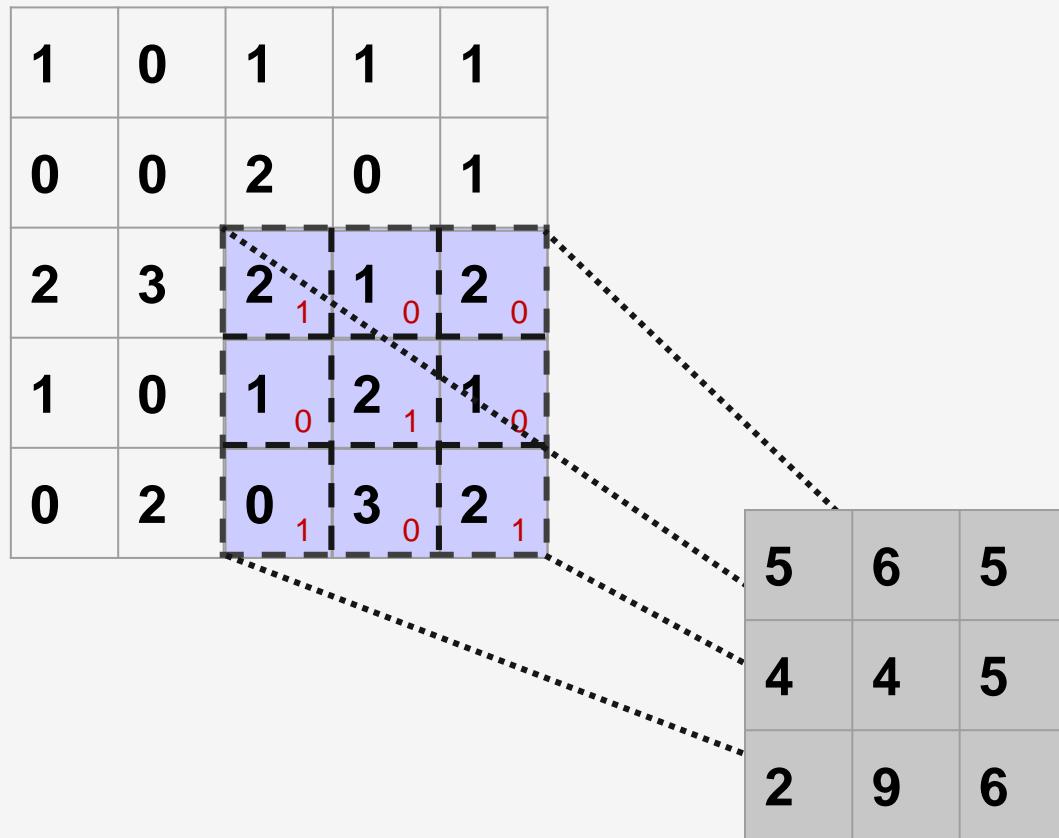
Convolutional Neural Networks

Convolution as core mechanism of CNNs



Convolutional Neural Networks

Convolution as core mechanism of CNNs



Convolutional Neural Networks

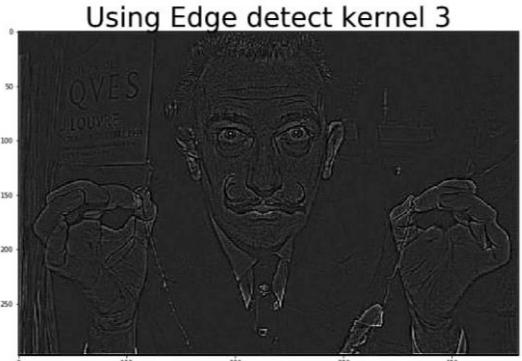
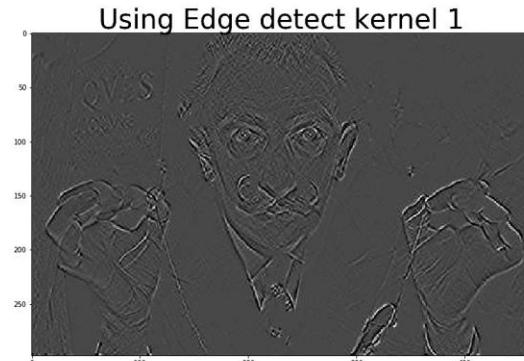
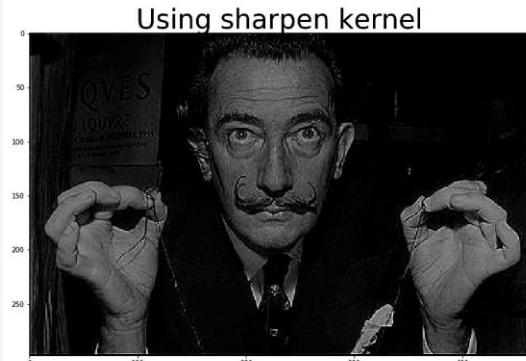
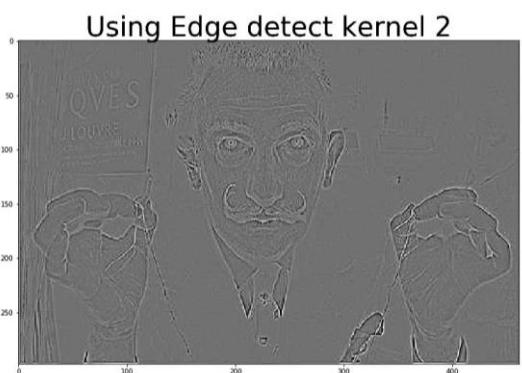
Convolution as core mechanism of CNNs

- Notion of **moving average**
- Represents a **filter**
- Importance of *shape* and *weights* of the filter matrix (kernel)
- Each kernel extracts a specific kind of features (edge, color, shape ...)

Convolutional Neural Networks

Convolution as core mechanism of CNNs

Using different kernels



Convolutional Neural Networks

CNN

- Particular kind of deep forward networks
→ **Weights** and **bias** to learn as well
- Convolutional layers filter inputs for useful information
- They automatically adjust to extract the information useful for the Computer Vision task (classification, object detection, ...)
- Basic CNN are based on three main ideas:
 - Local receptive fields
 - Shared weights and bias
 - Pooling layers

Convolutional Neural Networks

CNN – Basics

Input image

1	0	1	1	1
0	0	2	0	1
2	3	2	1	2
1	0	1	2	1
0	2	0	3	2

Local receptive field

MaxPooling

Shared weights
and bias

Output after
convolution

Output of MaxPooling

6	

Convolutional Neural Networks

CNN – Basics

Input image

1	0	1	1	1
0	0	2	0	1
2	3	2	1	2
1	0	1	2	1
0	2	0	3	2

Local receptive field

MaxPooling

Shared weights
and bias

Output after
convolution

Output of MaxPooling

5	6	5
4	4	5
2	9	6

6	6

Convolutional Neural Networks

CNN – Basics

Input image

1	0	1	1	1
0	0	2	0	1
2	3	2	1	2
1	0	1	2	1
0	2	0	3	2

Local receptive field

Shared weights
and bias

MaxPooling

5	6	5
4	4	5
2	9	6

Output after
convolution

6	6
9	

Output of MaxPooling

Convolutional Neural Networks

CNN - Basics

Input image

1	0	1	1	1
0	0	2	0	1
2	3	2	1	2
1	0	1	2	1
0	2	0	3	2

Local receptive field



MaxPooling

6	6
9	9

Shared weights
and bias



Output after
convolution

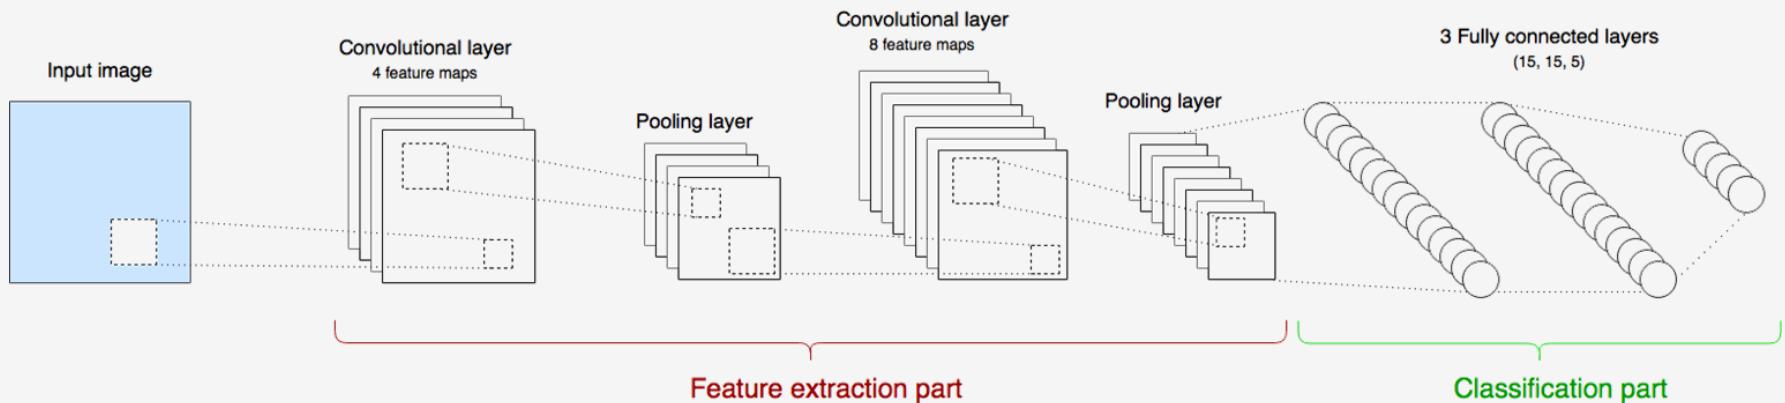
Output of MaxPooling

5	6	5
4	4	5
2	9	6

Convolutional Neural Networks

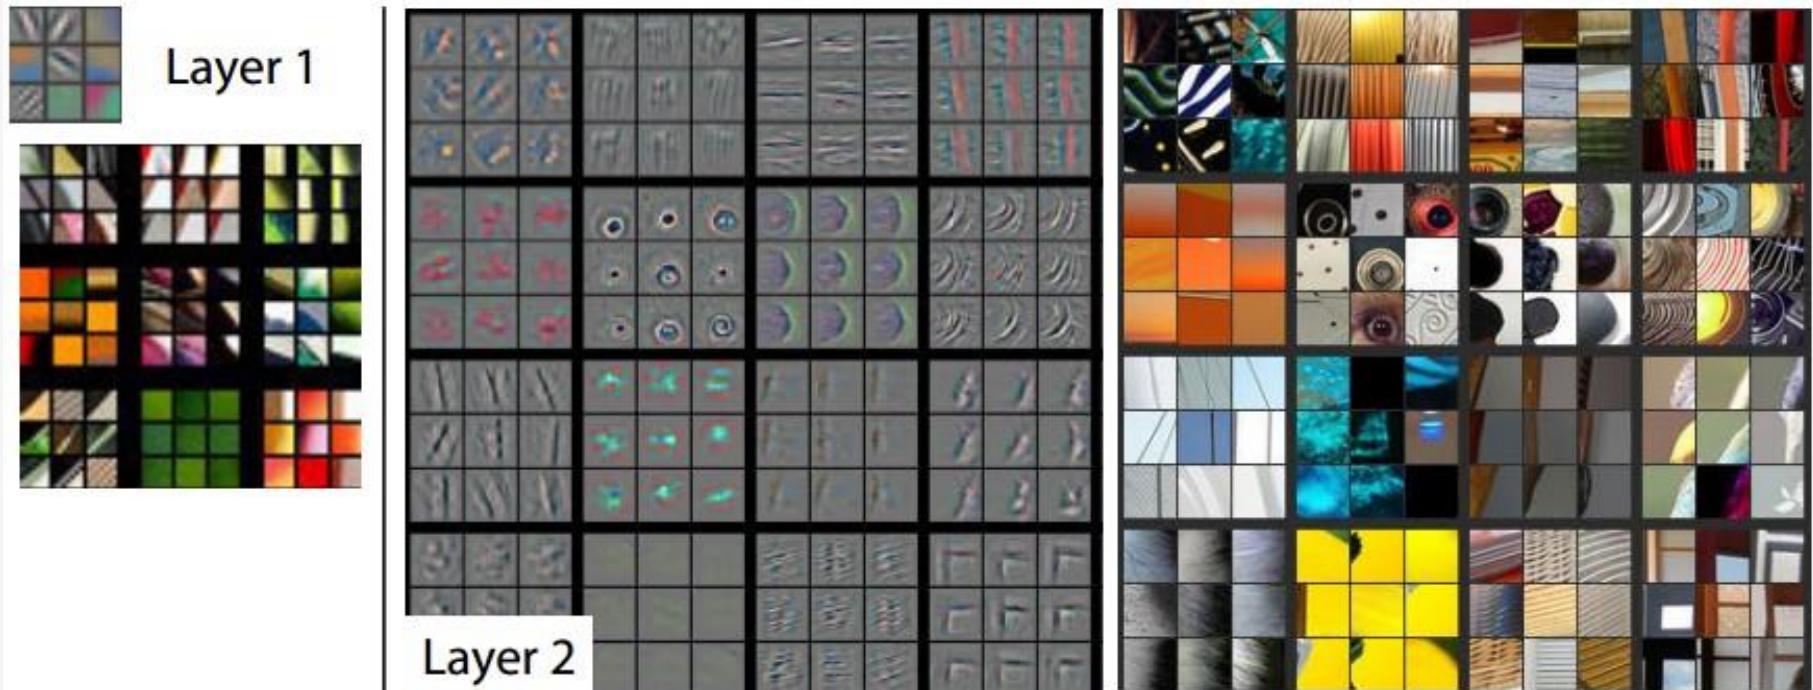
CNN

Illustration of a small CNN



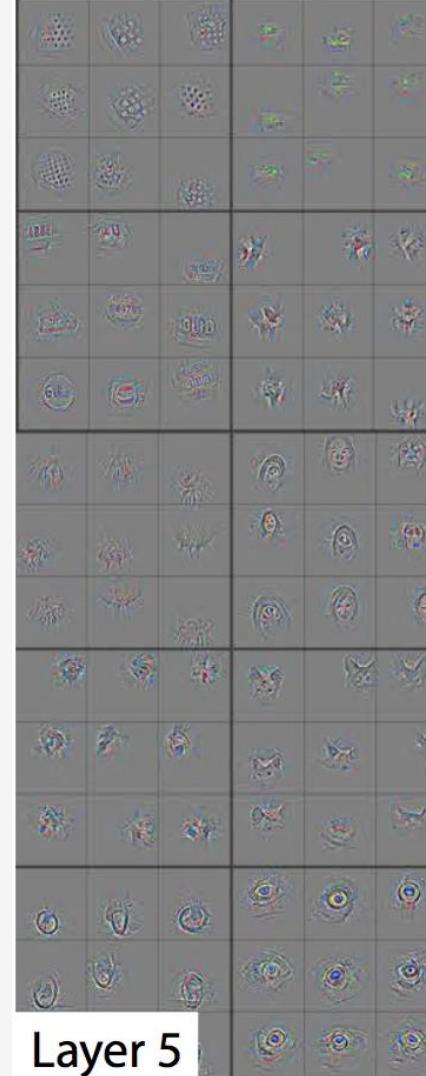
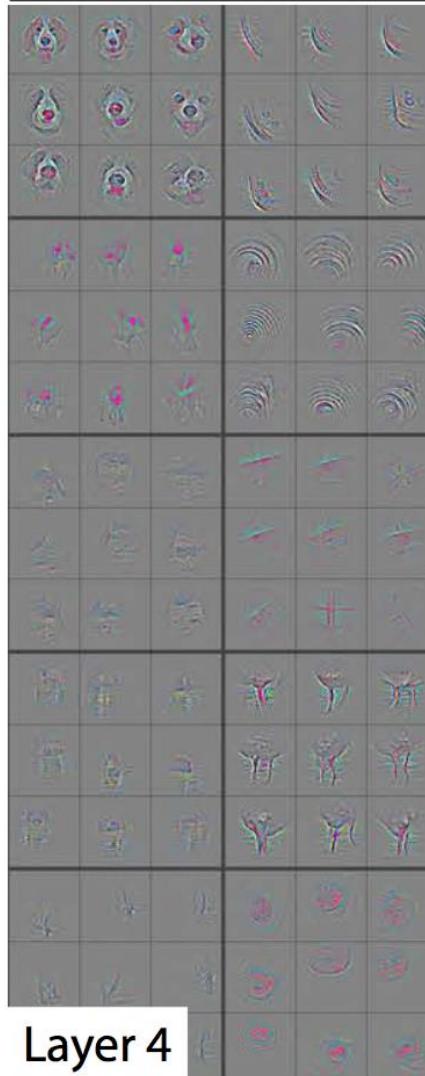
Convolutional Neural Networks

CNN – illustration of what layers see



Convolutional Neural Networks

CNN – illustration of what layers see



Convolutional Neural Networks

Famous CNN architectures



Large Vision Recognition Challenge (ILSVRC 2010 – 2017)

<http://www.image-net.org/>

1000
Object
classes

+ 1,2 M
Training
images

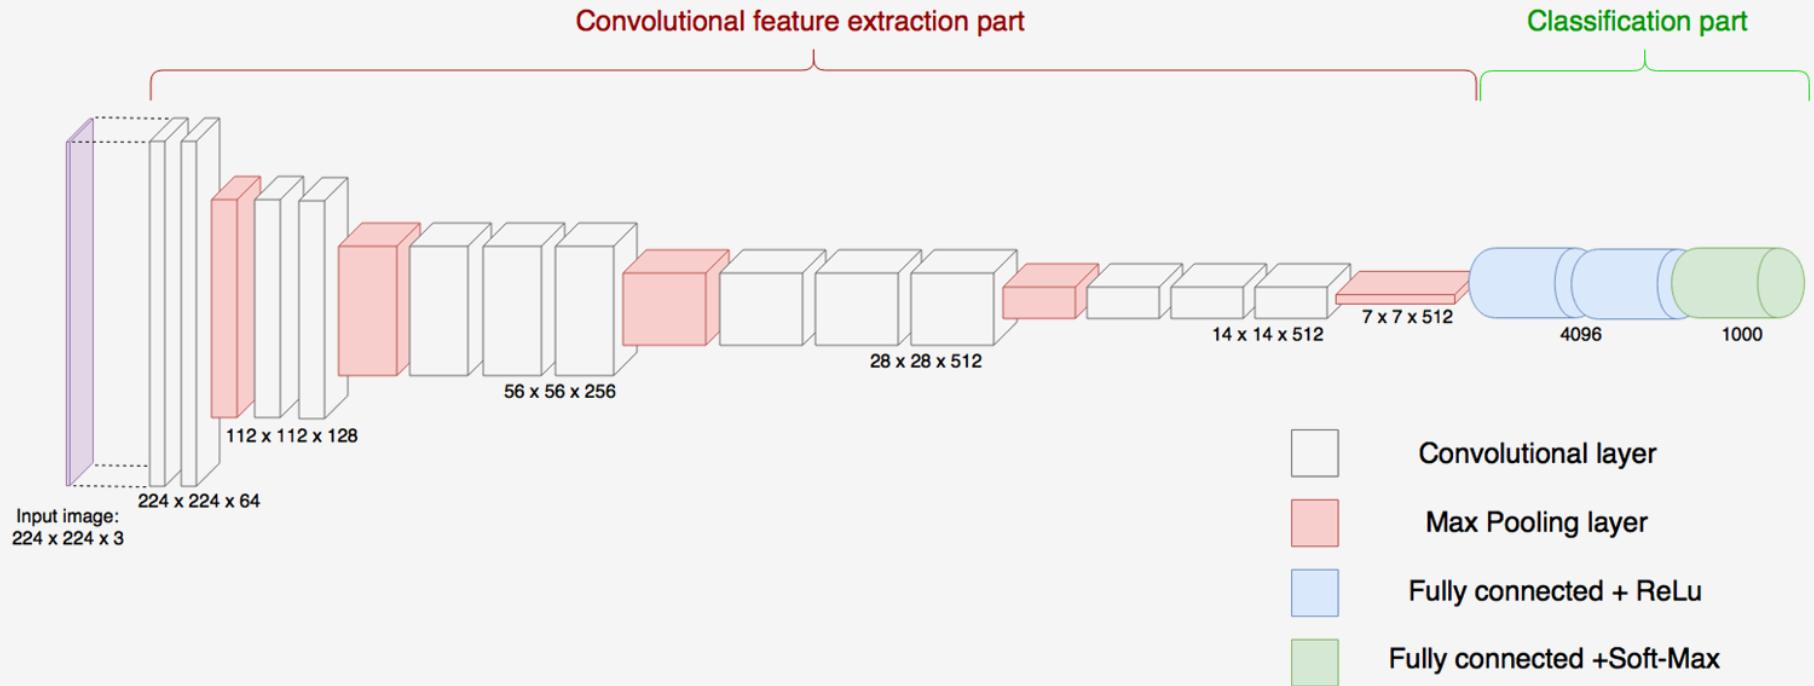
+ 100 k
Test
images

Convolutional Neural Networks

Famous CNN architectures

Very Deep CNN – VGG16

top-5 test error of around 8% (2014)

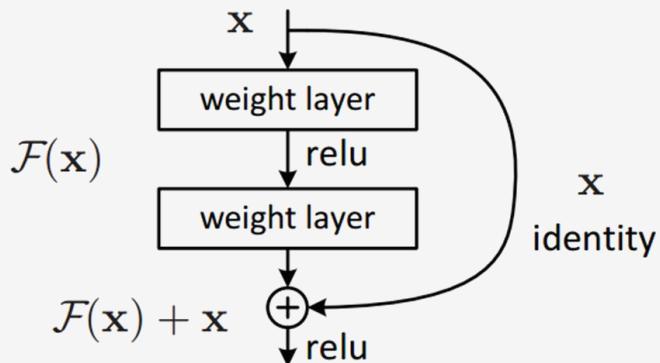


Convolutional Neural Networks

Famous CNN architectures

Residual Networks - ResNet

top-5 test error of around 3.57 % (2015)



- with 152 layers it is **8 times** deeper than VGG16
- Residual Learning intuition: multiple non-linear function can asymptotically approximate the residual function $F(x) = H(x) - x$
- It is then possible to fit: $H(x) := F(x) + x$
- Ease training of **critically Deep models**
- Tackles the degradation of the learning process we usually meet with critically deep networks

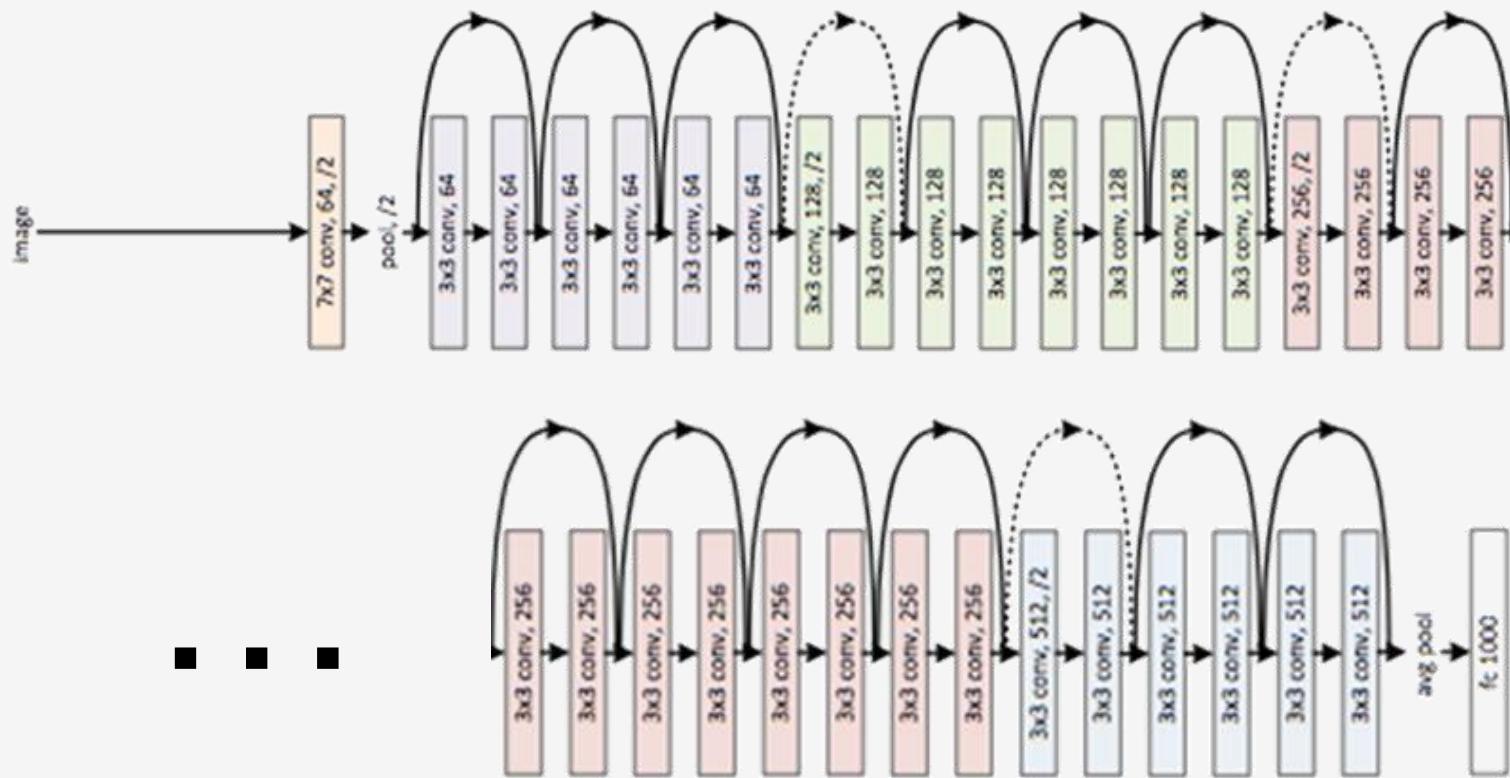
Convolutional Neural Networks

Famous CNN architectures

Residual Networks - ResNet

top-5 test error of around 3.57 % (2015)

34-layer residual



Convolutional Neural Network

Building a CNN

```
from keras.layers import Dense, MaxPooling2D, Conv2D, Flatten

def deep_conv_model():
    model = Sequential()
    model.add(Conv2D(64, (5,5), input_shape = (28,28,1), activation = 'relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

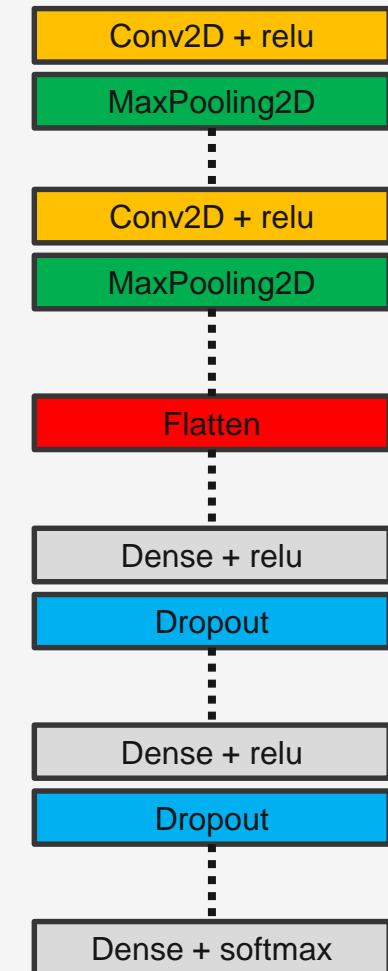
    model.add(Conv2D(64, (5, 5),activation = 'relu'))
    model.add(MaxPooling2D(pool_size=(2, 2)))

    model.add(Flatten())
    model.add(Dense(128, activation='relu'))
    model.add(Dropout(0.25))
    model.add(Dense(64, activation='relu'))
    model.add(Dropout(0.25))
    model.add(Dense(10, activation='softmax'))

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model

model = deep_conv_model()
model.summary()
```



Study case n°2

Design, train and predict

MNIST
Handwritten Digit
Recognizer

Using CNN



Transfer Learning



Transfer Learning

General intuition

"Inspired by human beings' capabilities to transfer knowledge across tasks, transfer learning aims to leverage knowledge from a source domain to improve the learning performance or minimize the number of labeled examples required in a target domain." (Wei et al. (2014))

If able to recognize lions and tigers



After training the
classification part of the
CNN

Transfer Learning approach

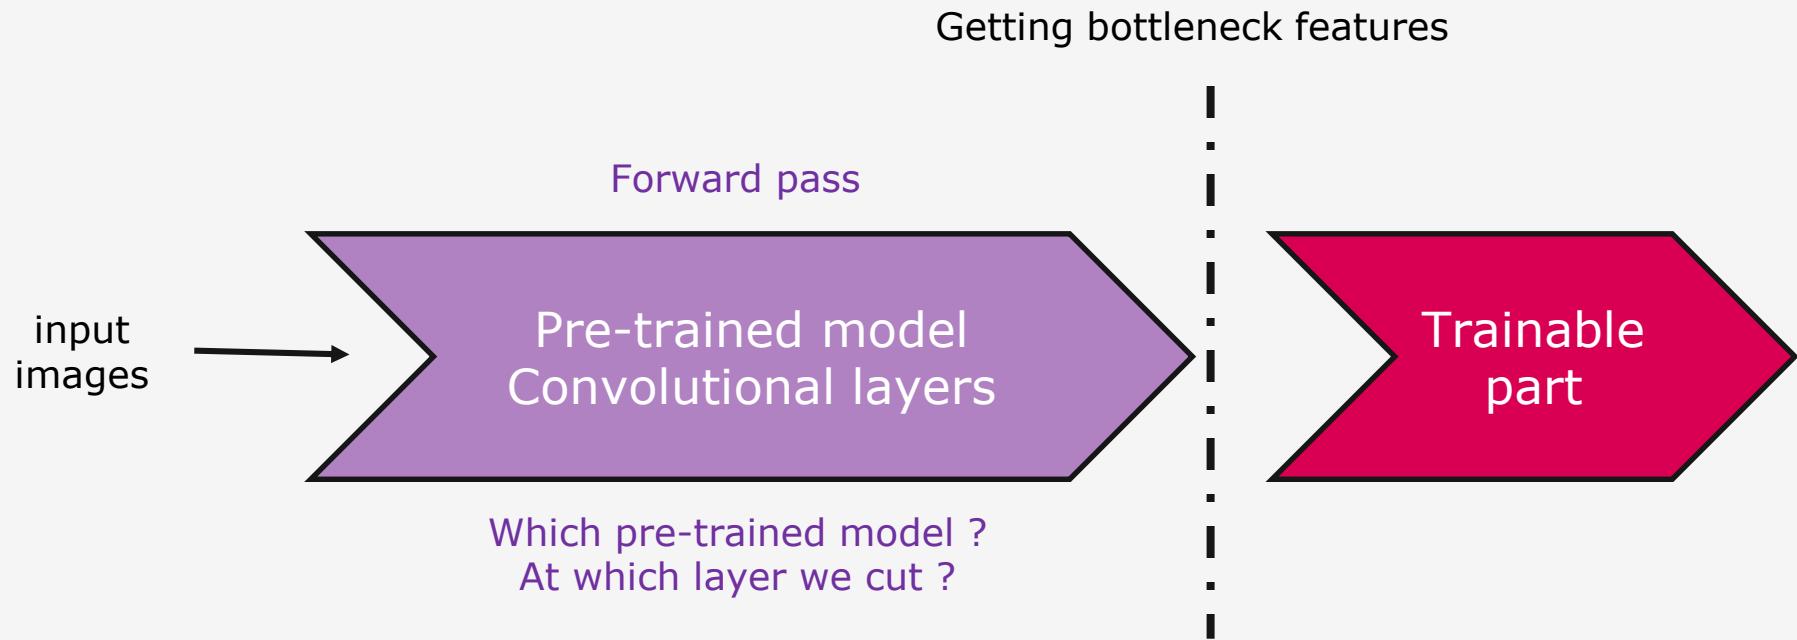
You can also recognize cats and dogs



Transfer Learning

General Scheme

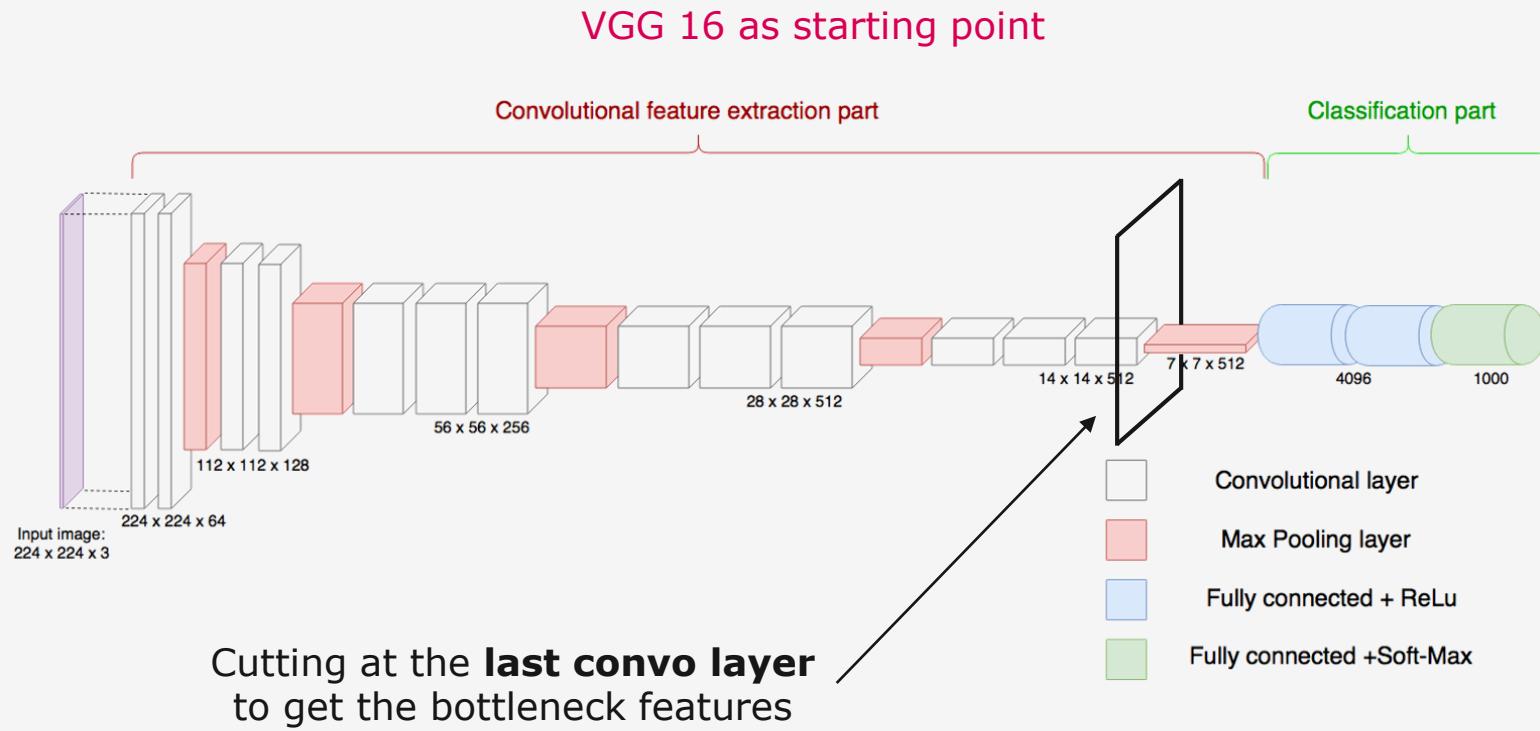
Transfer knowledge from one domain-to-another
in the case of CNN



Transfer Learning

General Scheme

- Models famous for having outperformed state-of-the-art results in ImageNet competition
- Weights for VGG16, VGG19, Inception, ResNet50 and Xception are available in keras



Study case n°3

Design, train and predict

MNIST
Handwritten Digit
Recognizer

Using Transfer
Learning



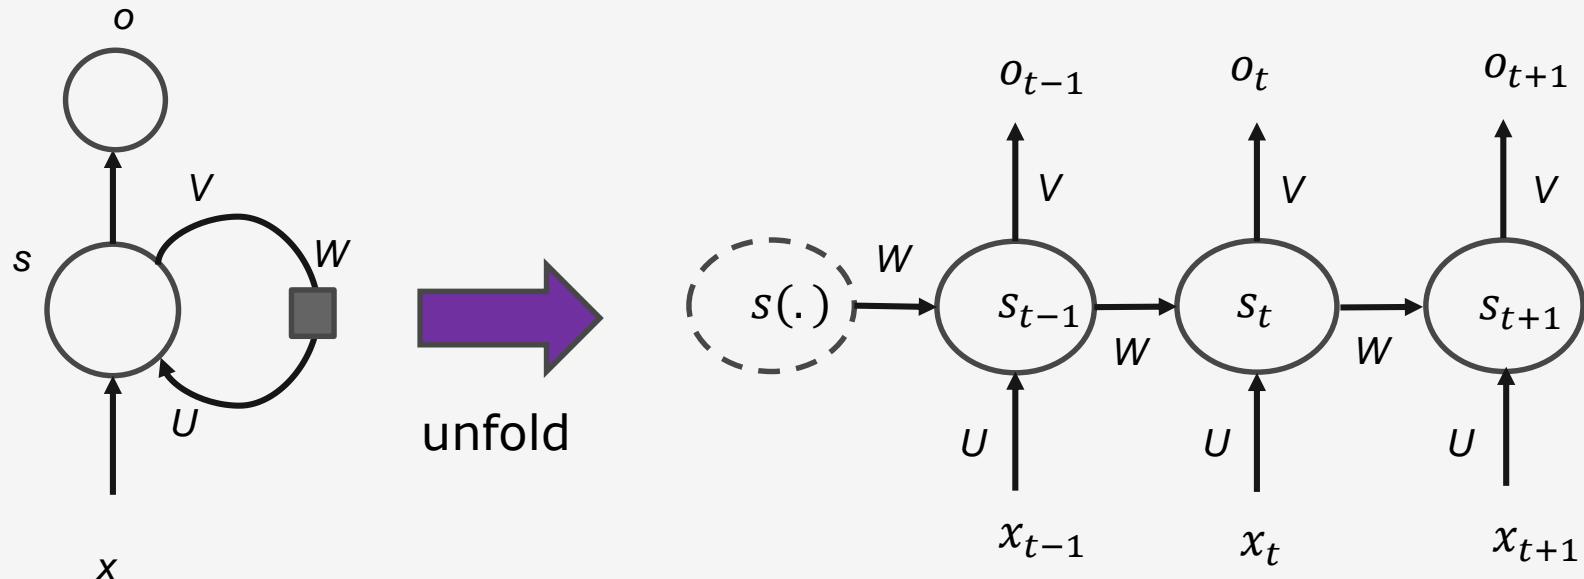
Recurrent Neural Networks



Recurrent Neural Networks

What's a RNN?

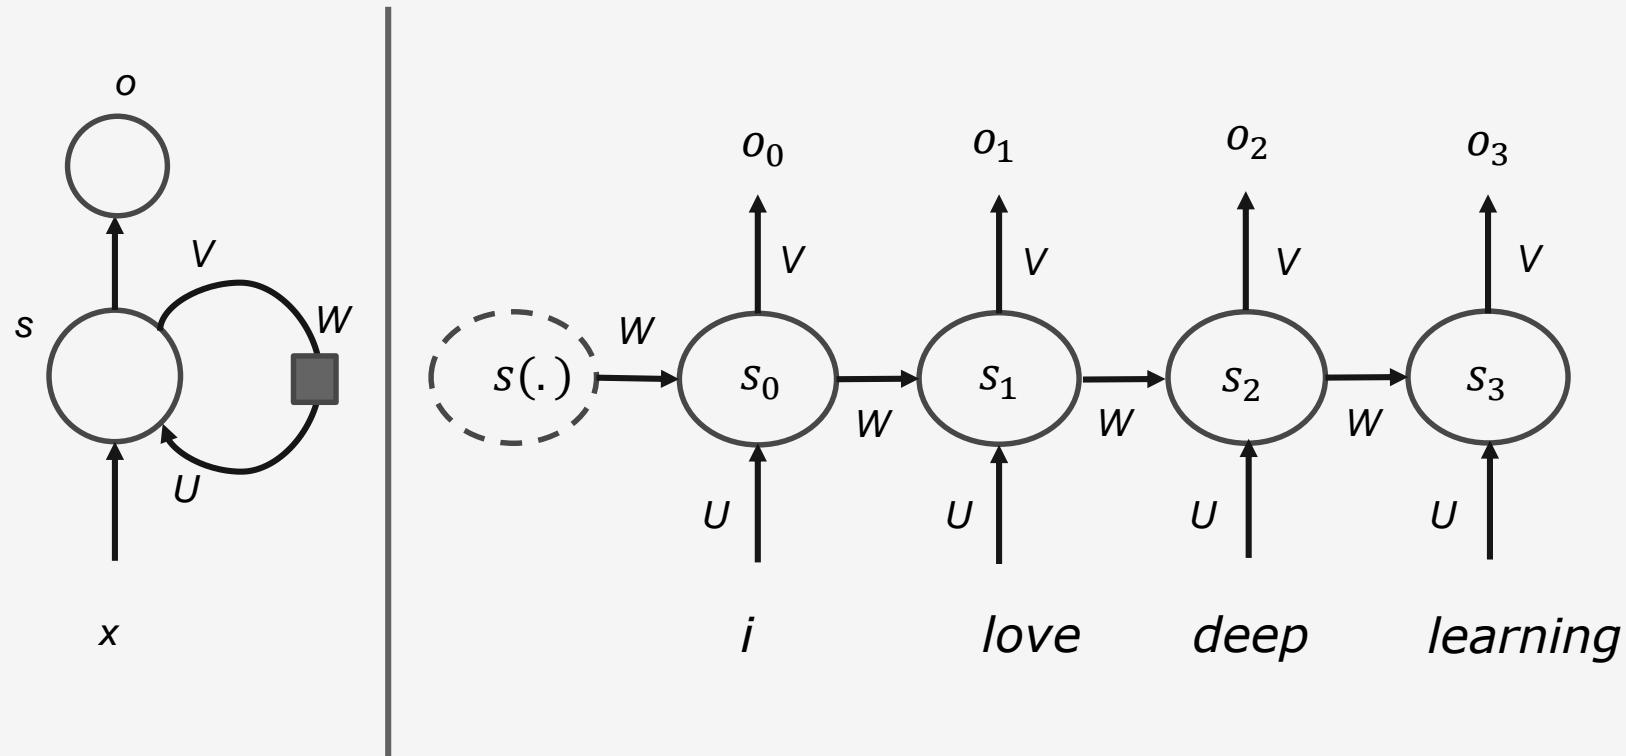
- Adapted to sequence treatment tasks: NLP, etc ...
- Inputs are not independent anymore → notion of memory



Recurrent Neural Networks

Illustration: example with a sentence

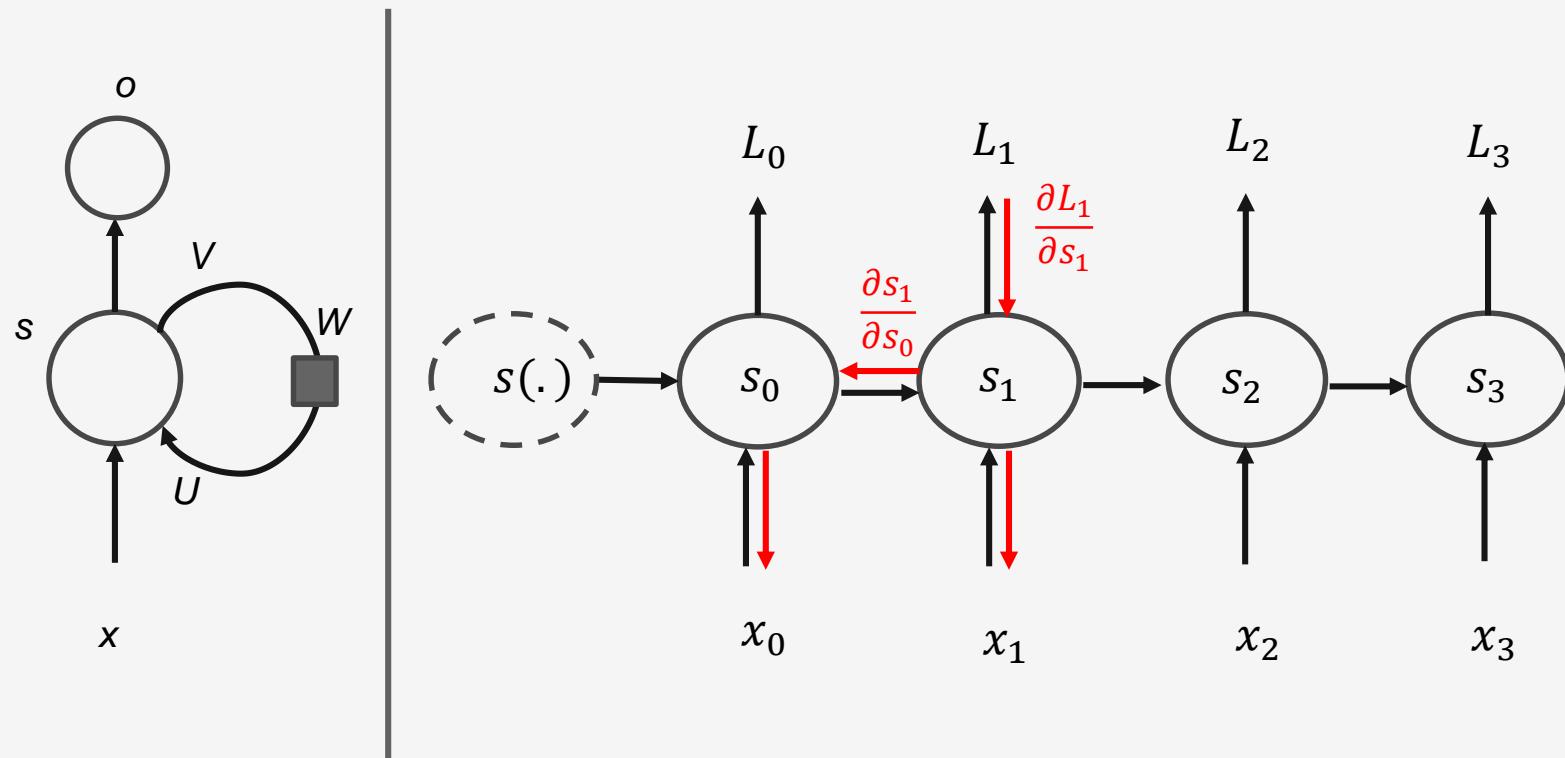
Sentence: «i love deep learning »



Training a Recurrent Neural Networks

Back-Propagation Through Time (BPTT)

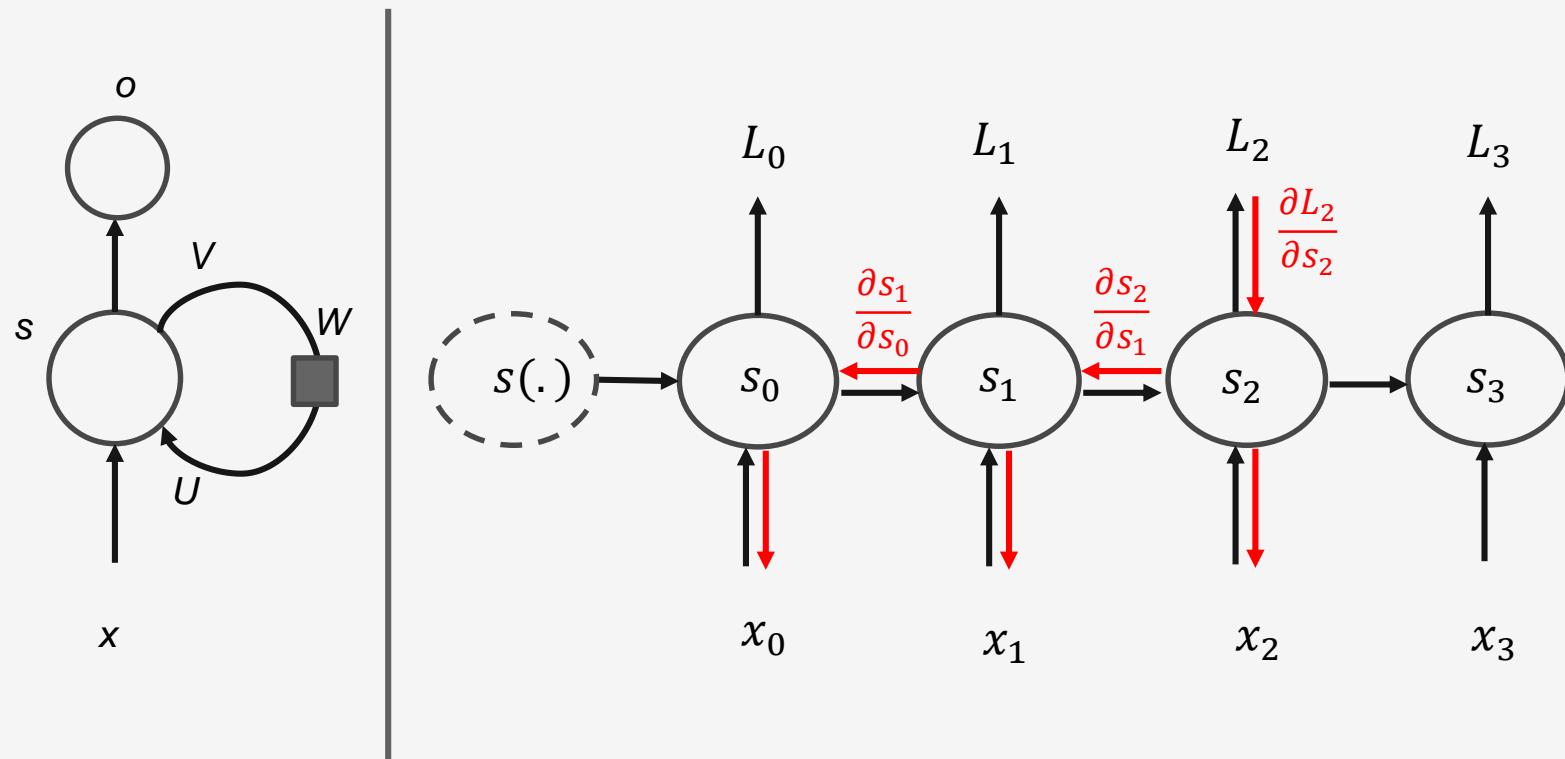
Sentence: «i love deep learning »



Training a Recurrent Neural Networks

Back-Propagation Through Time (BPTT)

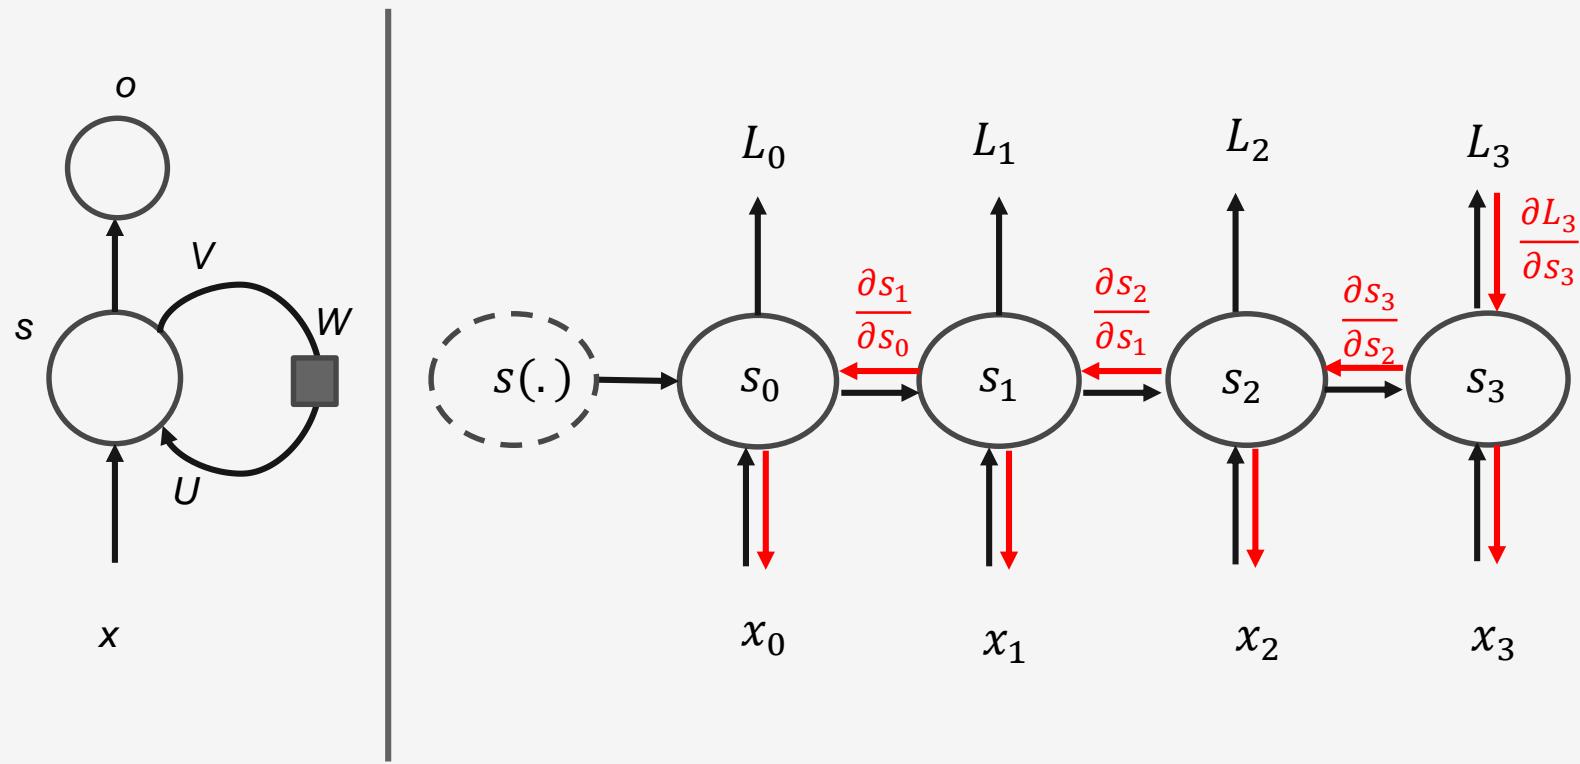
Sentence: «i love deep learning »



Training a Recurrent Neural Networks

Back-Propagation Through Time (BPTT)

Sentence: «i love deep learning »



Complexified RNNs

Gated Recurrent Units (GRU)

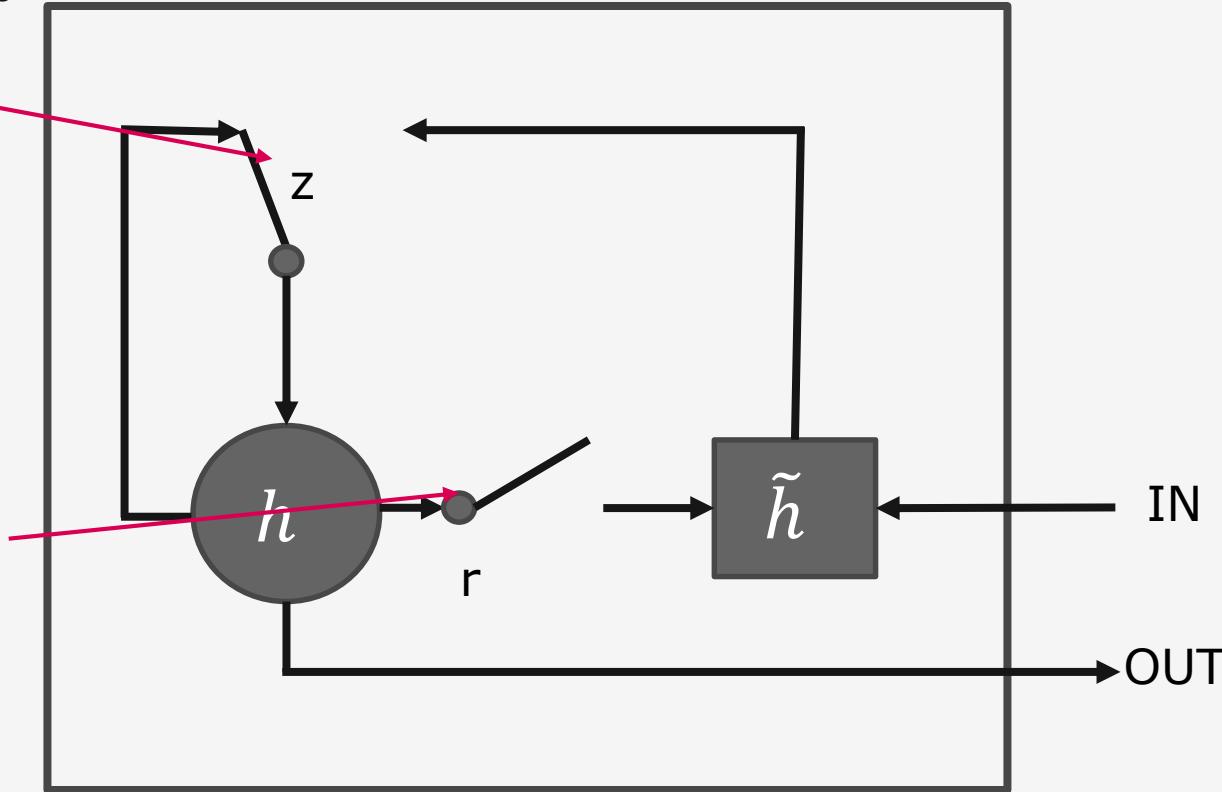
Z = update gate

How much of the previous memory to keep around

r = reset gate

How much of the new input to combine with the previous memory

Illustration of a GRU



Complexified RNNs

Long Short Term Memory (LSTM)

Illustration of a GRU

f = forget gate

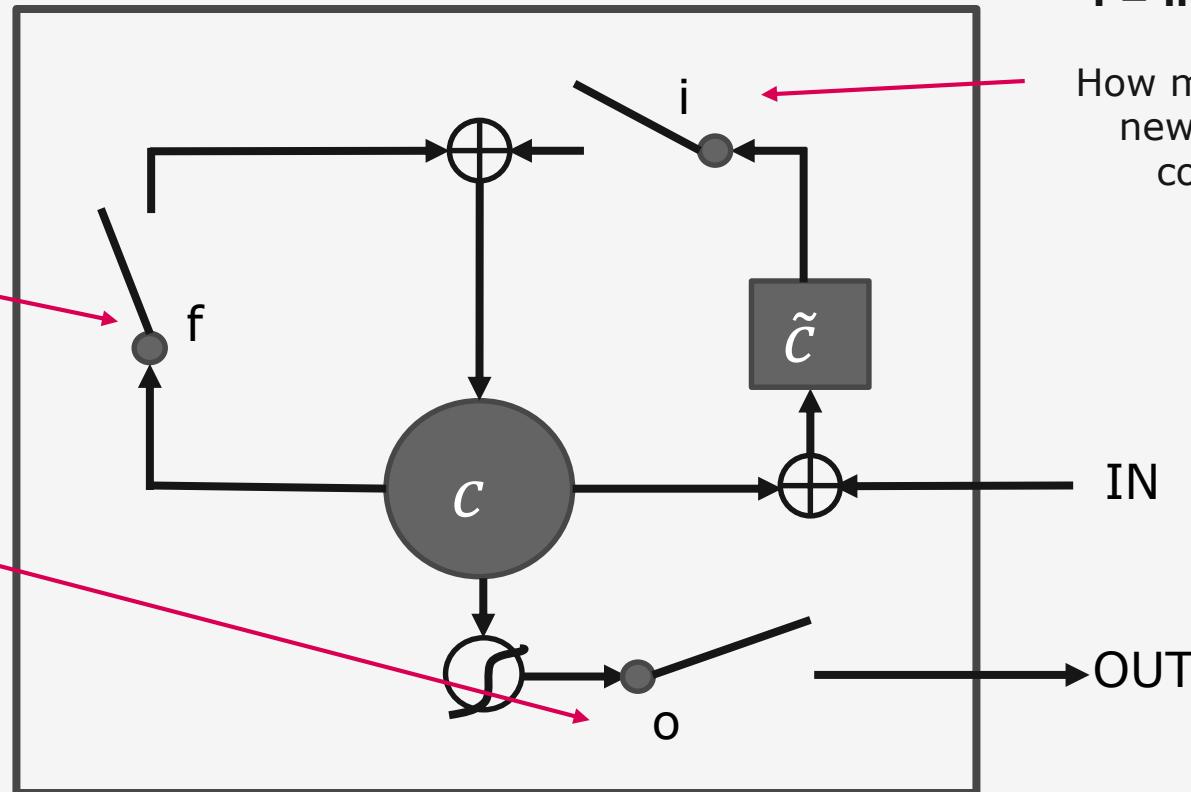
How much of the previous memory to forget

o = output gate

Computes the output hidden state s

i = input gate

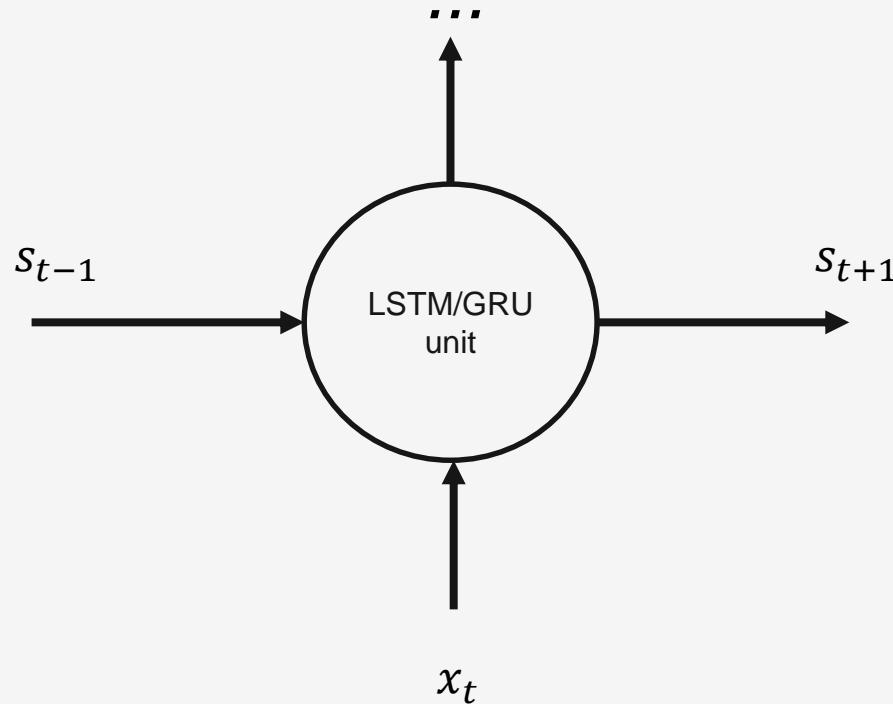
How much of the new state to consider



Complexified RNNs

LSTM/GRU units

You can also stick to the
simplest version !!!



Complexified RNNs

Sentiment Analysis

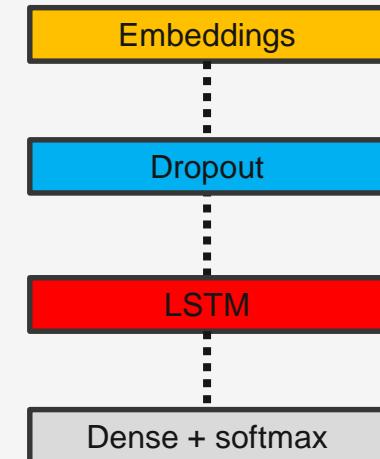
```
from keras.layers import Dense, MaxPooling2D, Conv2D, Flatten

def sentiment_analysis_model():
    model = Sequential()
    model.add(Embedding(vocab_size=2000, 5, input_length=500))
    model.add(Dropout(0.2))
    model.add(LSTM(50))
    model.add(Dense(1, activation='sigmoid'))

    model.compile(optimizer='adam',
                  loss='binary_crossentropy',
                  metrics=['accuracy'])

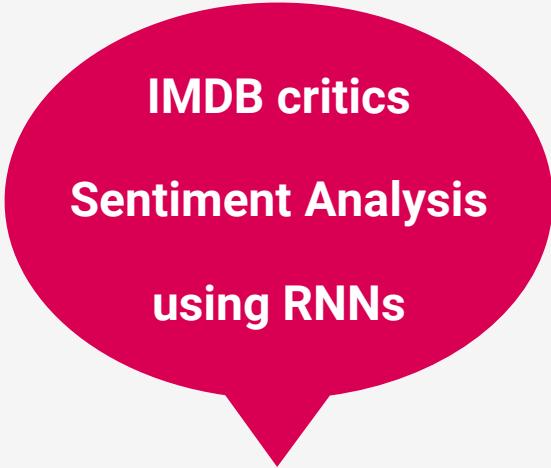
    return model

model = sentiment_analysis_model()
model.summary()
```



Study case n°4

Design, train and predict



IMDB critics
Sentiment Analysis
using RNNs



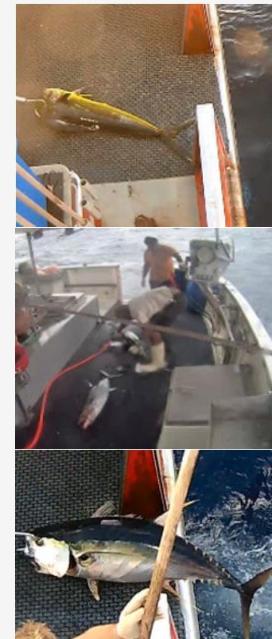


4

Practice:
Design and train
your own
network



-  ALB: Albacore tuna (*Thunnus alalunga*)
-  BET: Bigeye tuna (*Thunnus obesus*)
-  DOL: Dolphinfish, Mahi Mahi (*Coryphaena hippurus*)
-  LAG: Opah, Moonfish (*Lampris guttatus*)
-  SHARK: Various: Silky, Shortfin Mako
-  YFT: Yellowfin tuna (*Thunnus albacares*)





5

Conclusion

Conclusion

- Useful in AI
- The most powerful tool in the rest years in the AI environment
- Evaluating very fast ... need to constantly update the knowledge

Tips

- Courses online available:
 - <http://course.fast.ai/> , free course given by Jeremy Howard
 - <https://fr.coursera.org/specializations/deep-learning> course given by Andrew NG
- Books:
 - “The Elements of Statistical Learning”, by T. Hastie, R. Tibshirani and J. Friedman
 - “Deep Learning”, by I. Goodfellow, Y. Bengio and A. Courville
- Research papers at: <http://www.arxiv-sanity.com/>

Thanks

It was a big pleasure to animate this course!
Sincere thanks!



Thank you

